

Ch3-3

ILP:

Branch Prediction
& Hardware-based
speculation

Review

❑ Structural Hazard

- Separate/replicate function unit, fully-pipelined function units

❑ Data Hazard

- RAW: double-bump, forwarding path, simple compiler scheduling
- WAR, WAW: implicit register renaming

❑ Control Hazard

- Resolve branch asap, Flushing, Predict-not-taken, Predict-taken, Delayed-Branch

Control Hazard (example P_{A-25})

❑ **Flushing:** stall until the branch is resolved.

- Stall X cycles if delay slots=x

❑ **Predict-not-taken:**

- Hardware takes all branches as not-taken, compiler need to put frequent case in the not-taken branch
- Stall 0 cycles if not-taken, stall X cycles if taken

❑ **Predict-taken:**

- Hardware takes all branches as taken, can go forward only when get the branch target . Compiler need to put frequent case in the taken branch
- Stall y cycles waiting for branch target; stall X cycles if not taken.

❑ **Delayed Branch:**

Static branch prediction

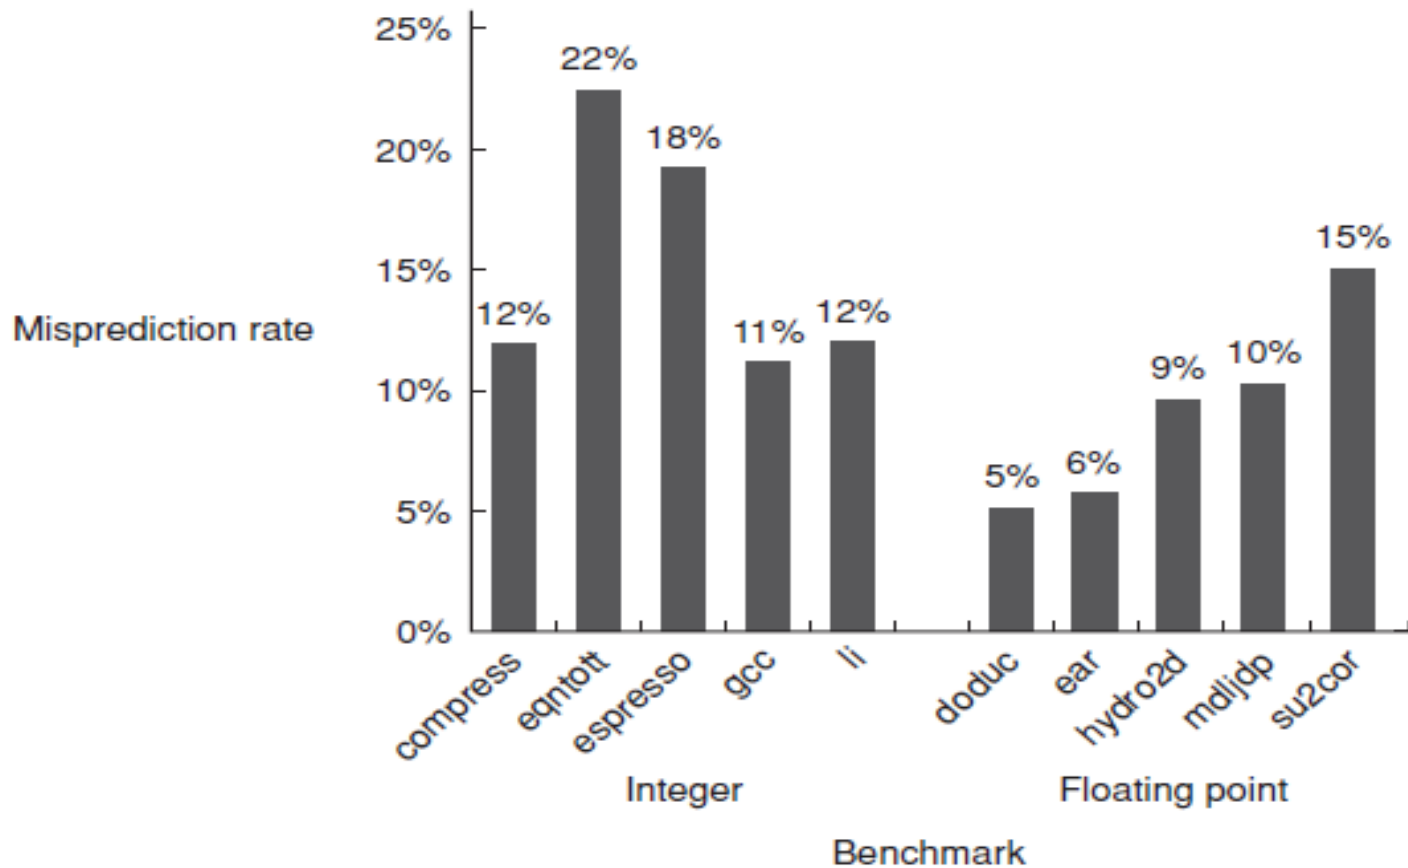
❑ Stall for control hazards

= f (branch frequency, prediction accuracy, misprediction penalty)

❑ Predict according to statistics: (predict taken)

➤ Misprediction: 9% ~ 59%, average 34%

Predict based on profile information



Average misprediction rate 9% for FP programs,
Average misprediction rate 15% for integer programs

Dynamic Hardware Prediction

--reducing branch costs

1. 1-bit Branch-Prediction Buffer
2. 2-bit Branch-Prediction Buffer
3. Correlating Branch Prediction Buffer
4. Tournament Branch Predictor
5. Branch Target Buffer
6. Integrated Instruction Fetch Units
7. Return Address Predictors

1-bit Branch-Prediction Buffer

❑ Performance = $f(\text{accuracy, cost of misprediction})$

❑ Branch History Table: Lower bits of PC address
index table of 1-bit values

- Says whether or not branch taken last time
- No address check (saves HW, but may not be right branch)


❑ Problem: in a loop, 1-bit BHT will cause
2 mispredictions (avg $2/n$ iterations before exit):

- Only $(n-2)/n$ accuracy even if loop $(n-1)/n$ of the time
- 0111.....1110111... ..1110111.....1110111...

2-bit Branch-Prediction Buffer

❑ Solution: 2-bit scheme where change prediction only if get misprediction *twice*:

➤ (Jim Smith, 1981, [A study of branch prediction strategies](#))



❑ Red: stop, not taken


❑ Blue: go, taken

❑ Adds *hysteresis* to decision making process

➤ 01111111.....111110111111...1110111111

An alternative 2-bit predictor

- It can be implemented using a 2-bit counter




Generalize the 2-bit predictor to n-bit predictor


- ❑ N-bit counter ($0 \leq x < 2^n - 1$)
- ❑ Predict taken $X \geq 2^{n-1}$
- ❑ Predict untaken $X < 2^{n-1}$
- ❑ Counter + 1 when a taken branch
- ❑ Counter - 1 when a not taken branch
- ❑ Most system rely on 2-bit predictor rather than n-bit predictor.

Accuracy of the 2-bit predictor

4096 entries, 2-bit predictor



How about unlimited entries ?



How to improve accuracy ?

❑ Analyze the branch behavior

```
IF ( aa == 2)          SUBI   R3, R1, #2
    aa = 0;            BNEZ   R3, L1      ; br.b1
IF ( bb == 2)          (aa!=2)
    bb = 0;            ADD    R1, R0, R0 ; aa==0
IF (aa == bb ) L1: SUBI   F3, R2, #2
    {                  BNEZ   R3, L2      ; br.b2
    .....             (bb!=2)
    }                  ADD    R2, R0, R0 ; bb==0
                        L2: SUB   R3, R1, R2; R3=aa-bb
                        BEQZ   R3, L3      ; br.b3
```

Conclusion from the example

- ❑ B3 branch depends on b1 and b2. B3 is taken only when both b1 and b2 is untaken.
- ❑ It's impossible to predict B3 correctly only depending on its previous behavior.
- ❑ A branch's behavior not only rely on it's previous behavior but also on recent behavior of other branches.

Let's check an example

```
If (d==0)
    d=1;
if (d==1){
    .....
}
```

❑ Assume $\text{Reg}[R1] = d$

```
BNEZ      R1, L1      ; br b1, (d!=0)
DADDIU    R1, R0, #1  ; d==0, so d=1
L1: DADDIU  R3, R1, #-1 ;
BNEZ      R3, L2      ; br b2, (d!=1)
.....
L2:
```

Feature of the code:

| d的初值 | d==0? | B1 | 在b2以前的d值 | d==1? | b2 |
|------|-------|-----------|----------|-------|-----------|
| 0 | Yes | Not taken | 1 | Yes | Not taken |
| 1 | No | Taken | 1 | Yes | Not taken |
| 2 | No | Taken | 2 | no | Taken |

Conclusion: b2 will be **not taken** if **b1 is not taken**

So, we can get accurate prediction if using correlate predictor.

Always mispredict if using one-bit predictor

| d=? | b1 预测 | b1 动作 | 新的b1 预测 | b2 预测 | b2 动作 | 新的b2 预测 |
|-----|----------|----------|------------|----------|----------|------------|
| 2 | NT | T | T | NT | T | T |
| 0 | T | NT | NT | T | NT | NT |
| 2 | NT | T | T | NT | T | T |
| 0 | T | NT | NT | T | NT | NT |

预测与实际动作
总是相反

预测与实际动作
总是相反

Correlating predictor

❑ Correlating predictor is composed of two parts (two bits):

- 1st bit: prediction for previous br. is NT,
- 2nd bit: prediction for previous br. Is T

❑ Four combinations

| 预测组合 | 上一次Br为NT,预测本次位(看第一位) | 上一次Br为T,预测本次为(看第二位) |
|-------|----------------------|---------------------|
| NT/NT | NT | NT |
| NT/ T | NT | T |
| T /NT | T | NT |
| T / T | T | T |

Note, previous Br might be a different br.

Prediction with correlating predictor (1,1) with initial value =NT/NT.

| d=? | b1 预测 | b1 动作 | 新的 b1 预测 | b2 预测 | b2 动作 | 新的 b2 预测 |
|-----|-------|-------|----------|-------|-------|----------|
| 2 | NT/NT | T | T/NT | NT/NT | T | NT/T |
| 0 | T/NT | NT | T/NT | NT/T | NT | NT/T |
| 2 | T/NT | T | T/NT | NT/T | T | NT/T |
| 0 | T/NT | NT | T/NT | NT/T | NT | NT/T |


粗体字表示本次预测




预测出错




预测正确



新预测值



b1动作与b2预测值选用关系



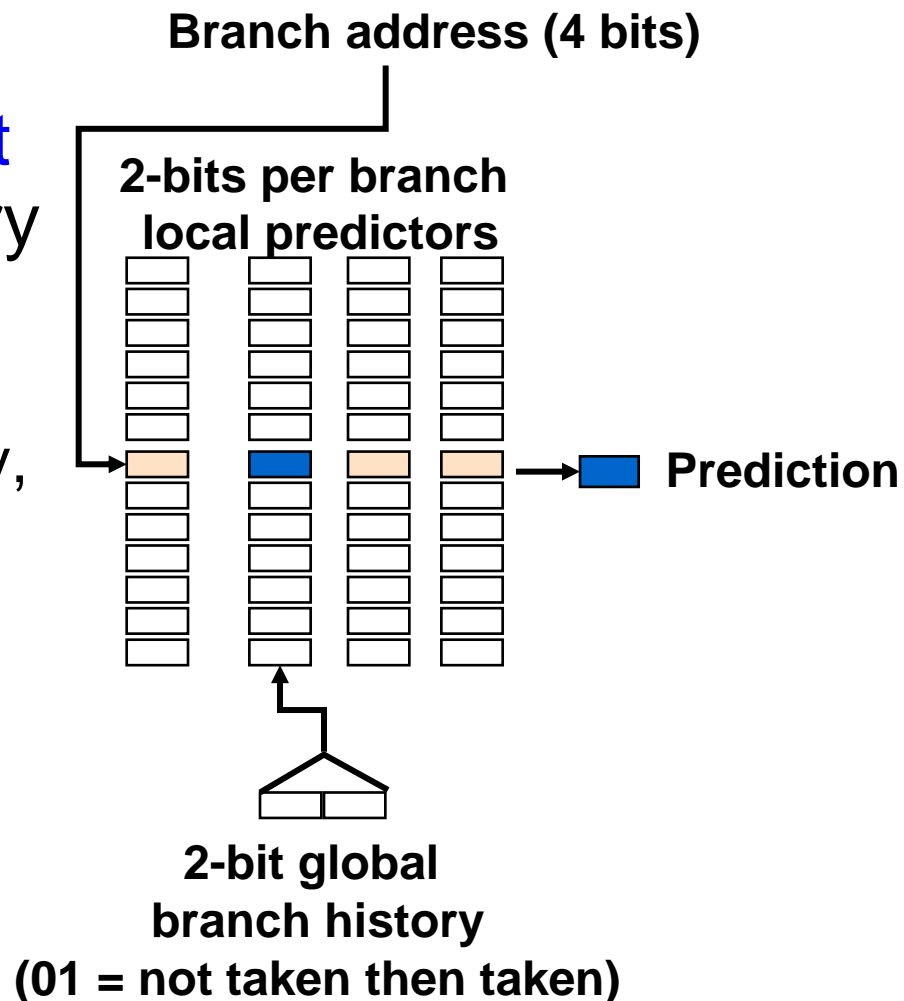
b2动作与b1预测值选用关系

Correlating Branches prediction buffer

❑ Idea: taken/not taken of **recently executed branches** is related to behavior of next branch (as well as the history of that branch behavior)

➤ Then behavior of recent branches selects between, say, 4 predictions of next branch, updating just that prediction

❑ (2,2) predictor: 2-bit global, 2-bit local




4. Correlating Branches prediction buffer

Assume branch history with branch i in blue box.


0 0 0 0 0 1 1 1 0 0 0 0 0 1 1 1 0 0 0 0 0 0 1 1 1

Predict: **v** **v** **x** **v** **v** **x** **v** **v** **v**



- ❑ (m, n) predictor
- ❑ m: previous m branch instruction
- ❑ 2^m

Comparison of 2-bit predictors



Accuracy analysis

❑ Mispredict because either:



- Wrong guess for the branch
- Got branch history of wrong branch when index the table.

4、Tournament Predictors

- ❑ Motivation for correlating branch predictors is 2-bit predictor failed on important branches; by adding global information, performance improved
- ❑ **Tournament predictors**: use 2 predictors, 1 based on global information and 1 based on local information, and combine with a selector
- ❑ Hopes to select right predictor for right branch

Tournament Predictor in Alpha 21264

- ❑ 4K 2-bit counters to choose from among a global predictor and a local predictor
- ❑ **Global predictor** also has 4K entries and is indexed by the history of the last 12 branches; each entry in the global predictor is a standard 2-bit predictor
 - 12-bit pattern: i th bit 0 \Rightarrow i th prior branch not taken;
 i th bit 1 \Rightarrow i th prior branch taken;




T-Predictor in Alpha 21264 (cont.)


Local predictor consists of a 2-level predictor:

- **Top level** a local history table consisting of 1024 10-bit entries; each 10-bit entry corresponds to the most recent 10 branch outcomes for the entry. 10-bit history allows patterns 10 branches to be discovered and predicted.
- **Next level** Selected entry from the local history table is used to index a table of 1K entries consisting a 3-bit saturating counters, which provide the local prediction

Total size: $4K \times 2 + 4K \times 2 + 1K \times 10 + 1K \times 3 = 29K \text{ bits!}$ (~180,000 transistors)



% of predictions from local predictor in Tournament Prediction Scheme



Accuracy of Branch Prediction





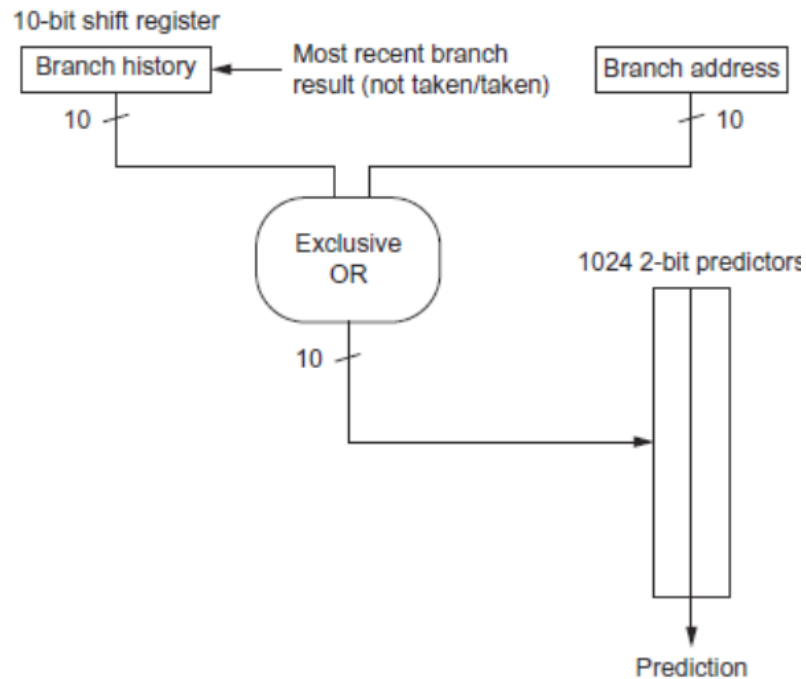
fig 3.40

□ Profile: branch profile from last execution
(static in that is encoded in instruction, but profile)

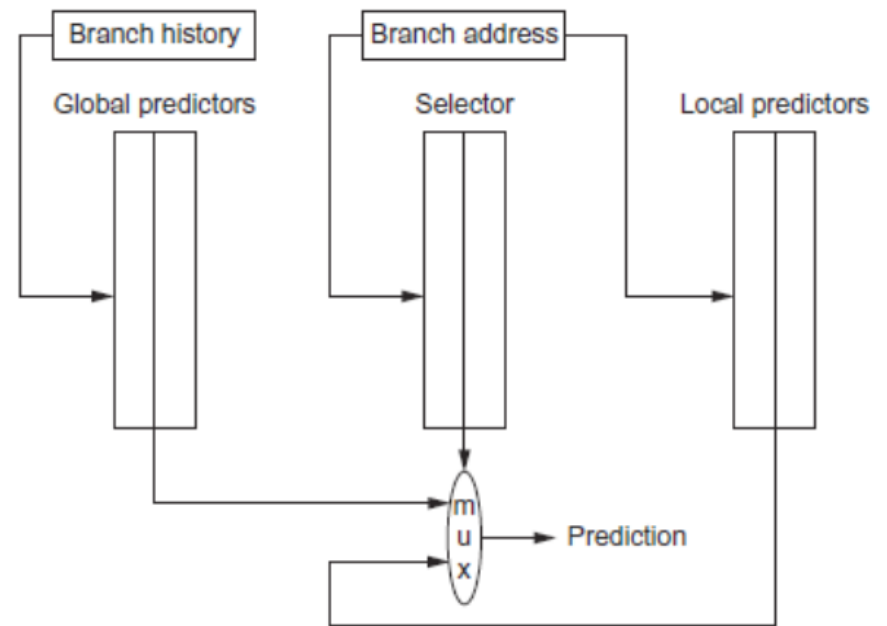
Accuracy v. Size (SPEC89)



Other Branch Prediction

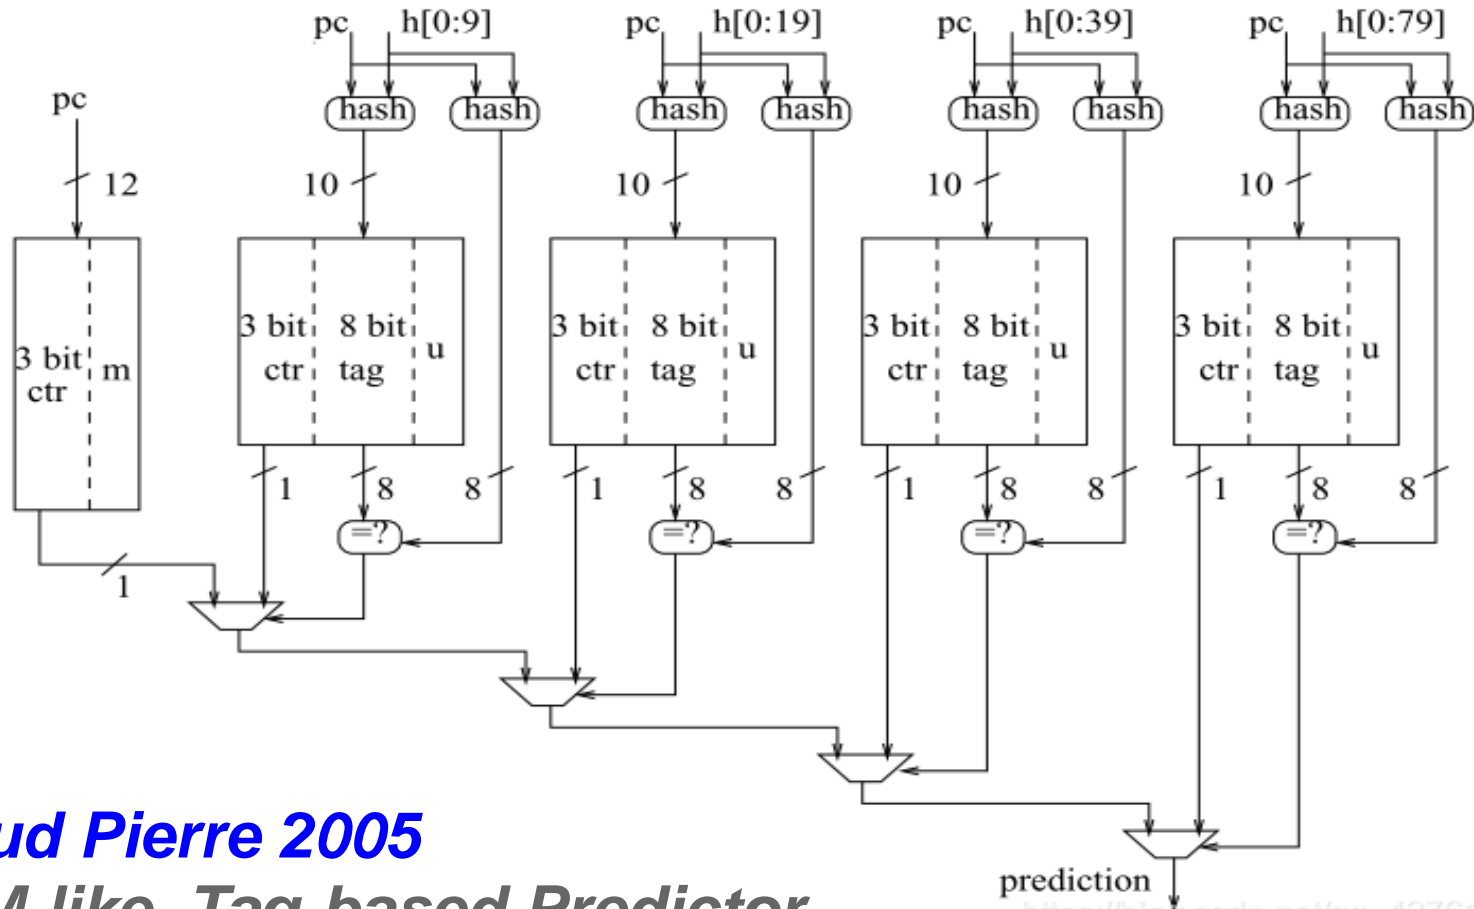


gshare



tournament

PPM (Prediction by Partial Matching)



https://blog.csdn.net/qq_42761840

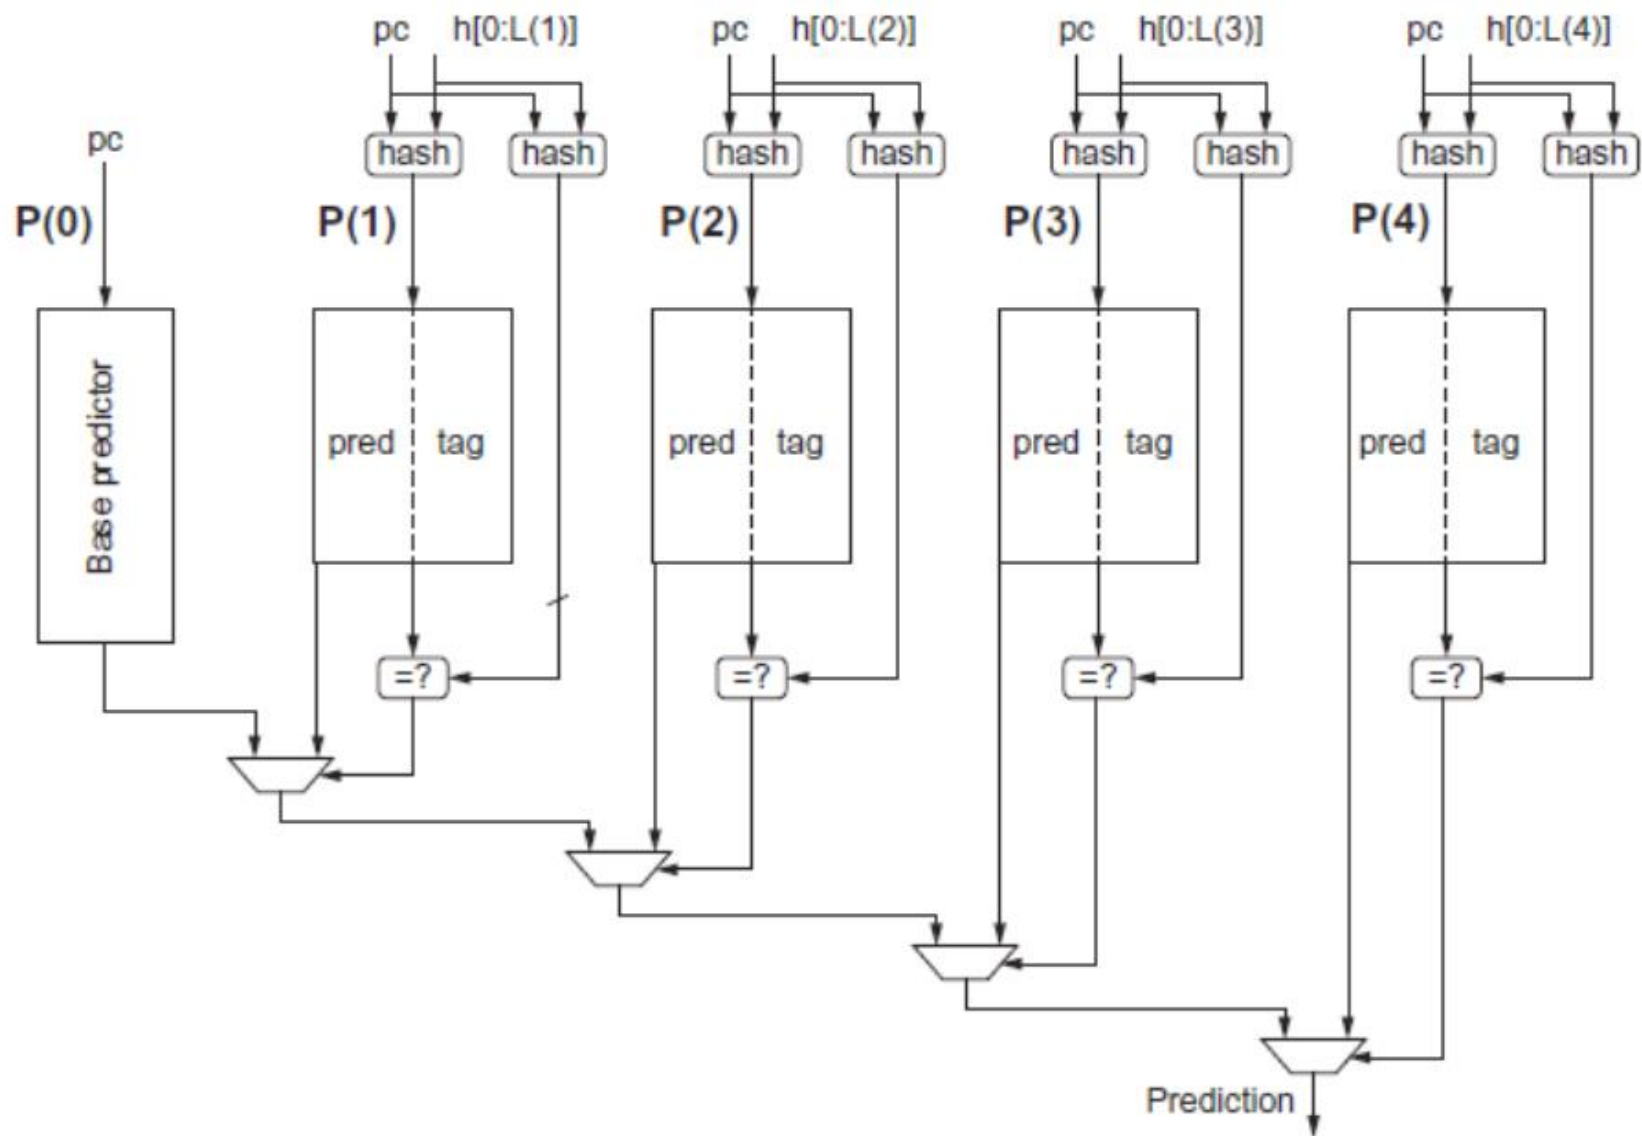
Michaud Pierre 2005

- *A PPM-like, Tag-based Predictor.*
- *J. Instr. Level Parallelism 7* (2005)
- 第一届CBP取得了第五名的成绩


Tagged Hybrid Predictors

- ❑ Need to have predictor for each branch and history
 - Problem: this implies huge tables
 - Solution:
 - Use hash tables, whose hash value is based on branch address and branch history
 - Longer histories may lead to increased chance of hash collision, so use multiple tables with increasingly shorter histories

TAGE (Tagged GEometric history length branch prediction):



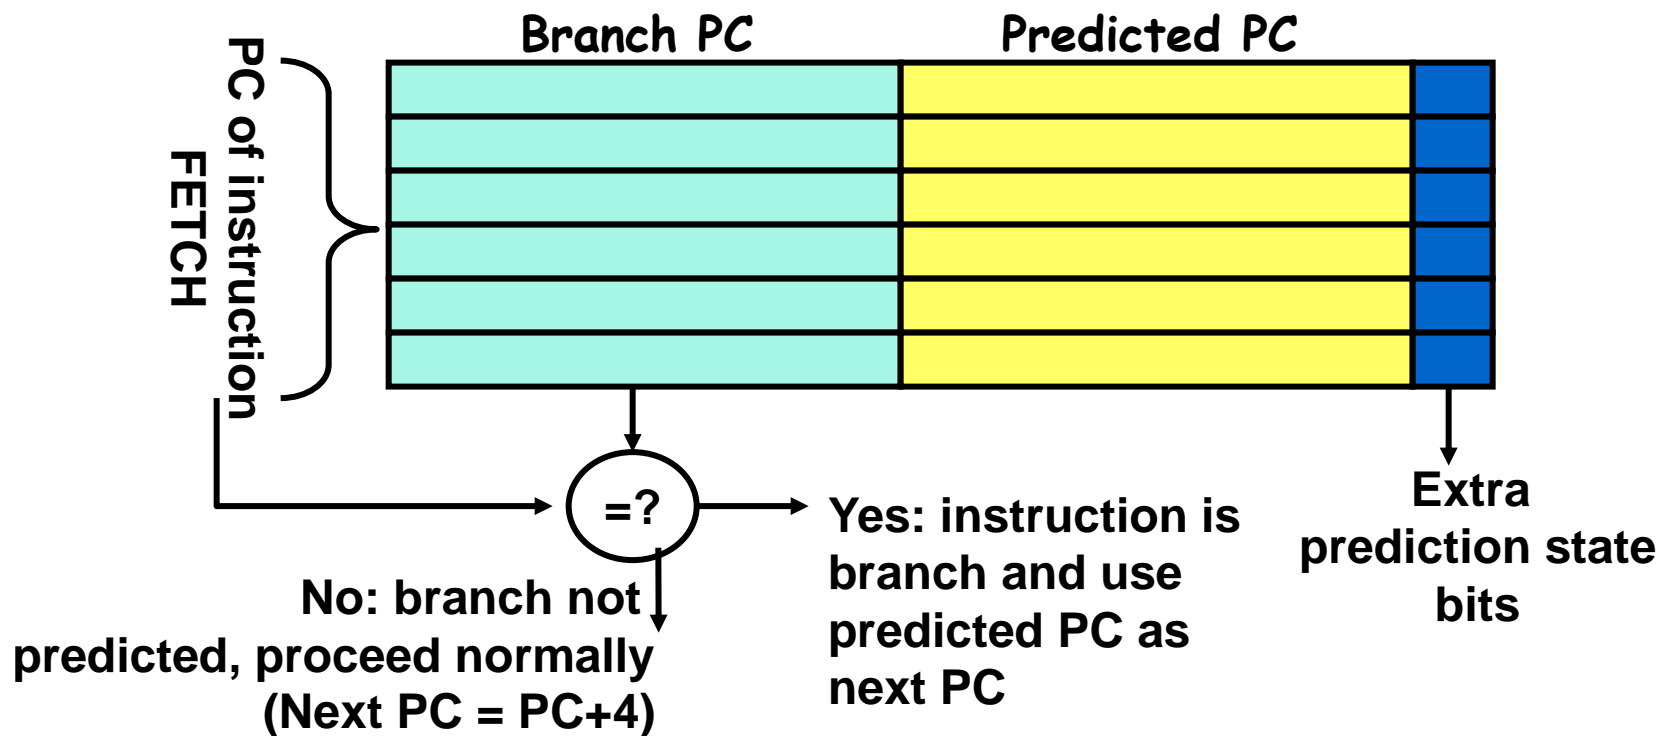
Tagged Hybrid Predictors



5. Branch Target Buffer

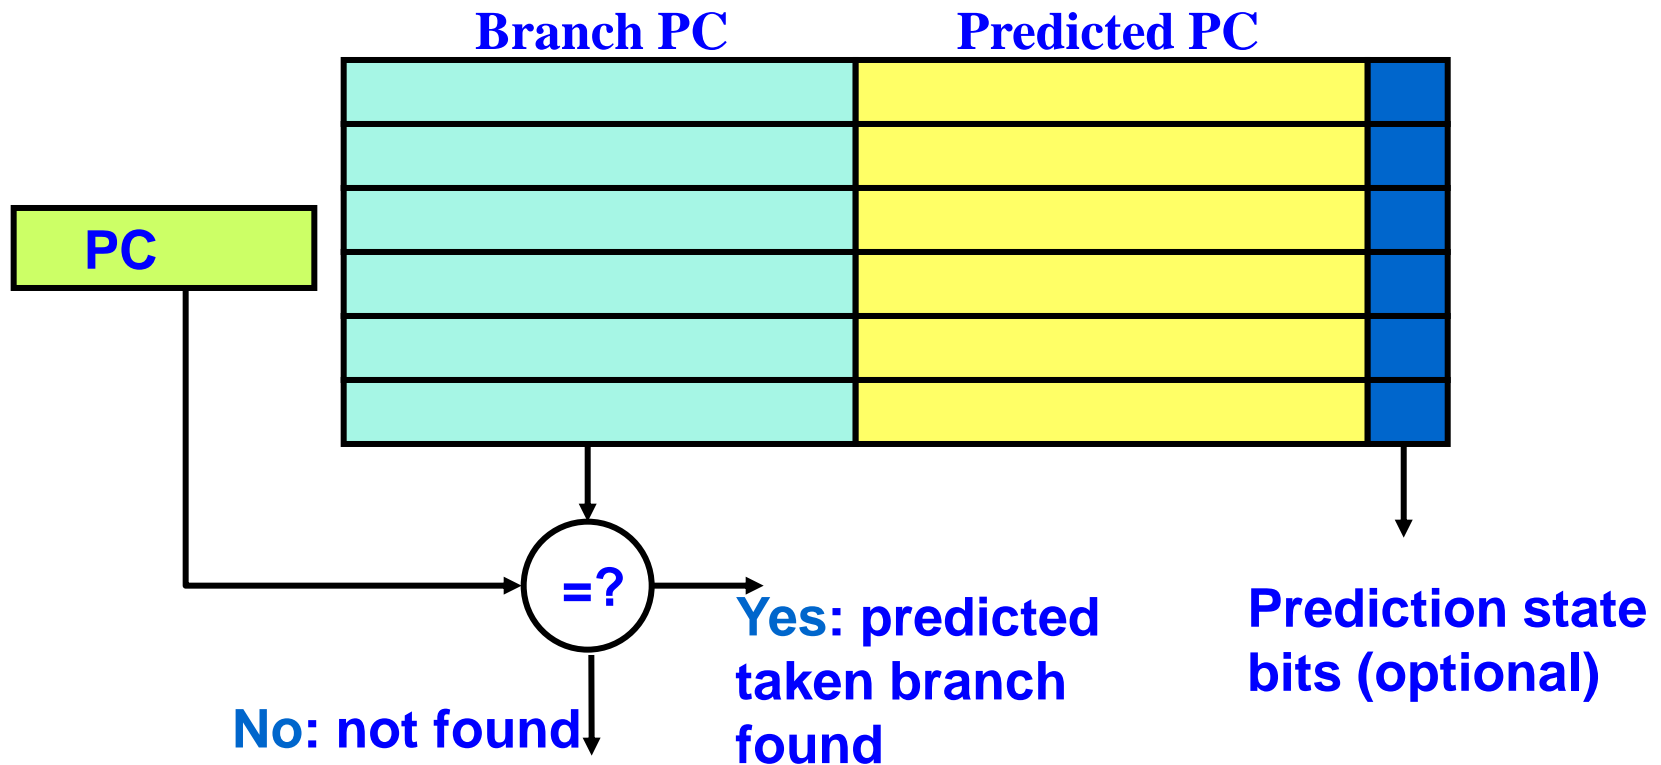
❑ Branch Target Buffer (Branch Target Cache):

- Address of branch index to get **prediction AND branch address** (if taken)
- Note: must check for branch match now, since can't use wrong branch address

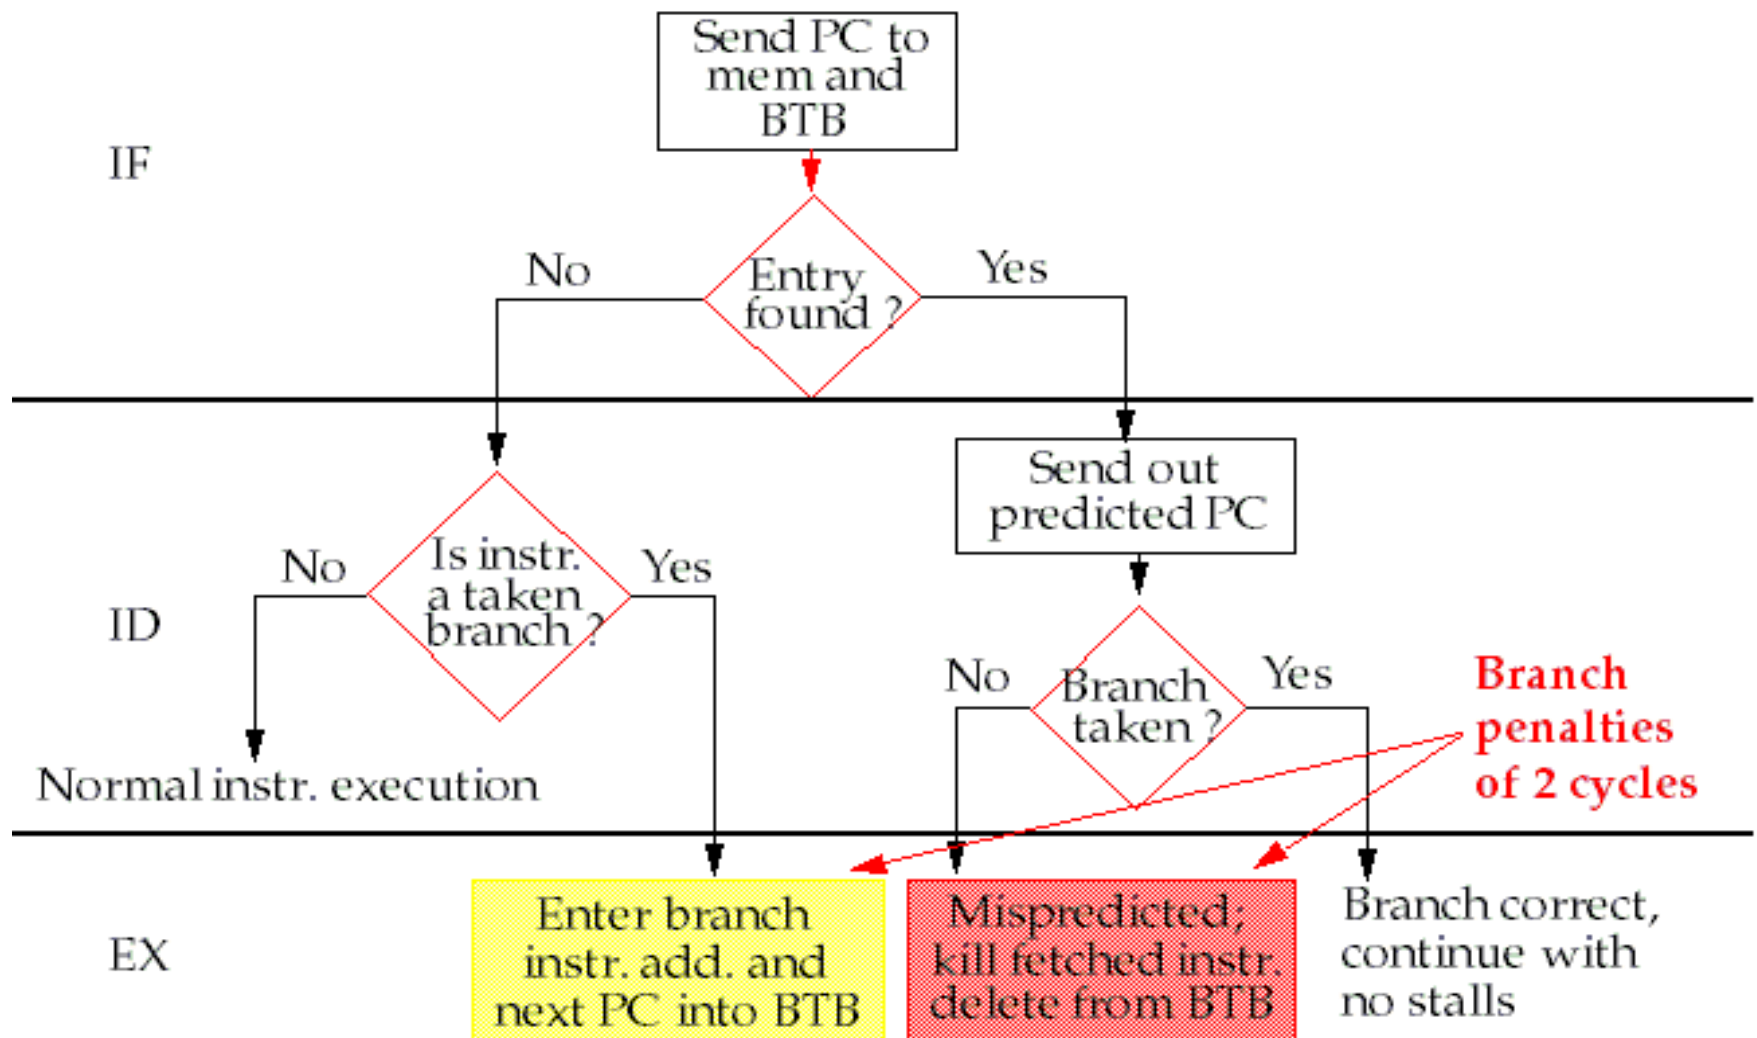


Branch Target “Cache”

- ❑ Branch Target cache - Only predicted taken branches
- ❑ “Cache” - Content Addressable Memory (CAM) or Associative Memory (see figure)
- ❑ Use a big Branch History Table & a small Branch Target Cache



Steps in handling an instruction with a Branch-Target Buffer



Variations of Branch Target Buffer

- ❑ Instead of storing just the branch address, the BTB can store the actual instruction as well.
- ❑ Storing both the branch address and the actual instruction in the buffer
- ❑ Very big branch Target buffer to cache addresses or instructions from multiple paths, ex, the predicted and unpredicted direction

6. Integrated Instruction Fetch Units

- ❑ Increasing the number of instructions executed per clock, instruction fetch will become an significant bottleneck.
- ❑ Integrated instruction fetch unit:
 - Integrated branch prediction
 - Instruction prefetch
(see P438 hardware prefetch & P483 ex. Instruction prefetcher)
 - Instruction memory access and buffering

7、Return Address Predictors

- ❑ Technique for predicting indirect jumps, whose destination address varies at run time.
- ❑ Register Indirect branch hard to predict address
- ❑ SPEC89: **procedure return** account for 85% of the indirect jumps.
- ❑ Branch-target buffer doesn't work well for procedure return

Return Addresses Buffer

- ❑ Since stack discipline for procedures, **save return address in small buffer that acts like a stack:**
 - Caches the most recent return addresses.
 - Push return address into stack at a call;
 - pop an address off at a return

- ❑ SPEC benchmark shows a stack with 8-16 entries works quite well. (fig-3.22)

Dynamic Branch Prediction Summary

- ❑ Prediction becoming important part of scalar execution
- ❑ Branch History Table: 2 bits for loop accuracy
- ❑ **Correlation**: Recently executed branches correlated with next branch.
 - Either different branches
 - Or different executions of same branches
- ❑ **Tournament Predictor**: more resources to competitive solutions and pick between them
- ❑ **Tagged Hybrid Predictors**: have predictor for each branch and history
- ❑ **Branch Target Buffer**: include branch address & prediction
- ❑ Predicated Execution can reduce number of branches, number of mispredicted branches
- ❑ **Return address stack** for prediction of indirect jump

Chapter3 ILP

Hardware Based Speculation

Tomasulo Drawbacks

❑ Complexity

- delays of 360/91, MIPS 10000, Alpha 21264, IBM PPC 620 in CA:AQA 2/e, but not in silicon!

❑ Many associative stores (CDB) at high speed, Performance limited by Common Data Bus

❑ Non-precise interrupts!

- We will address this later

❑ Partially overlapped basic blocks

- integer units get ahead, beyond FP operation
- the successor basic block can not start execution until the branch is resolved, in spite of that it can be issued.

Hardware-based Speculation

❑ Speculating on the outcome of branches and executing the program as if the predictions were correct.

- Need to handle the situation where the speculation is incorrect.

❑ Key ideas:

- Dynamic branch prediction to choose which instruction to execute
- Speculation to allow the execution of instructions before the control dependences are resolved.
- Dynamic scheduling to deal with the scheduling of different combinations of basic blocks.

❑ Data flow execution:

- Operations execute as soon as their operands are available.

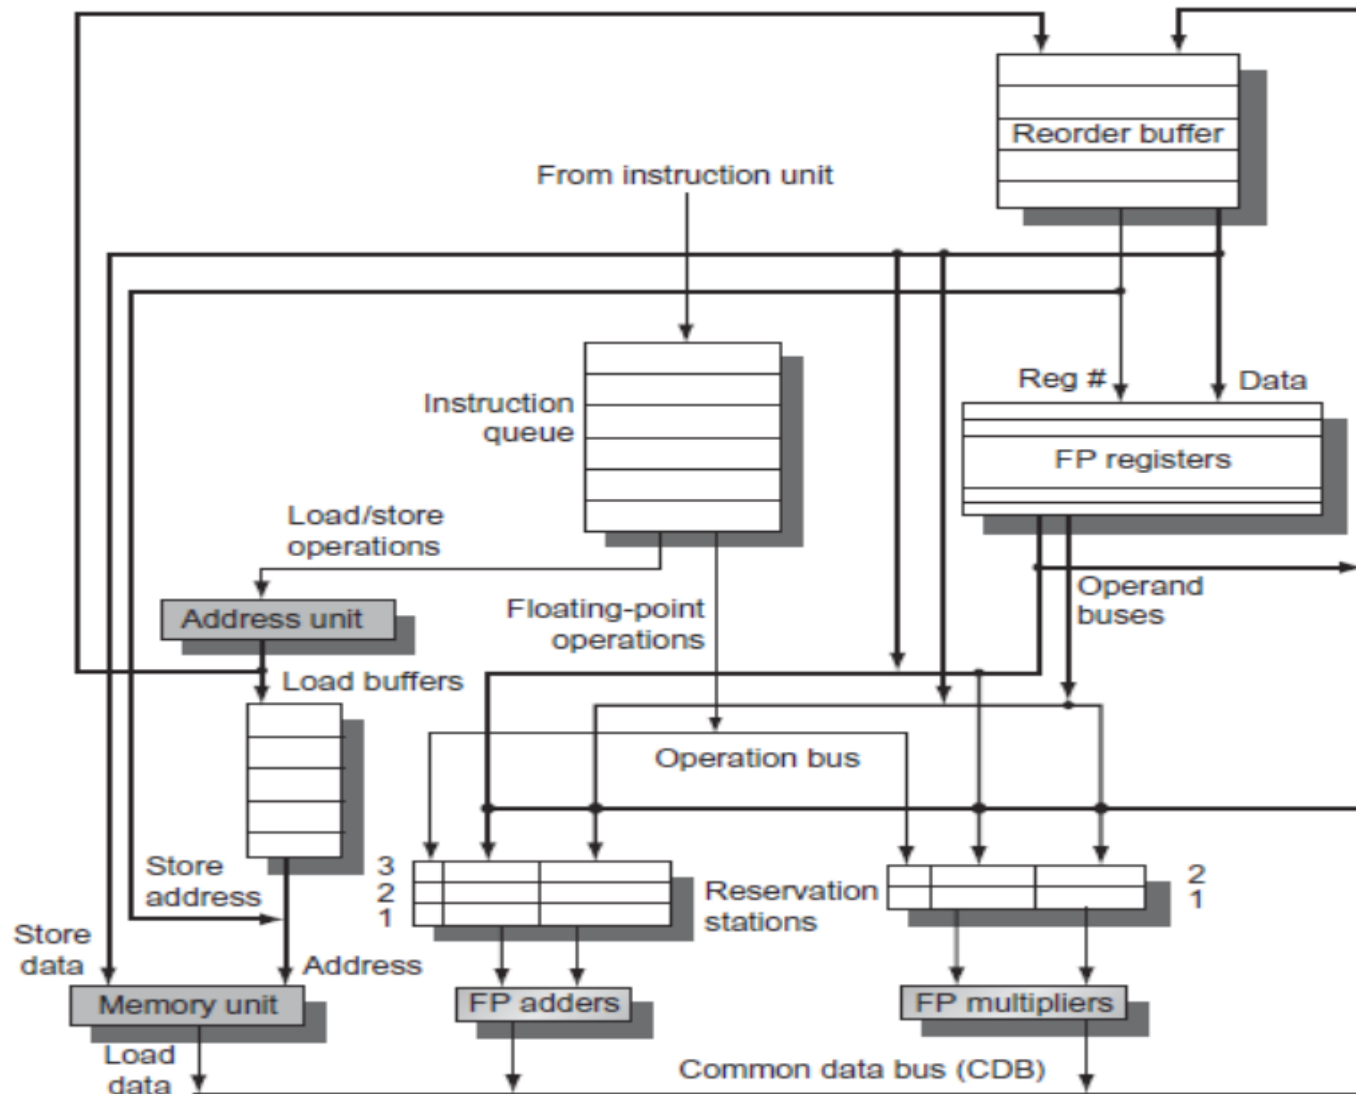
Hardware-Based Speculation

- ❑ Execute instructions along predicted execution paths but only commit the results if prediction was correct
- ❑ Instruction commit: allowing an instruction to update the register file when instruction is no longer speculative
- ❑ Need an additional piece of hardware to prevent any irrevocable action until an instruction commits
 - I.e. updating state or taking an executionb

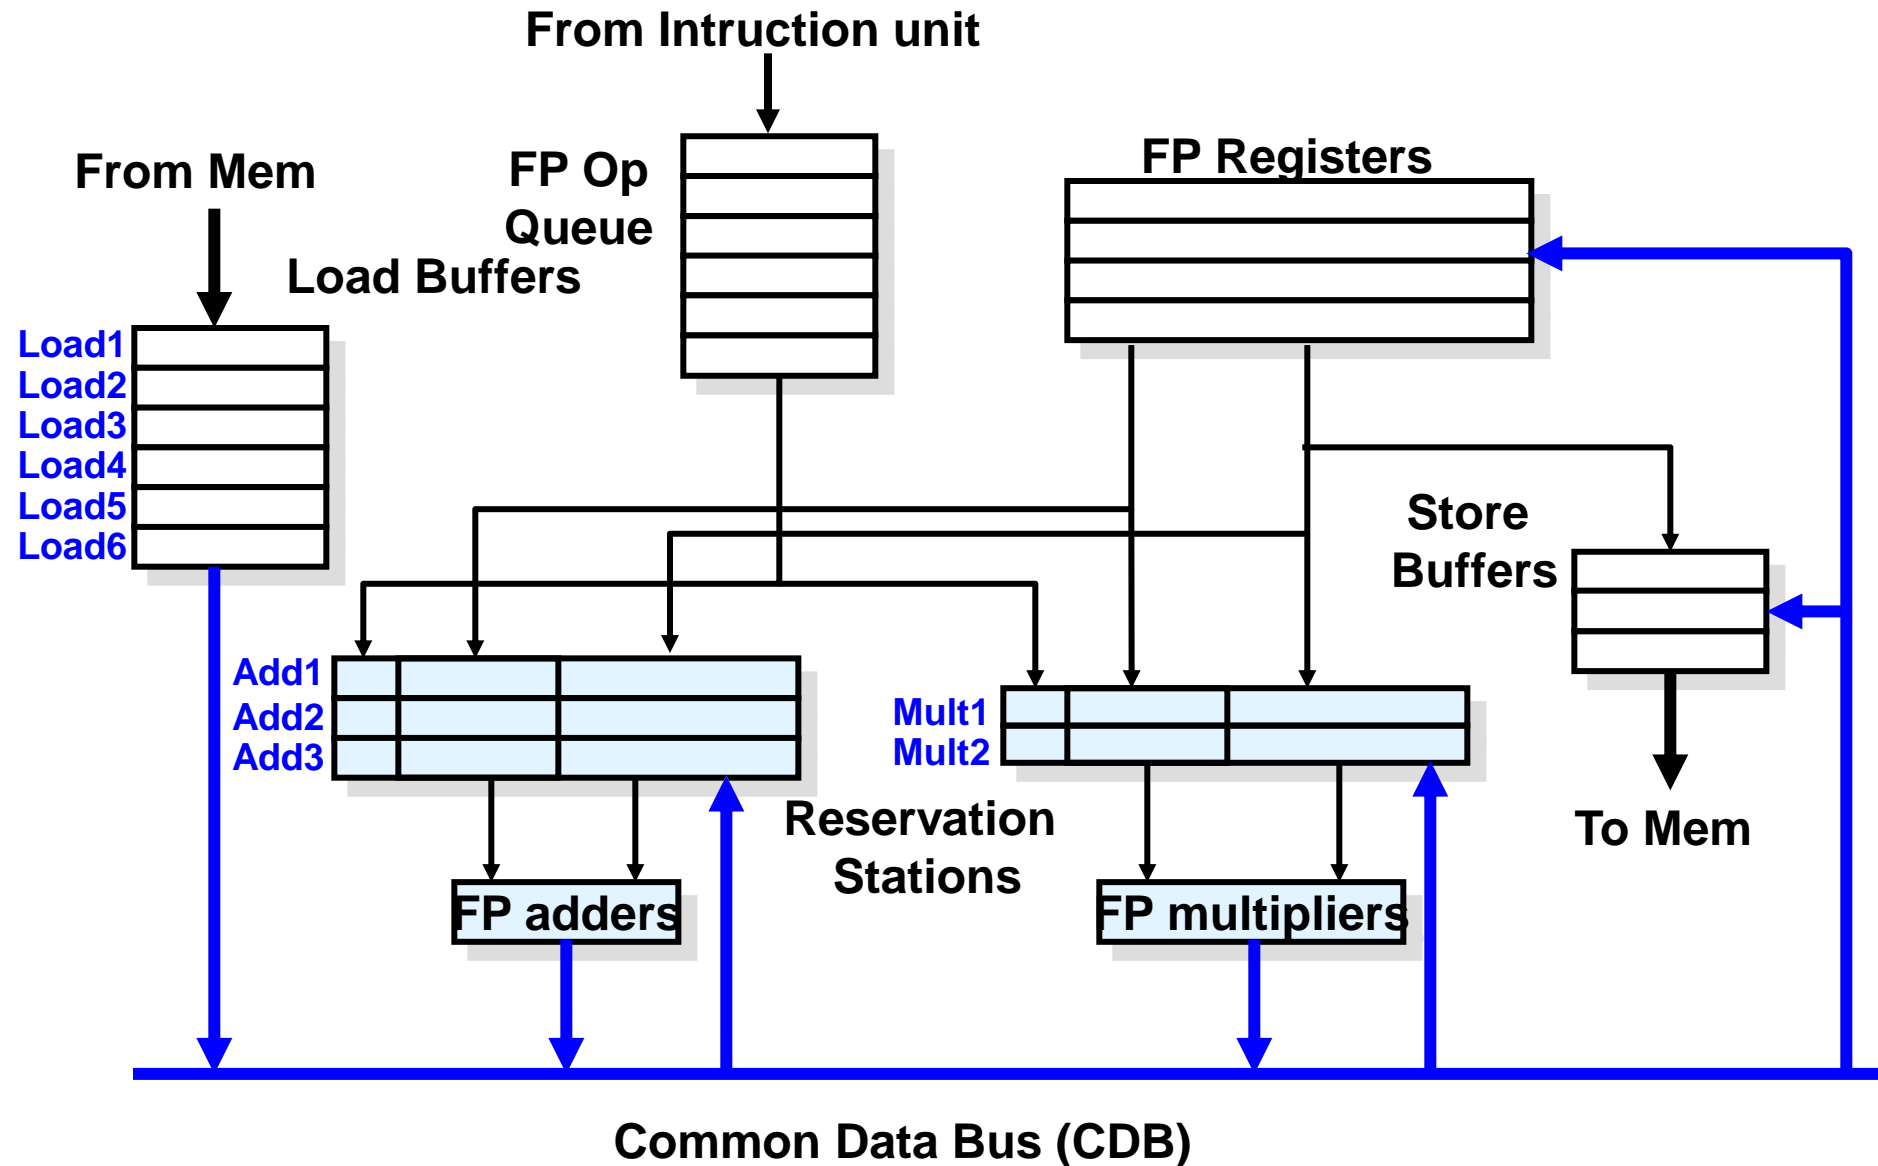
Speculative execution based on Tomasulo's algorithm

- ❑ Separate the process of completing execution and the bypassing of results among instructions from instruction commit (reg file or memory update).
- ❑ Allows other (speculative) instructions to execute, but no results are committed until we know the instruction is no longer speculative.
- ❑ Allow instructions to **execute out of order** but force them to **commit in order**, which helps with handling exceptions properly.

Extend Tomasulo's Algorithm to handle speculation



Review: Tomasulo Algorithm



Difference from Tomasulo Algorithm

- ❑ Reorder Buffer is added
- ❑ Eliminate store buffer (Mem write is ordered in Reorder buffer)
- ❑ Register renaming is done via reorder buffer instead of reservation station
- ❑ Reservation only for storing opcodes and operands during the instruction flying between issue and execution.
- ❑ ROB hold the instruction results and supply operands in the interval between completion of instruction execution and instruction commit.

Entry of ROB

- ☐ Instruction type
- ☐ Destination field
- ☐ Value field
- ☐ Ready field

- ☐ Exception vector

4 Steps of Speculative Tomasulo Algorithm

1. Issue—get instruction from FP Op Queue

If reservation station **and reorder buffer slot** free, issue instr & send operands **& reorder buffer no. for destination** (this stage sometimes called “**dispatch**”), update control entries “in use”.

2. Execution—operate on operands (EX)

When both operands ready then execute; if not ready, watch CDB for result; when both in reservation station, execute; **checks RAW**

3. Write result—finish execution (WB)

Write on Common Data Bus to all awaiting FUs **& reorder buffer**; mark the reservation station available.

4. Commit—**update register with reorder result**

When instruction at head of reorder buffer & result present, update register with result (or store to memory) and remove instr from reorder buffer. **Mispredicted branch flushes reorder buffer (sometimes called “graduation”)**

What to do when commit ?

when an instruction reaches the head of ROB and its result is present in the buffer, then

❑ For non-Branch instruction

- Update the register with the result
- Remove the instruction from ROB
- When it's a STORE, update the memory instead of Reg.

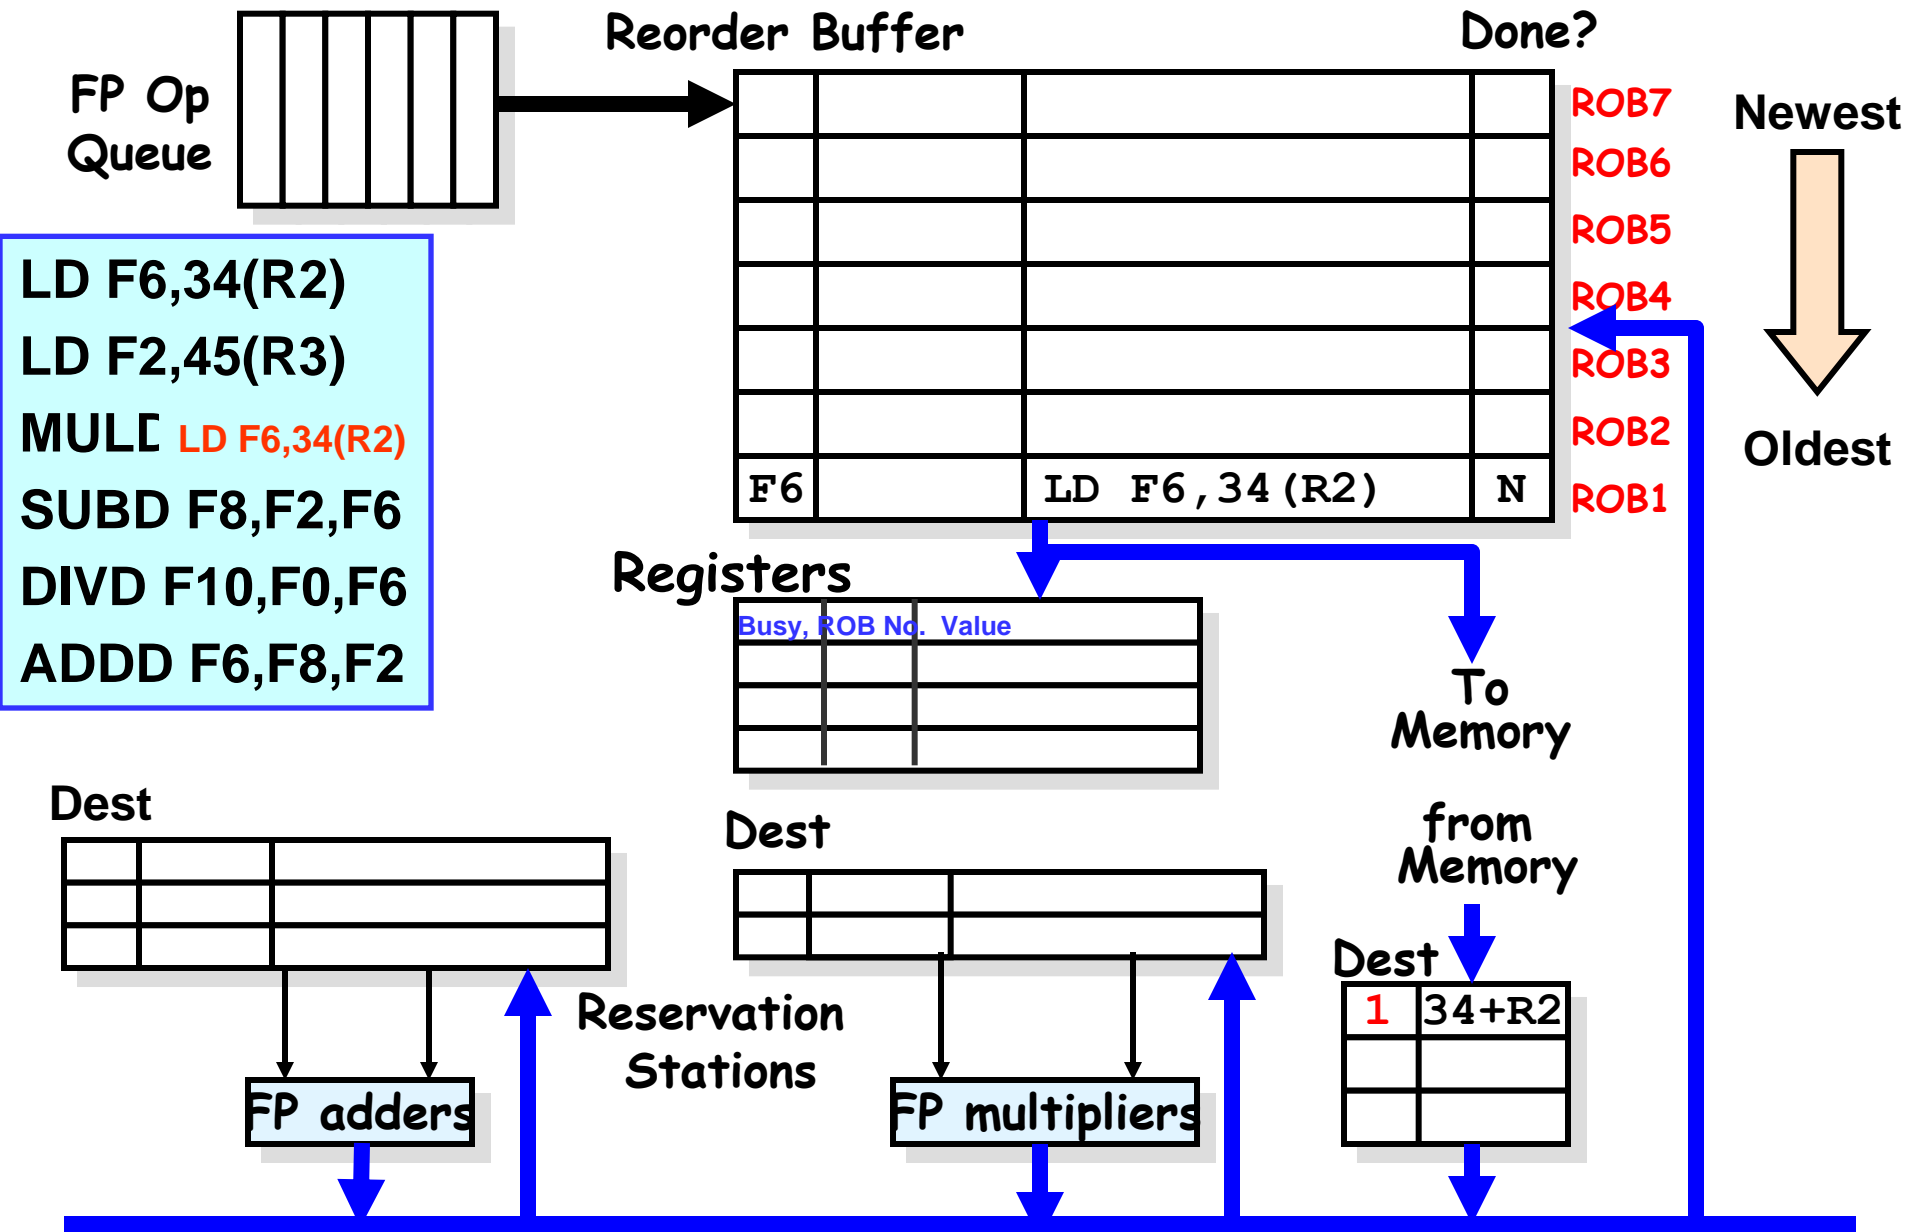
❑ For a mispredicted branch

- ROB is flushed and execution is restarted at the correct successor of the branch

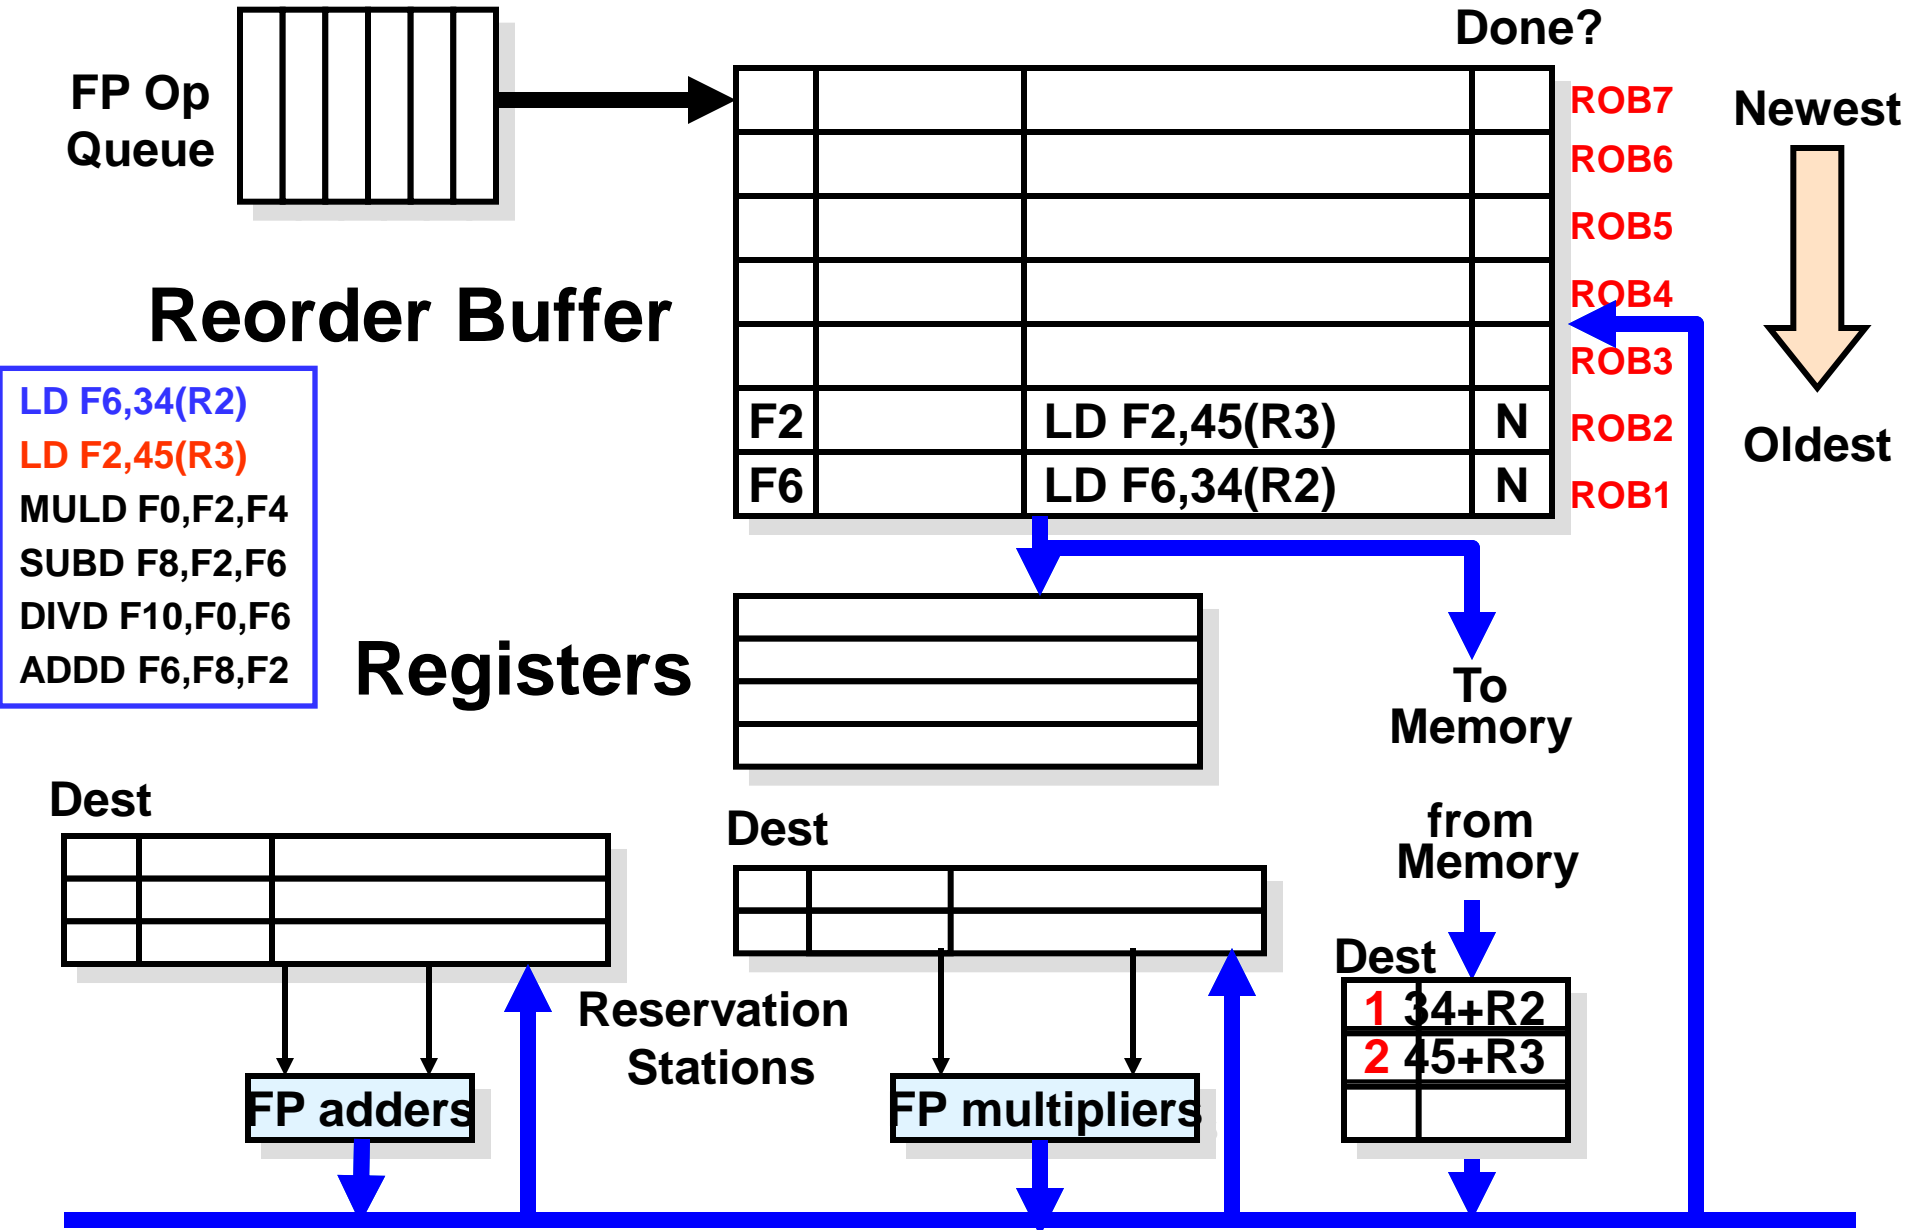
❑ For a right-predicted branch

- End of branch instruction.

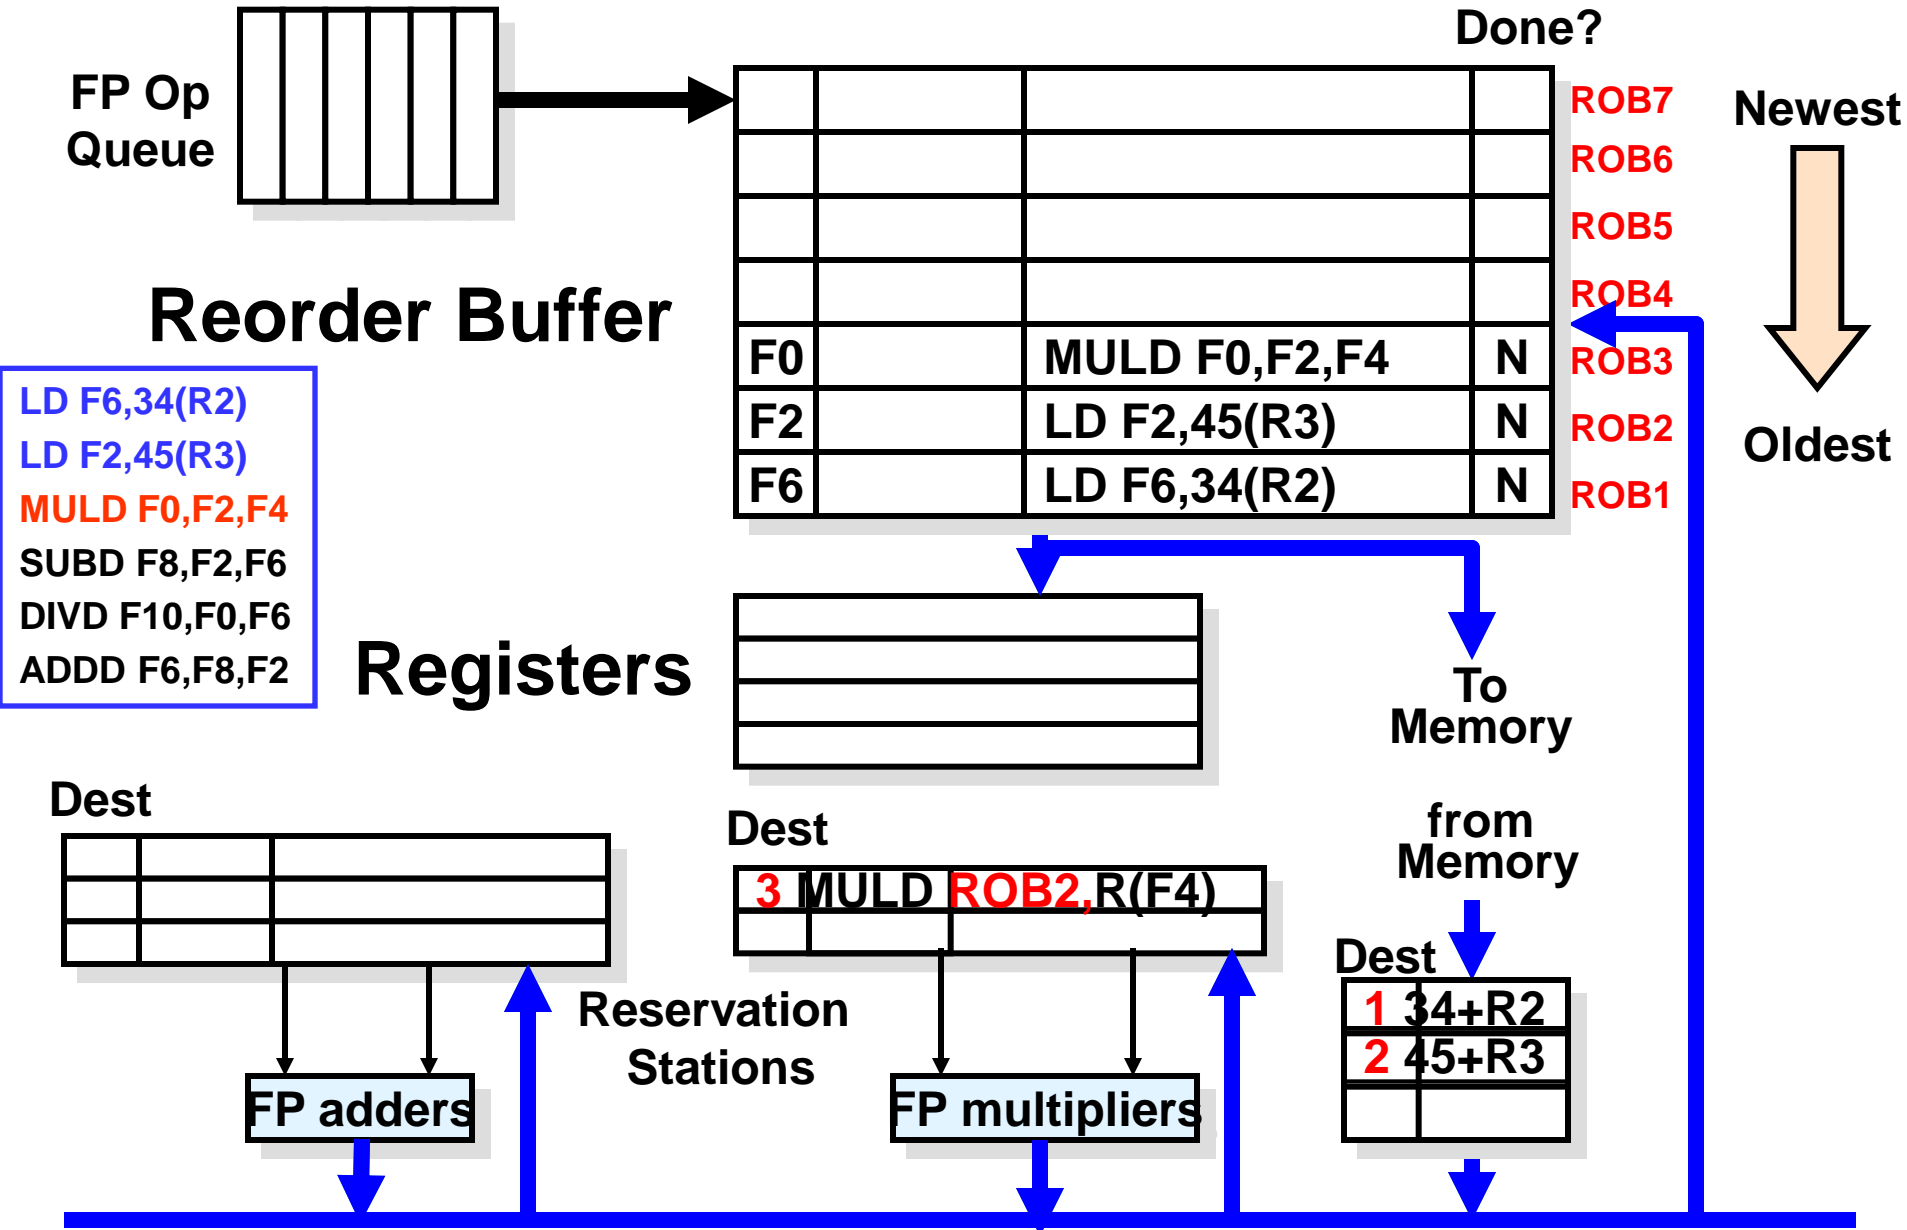
2 clock cycles for pipelined memory access



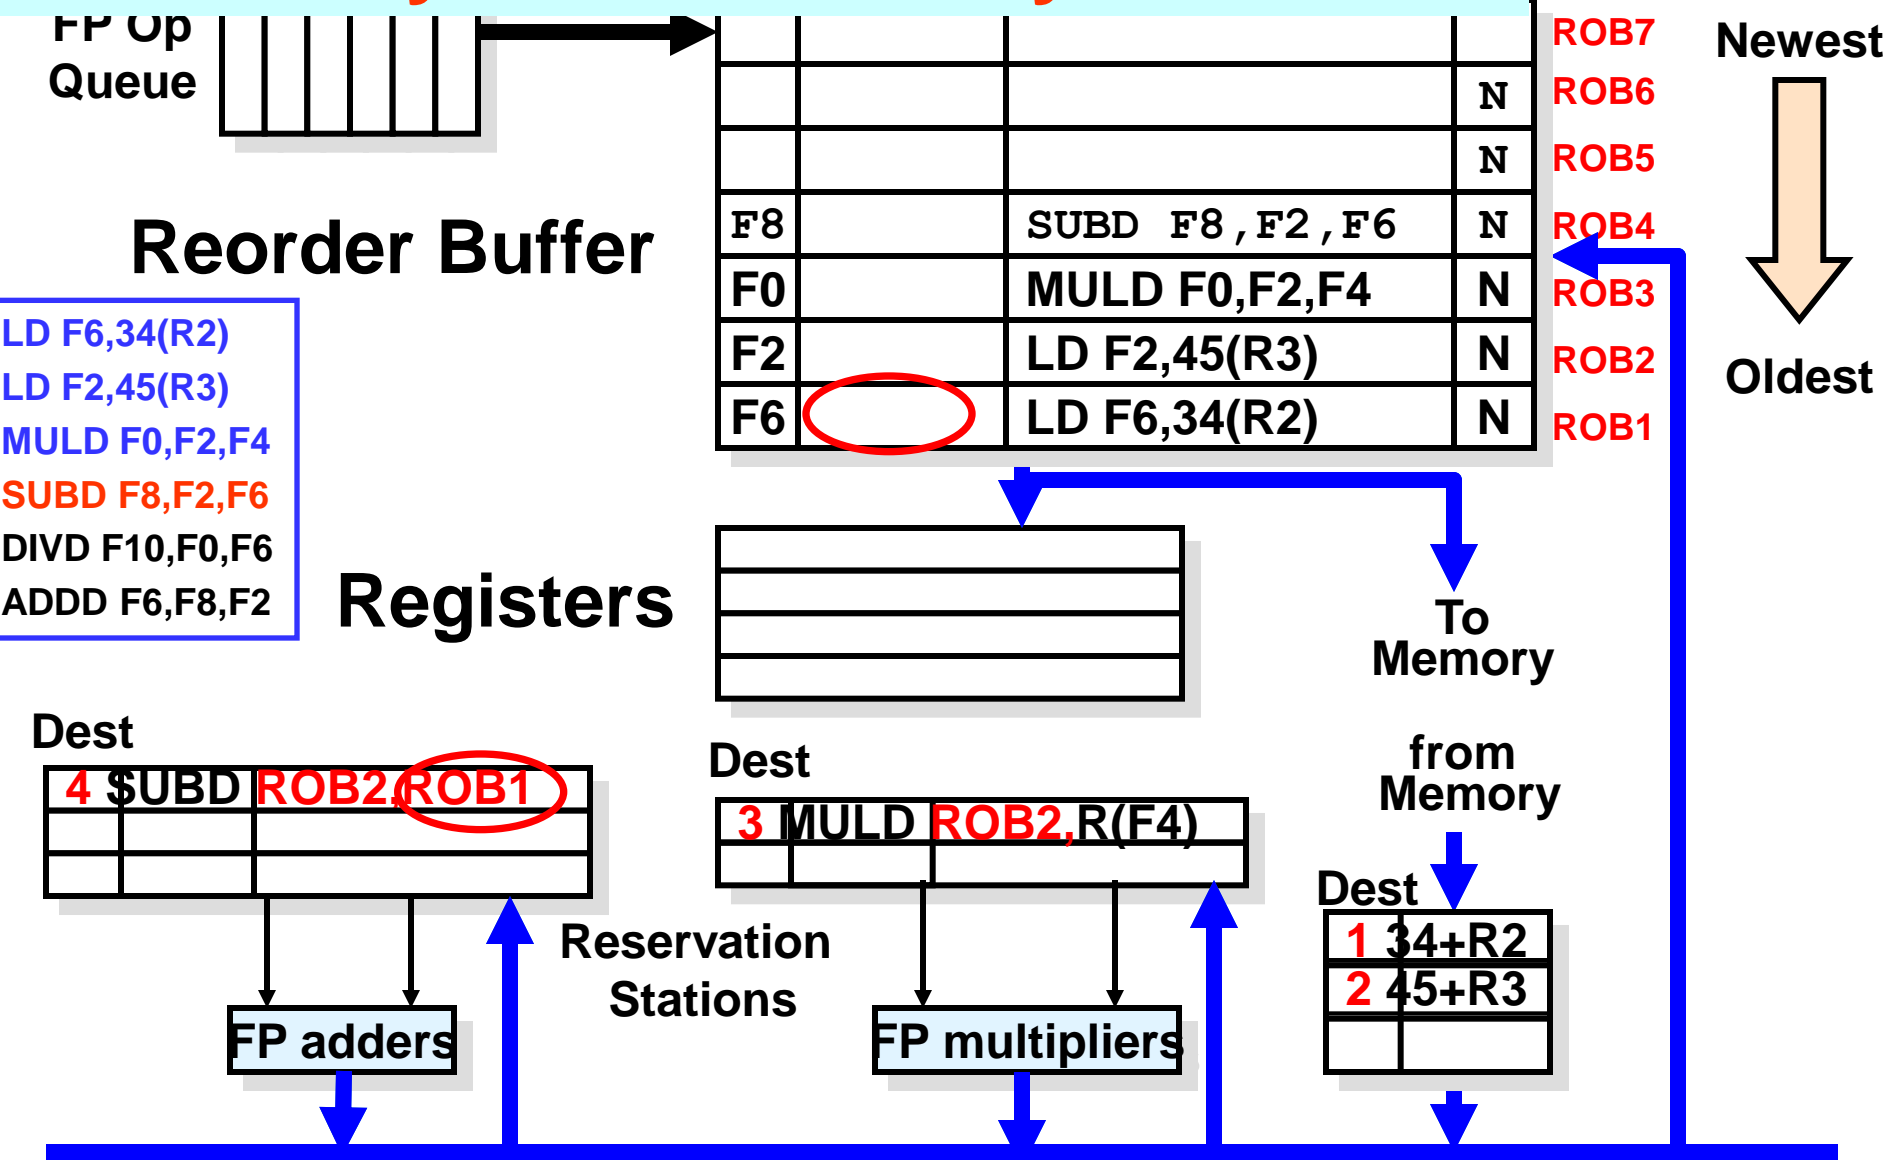
Tomasulo With Reorder buffer:



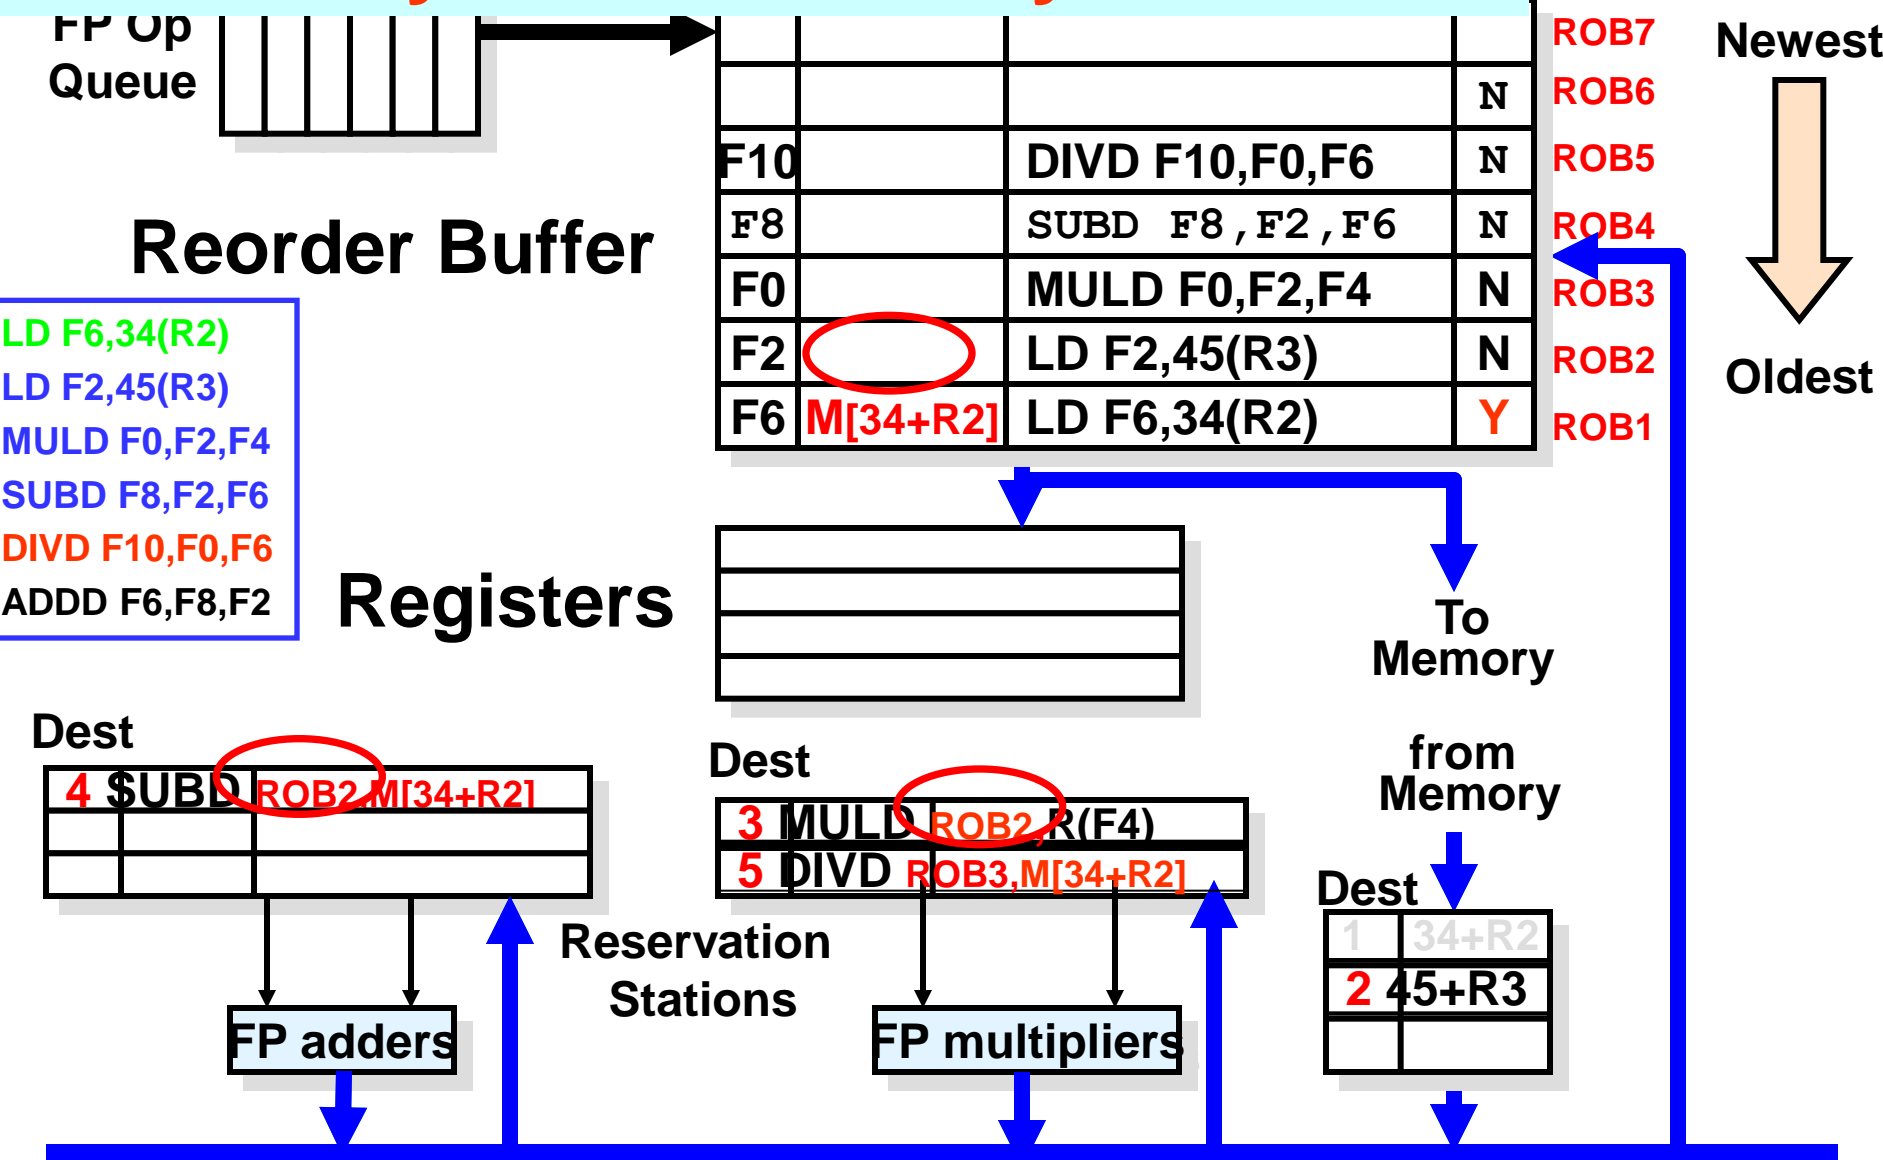
Tomasulo With Reorder buffer:



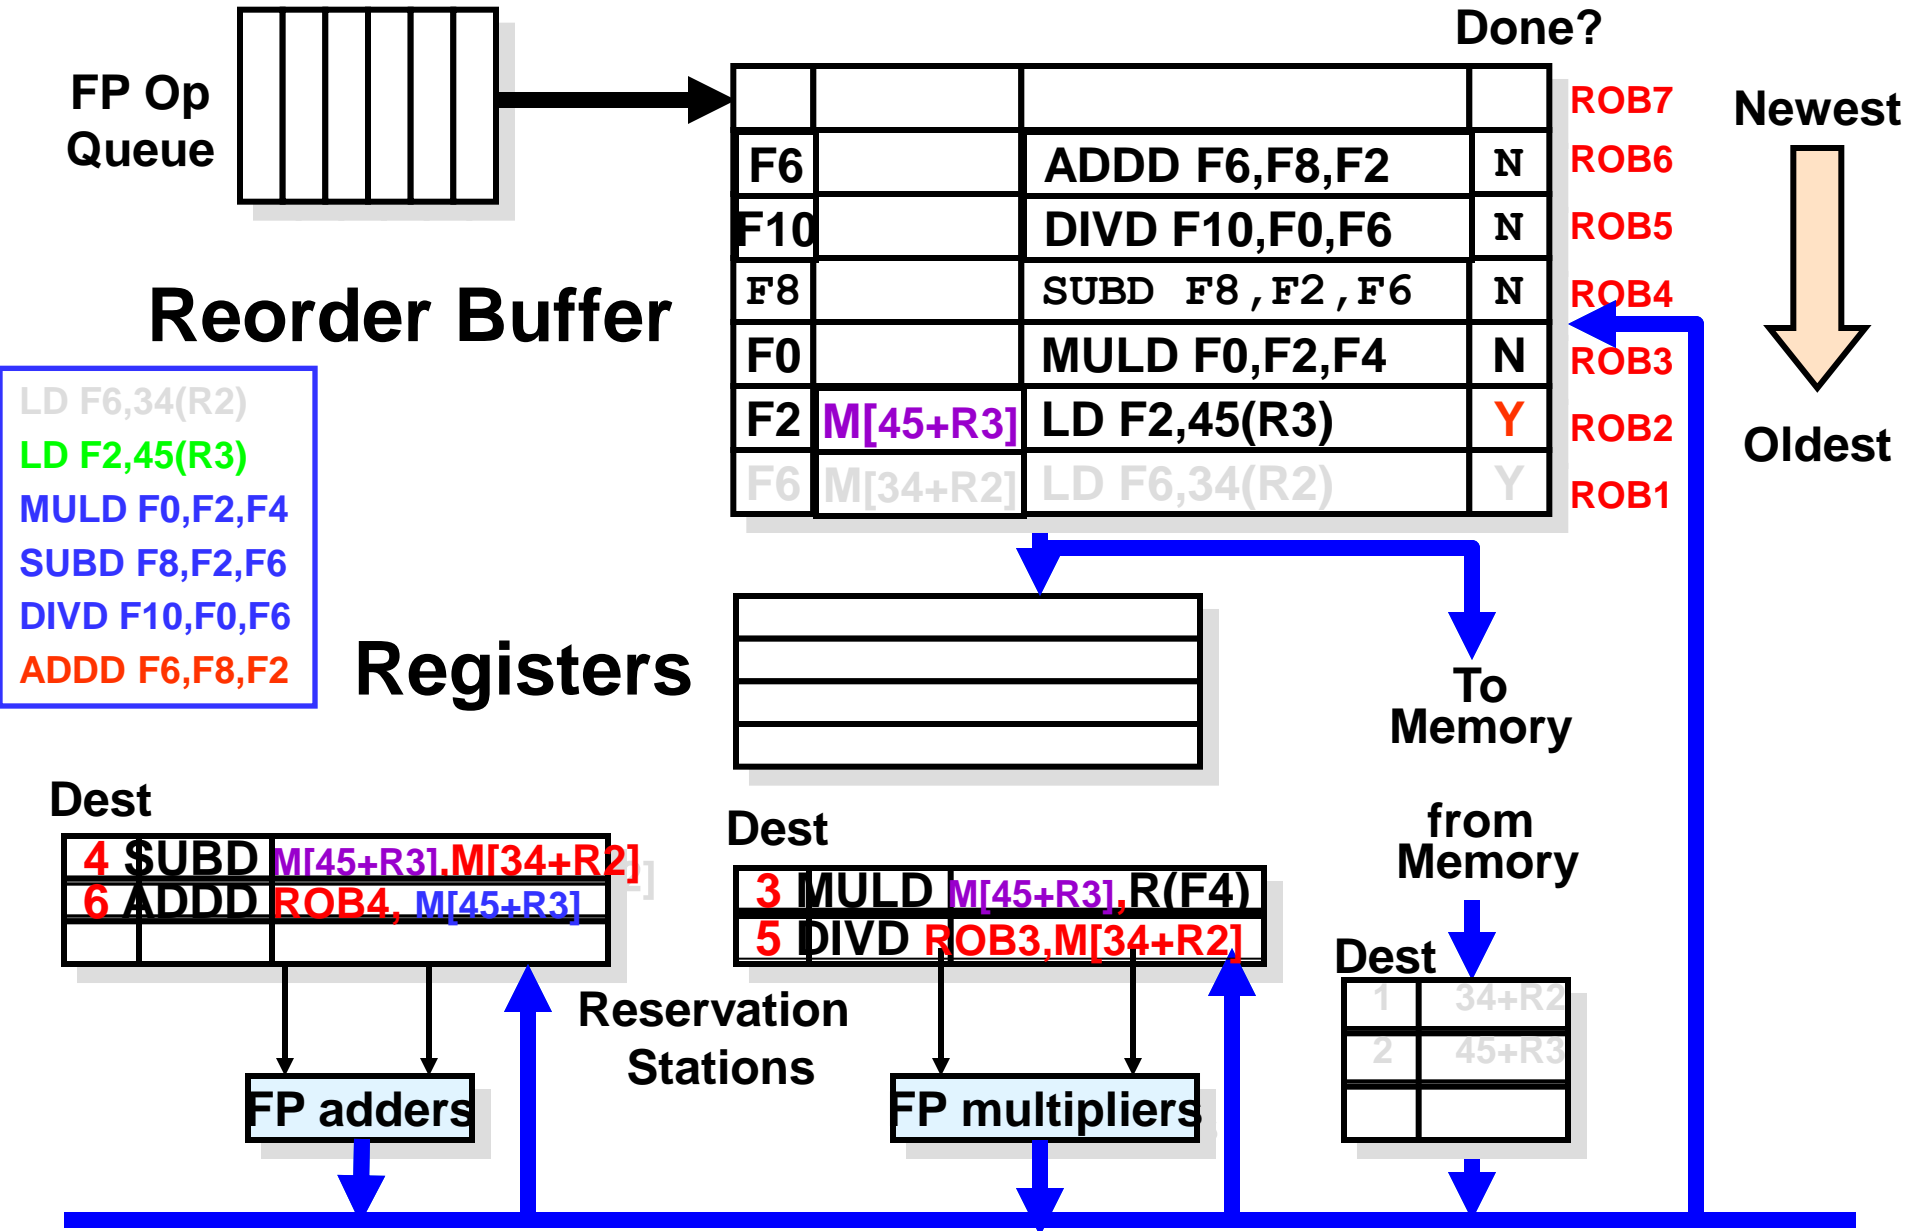
**Note: The first LD will finish
memory access next cycle**



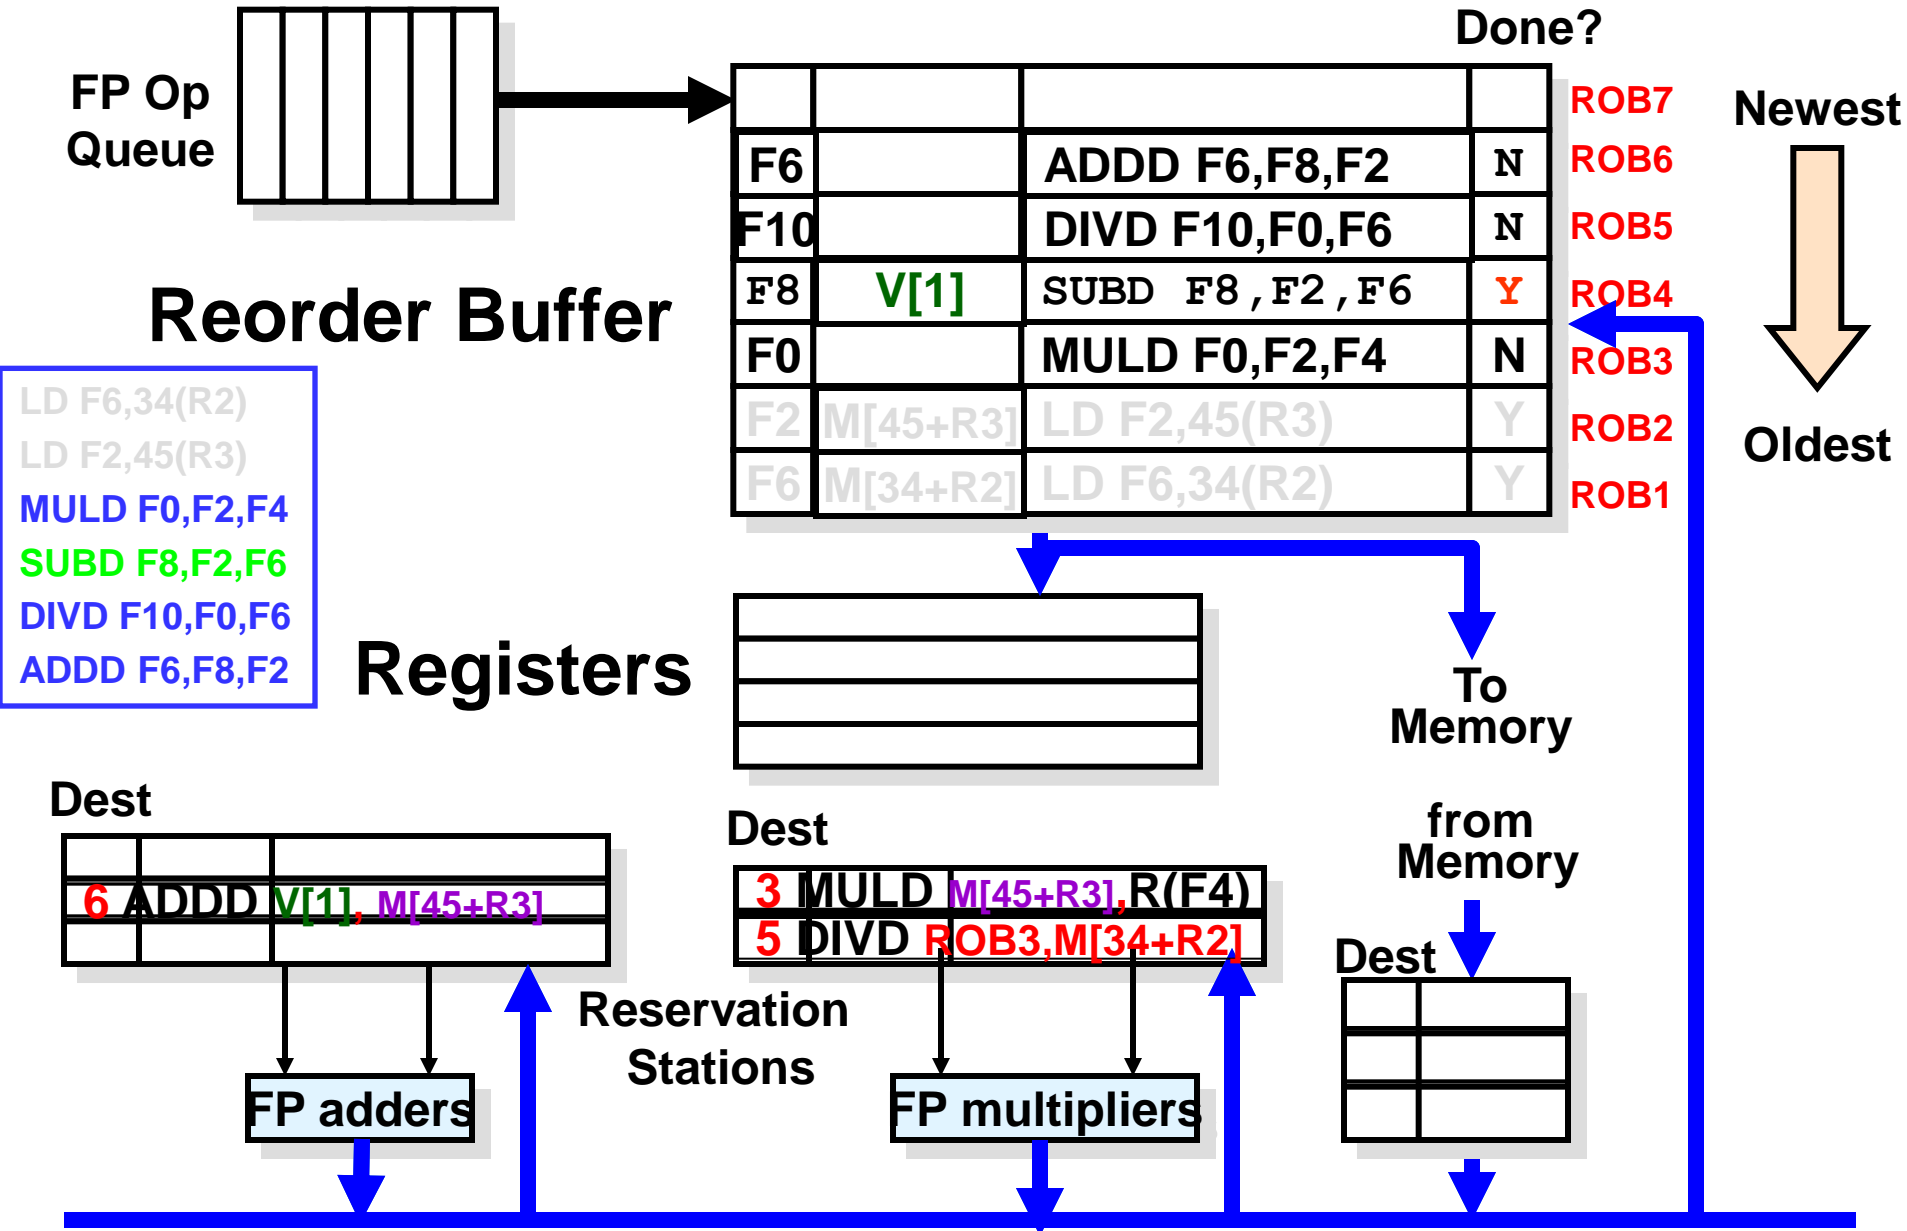
Note: The 2nd LD will finish
memory access next cycle



Tomasulo With Reorder buffer:



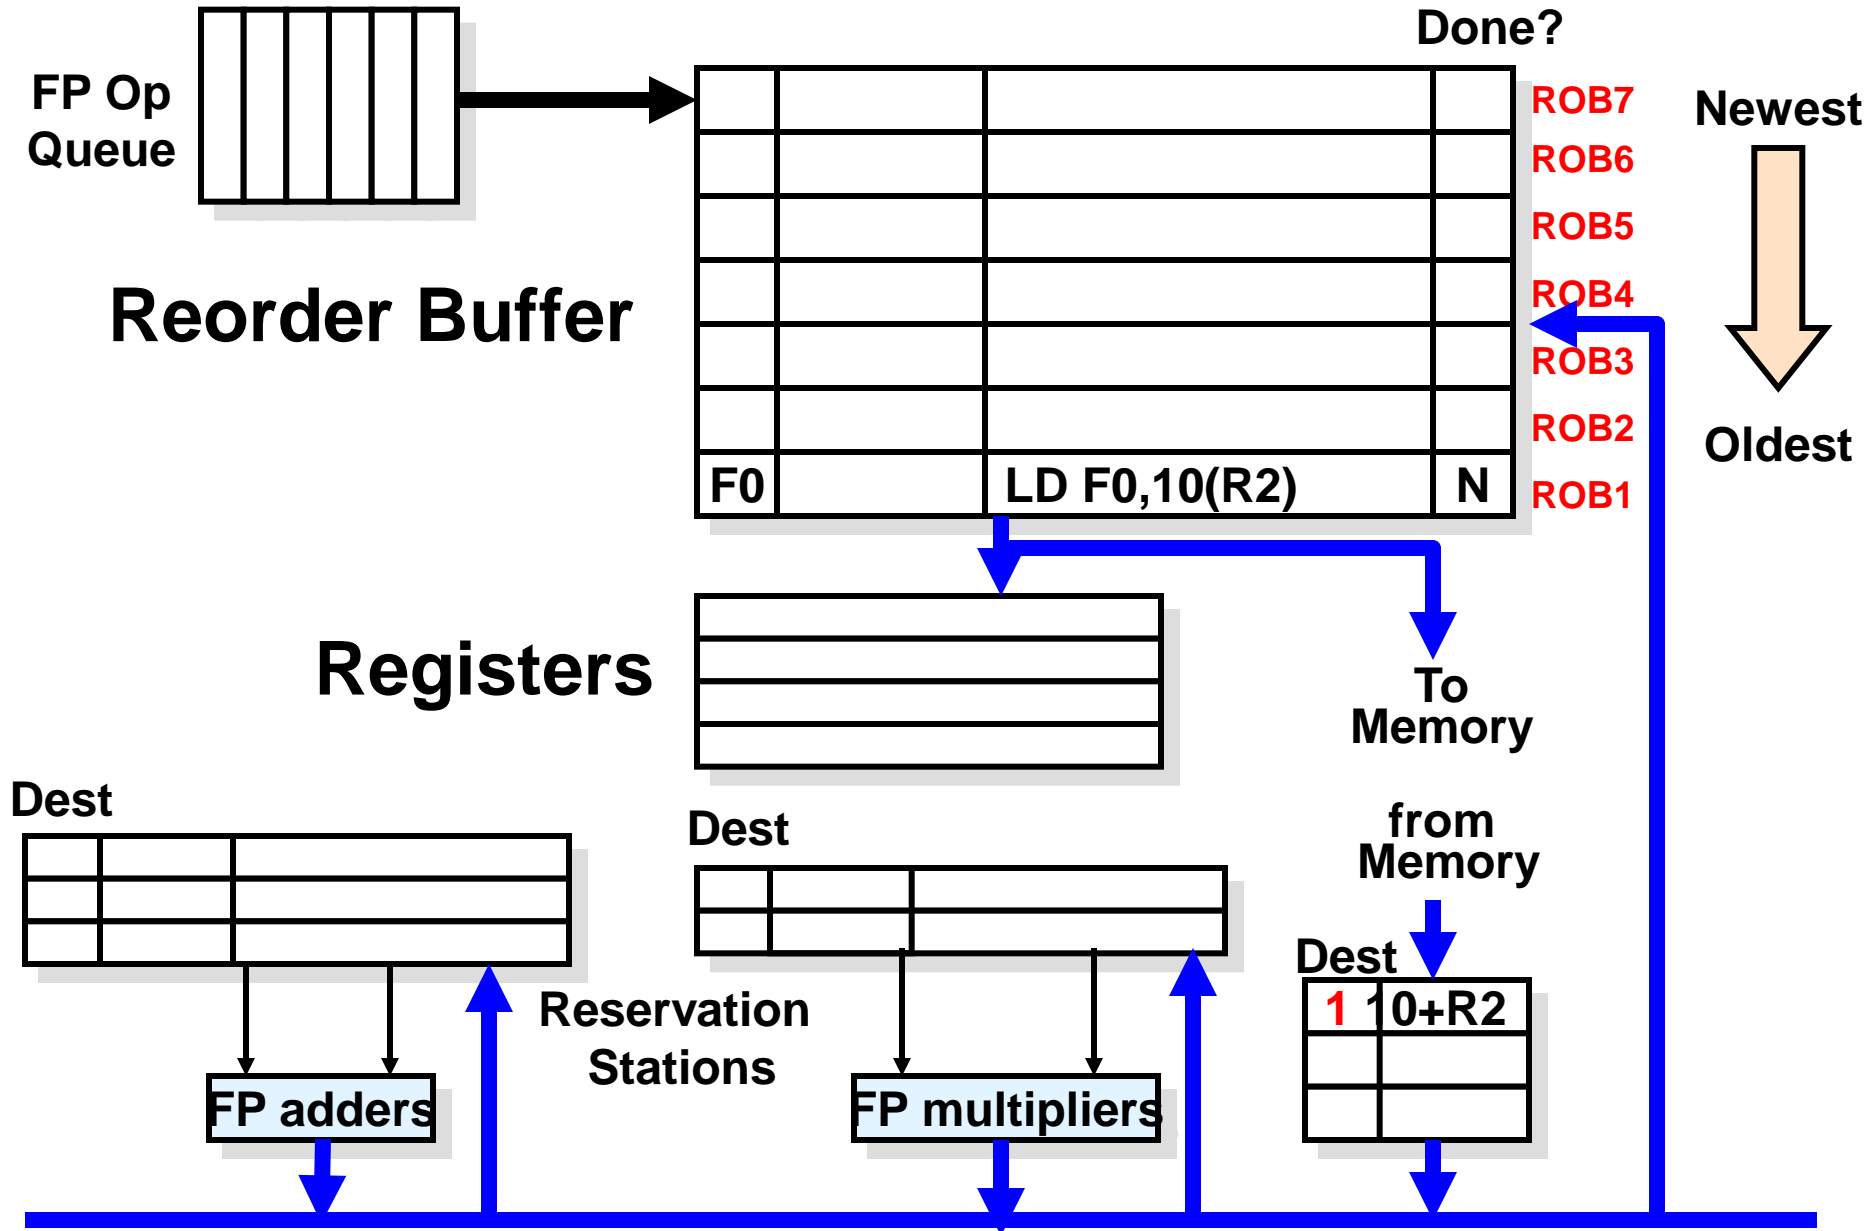
Tomasulo With Reorder buffer:



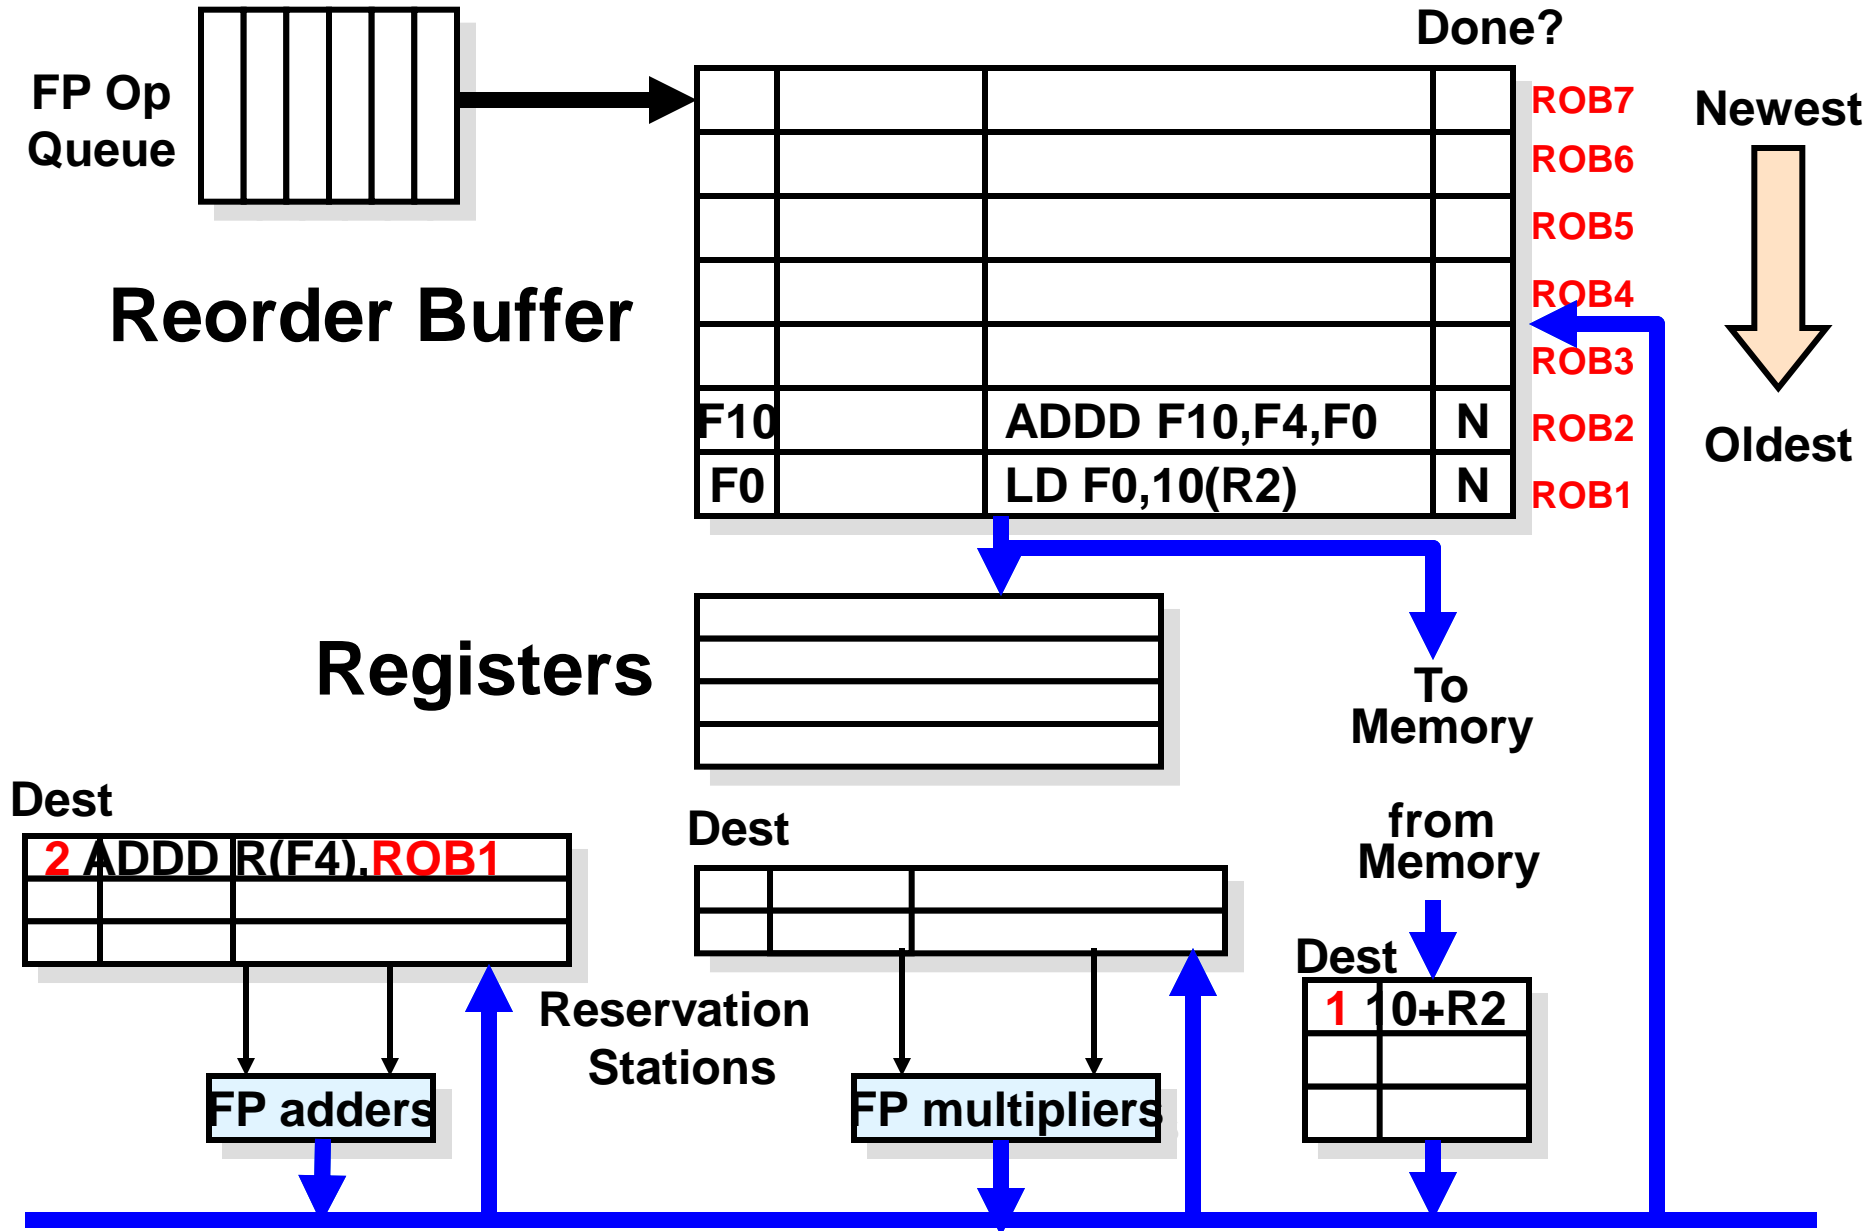
Speculation performance

- ❑ Overlapped basic block
- ❑ Is there any new problems ?
 - Pay attention to the memory accessing

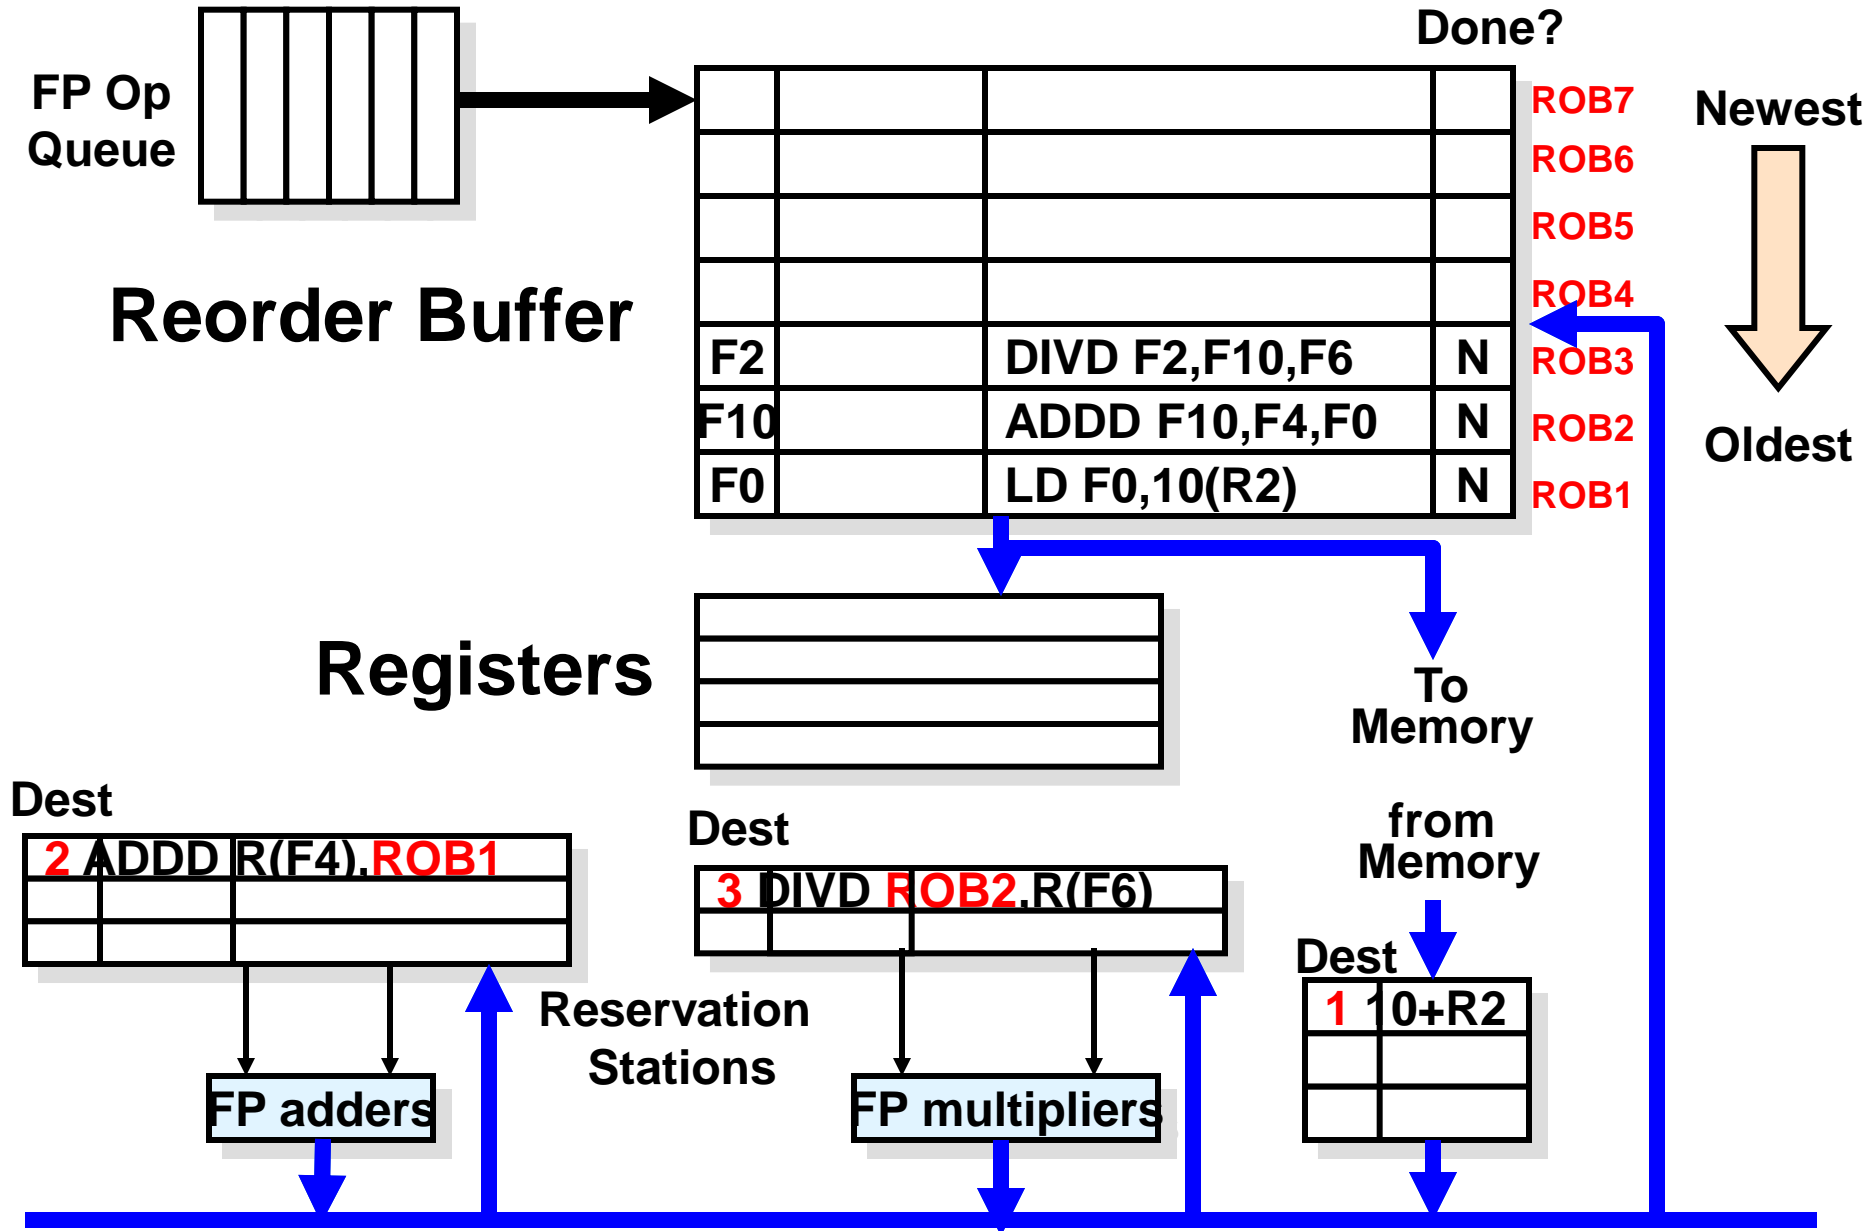
Tomasulo With Reorder buffer:



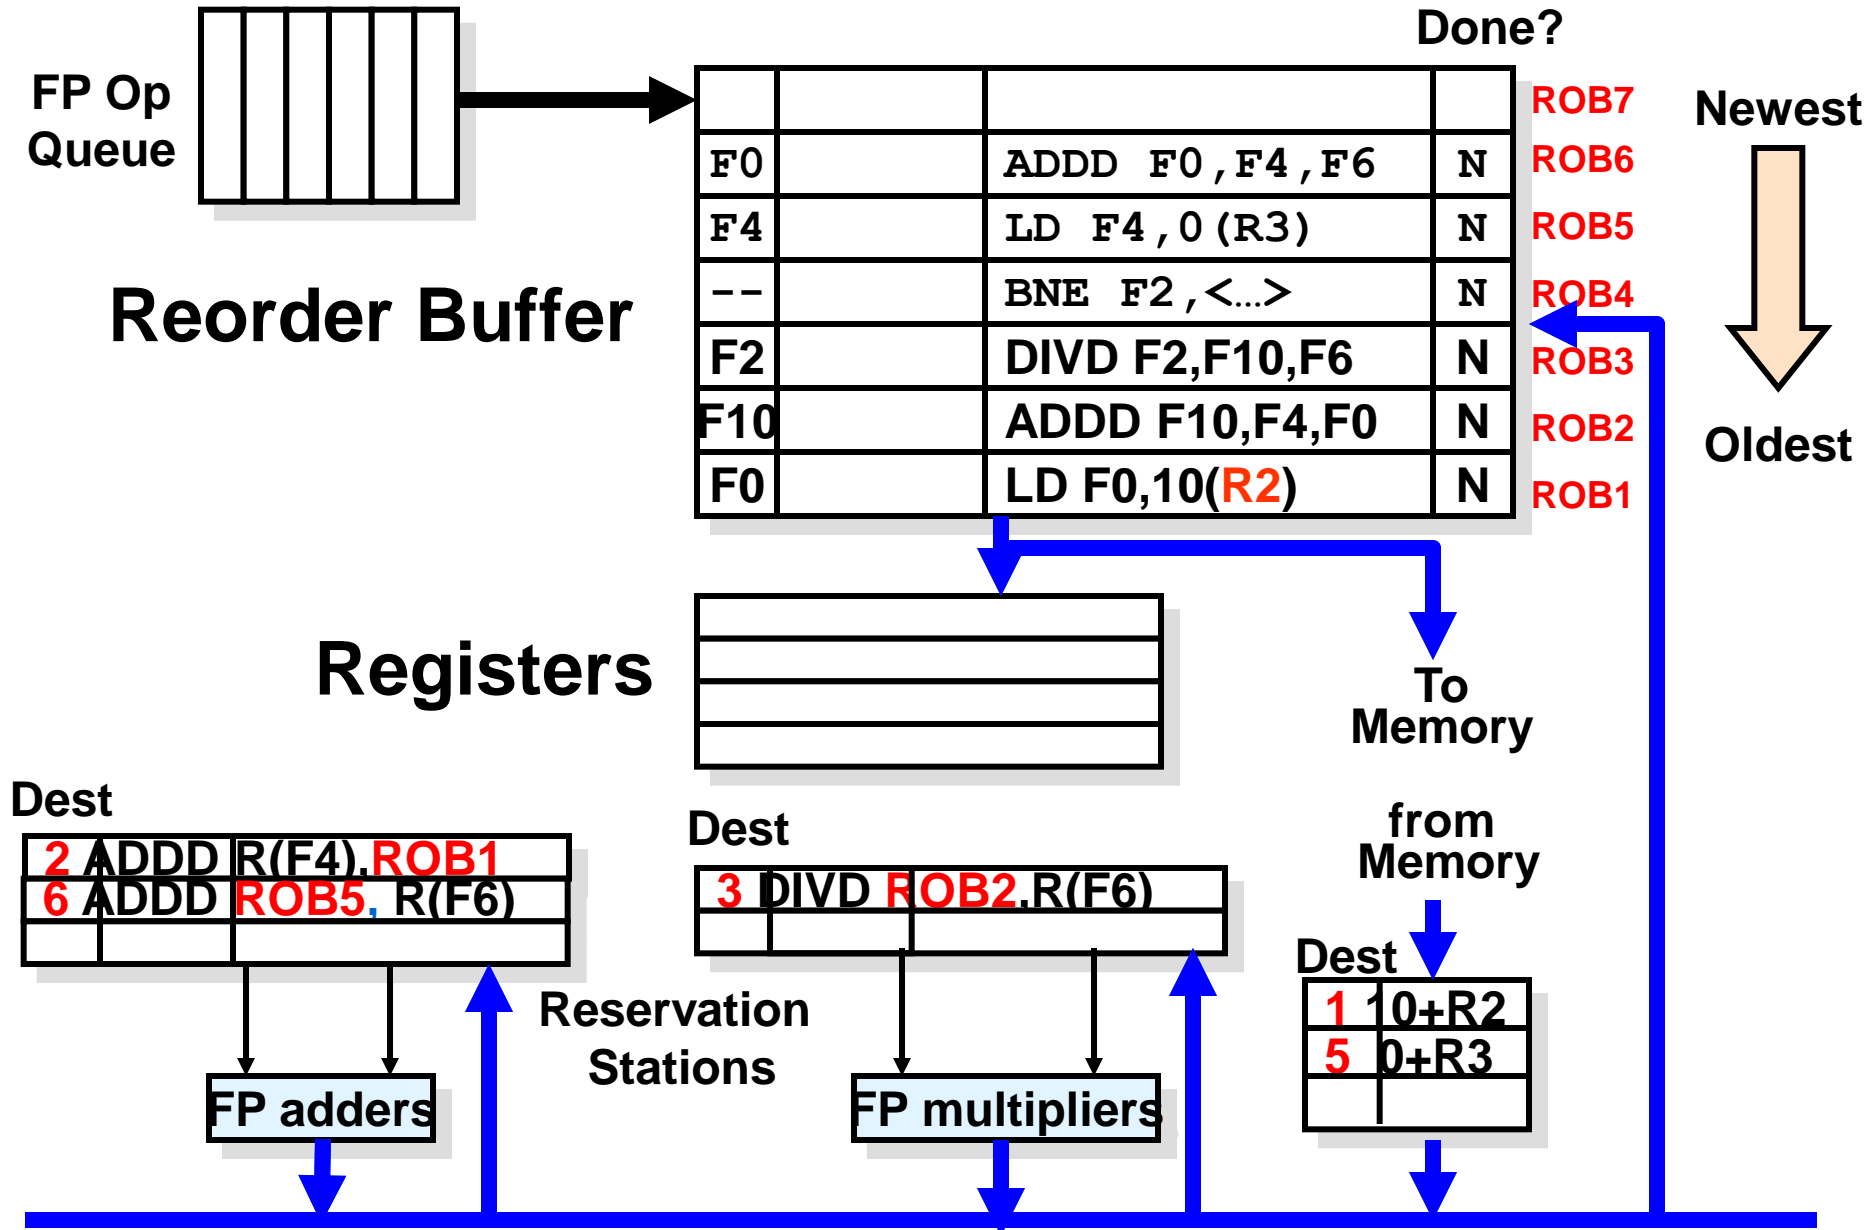
Tomasulo With Reorder buffer:



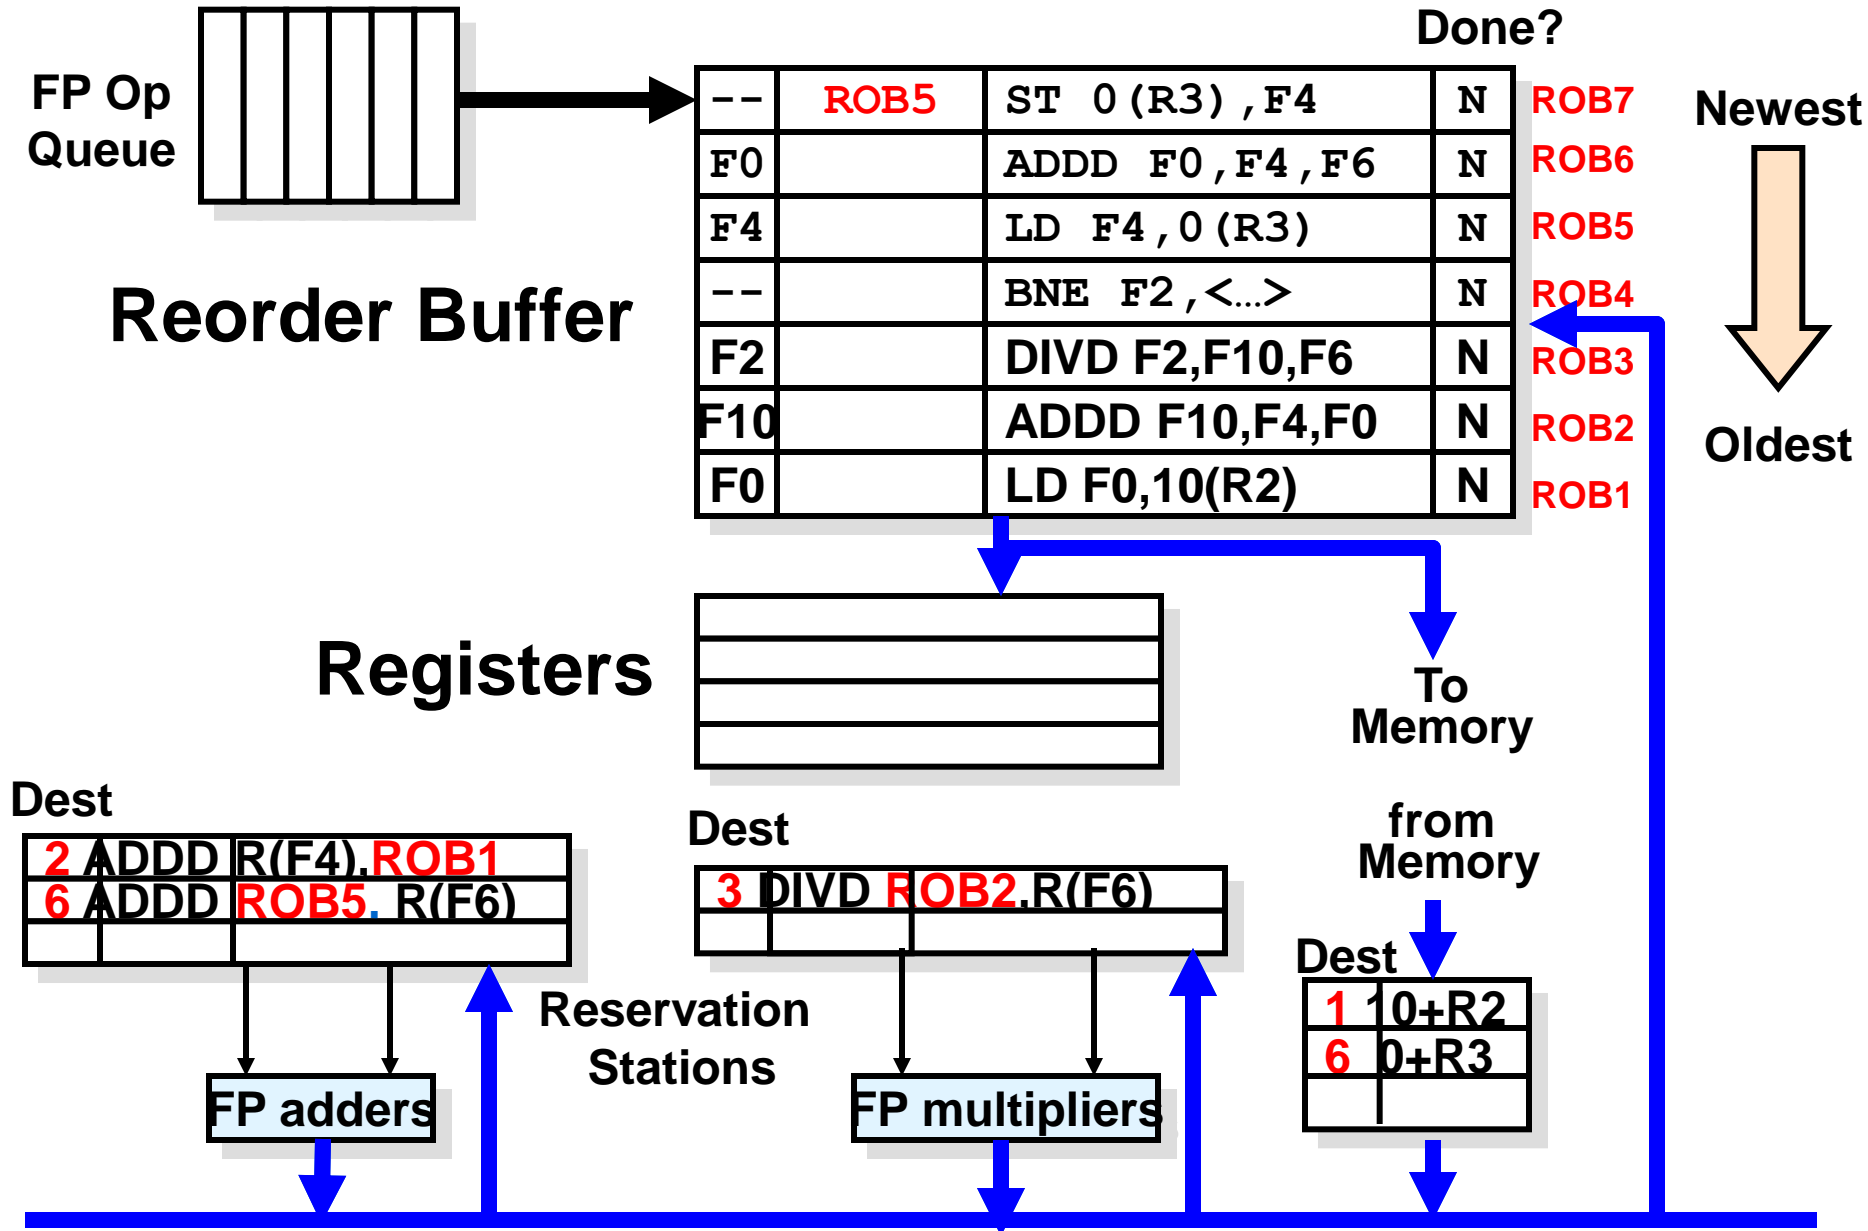
Tomasulo With Reorder buffer:



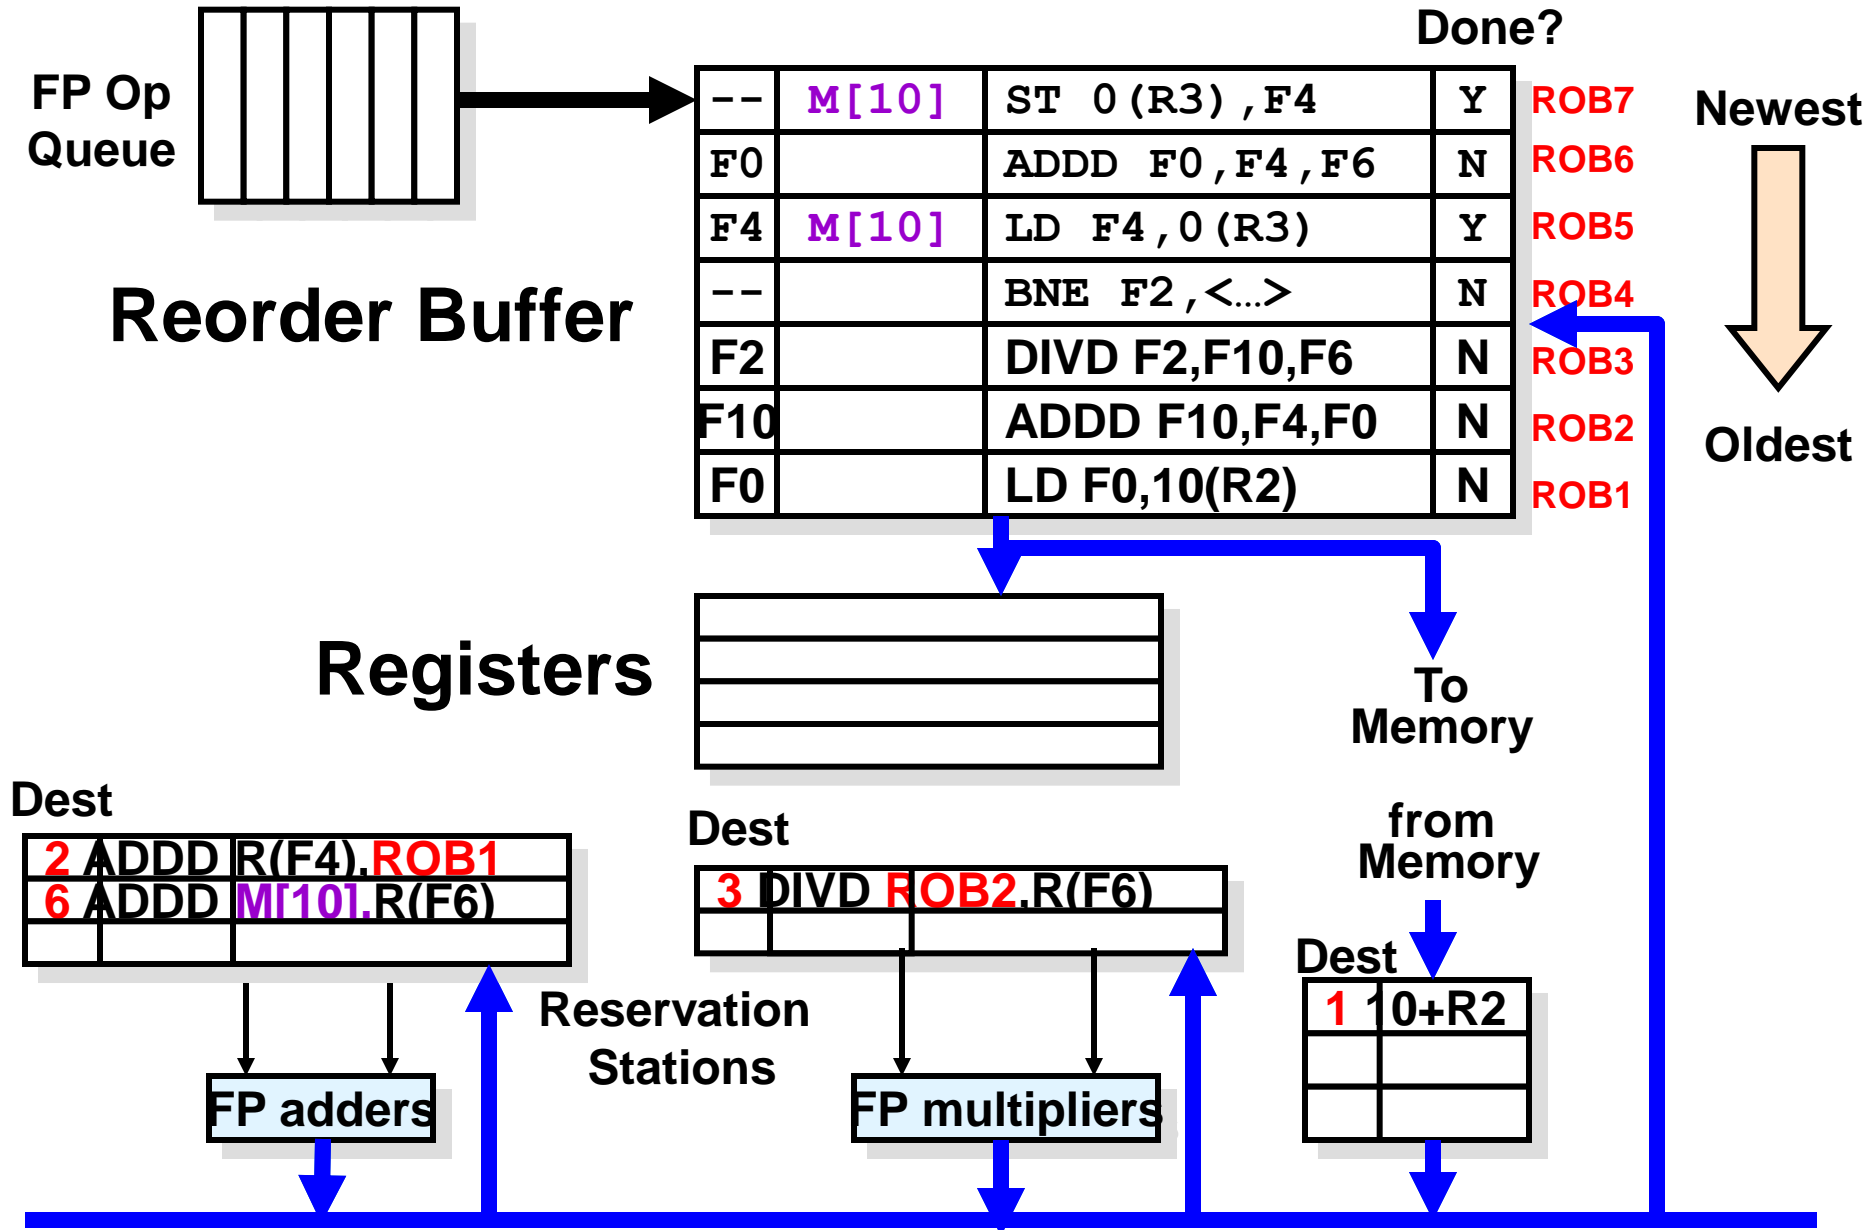
Tomasulo With Reorder buffer:



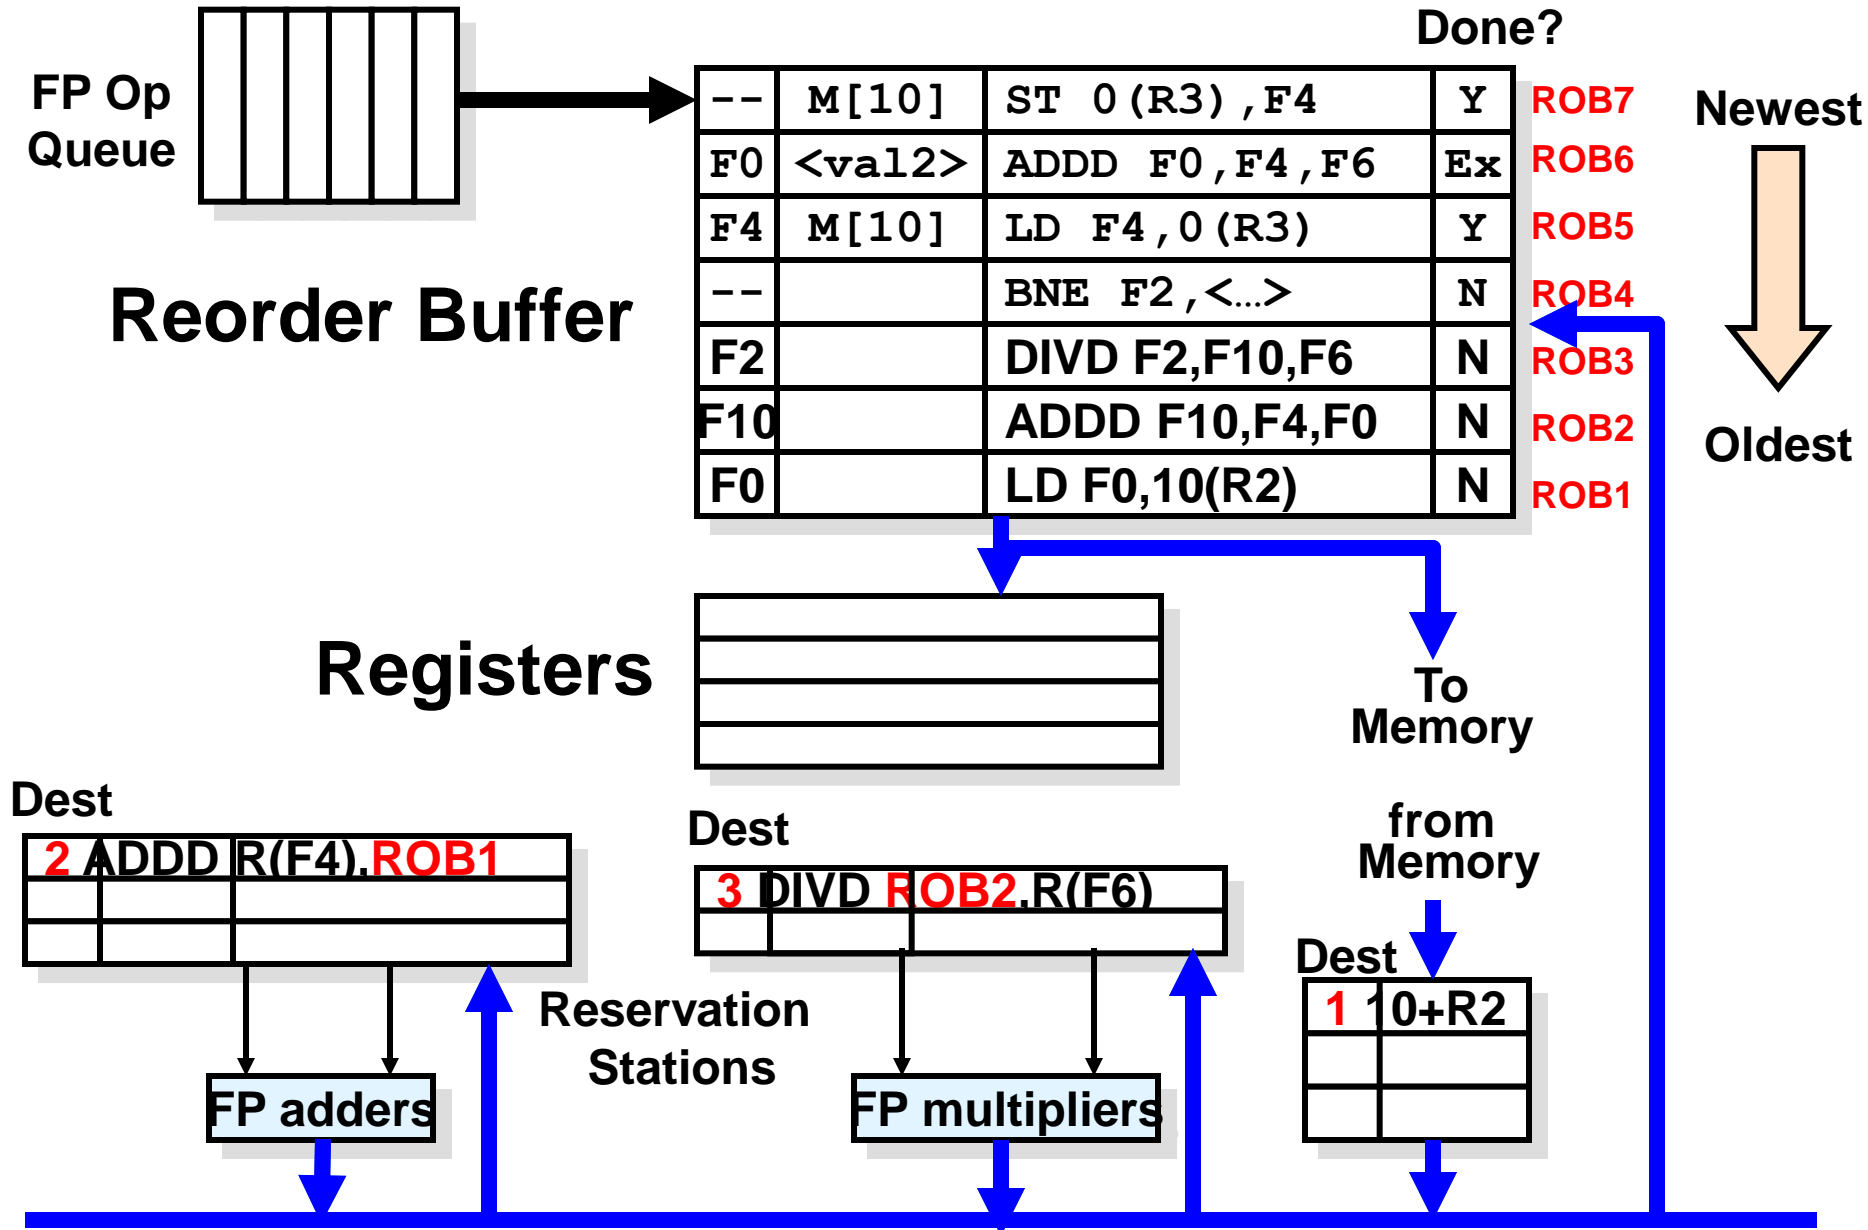
Tomasulo With Reorder buffer:



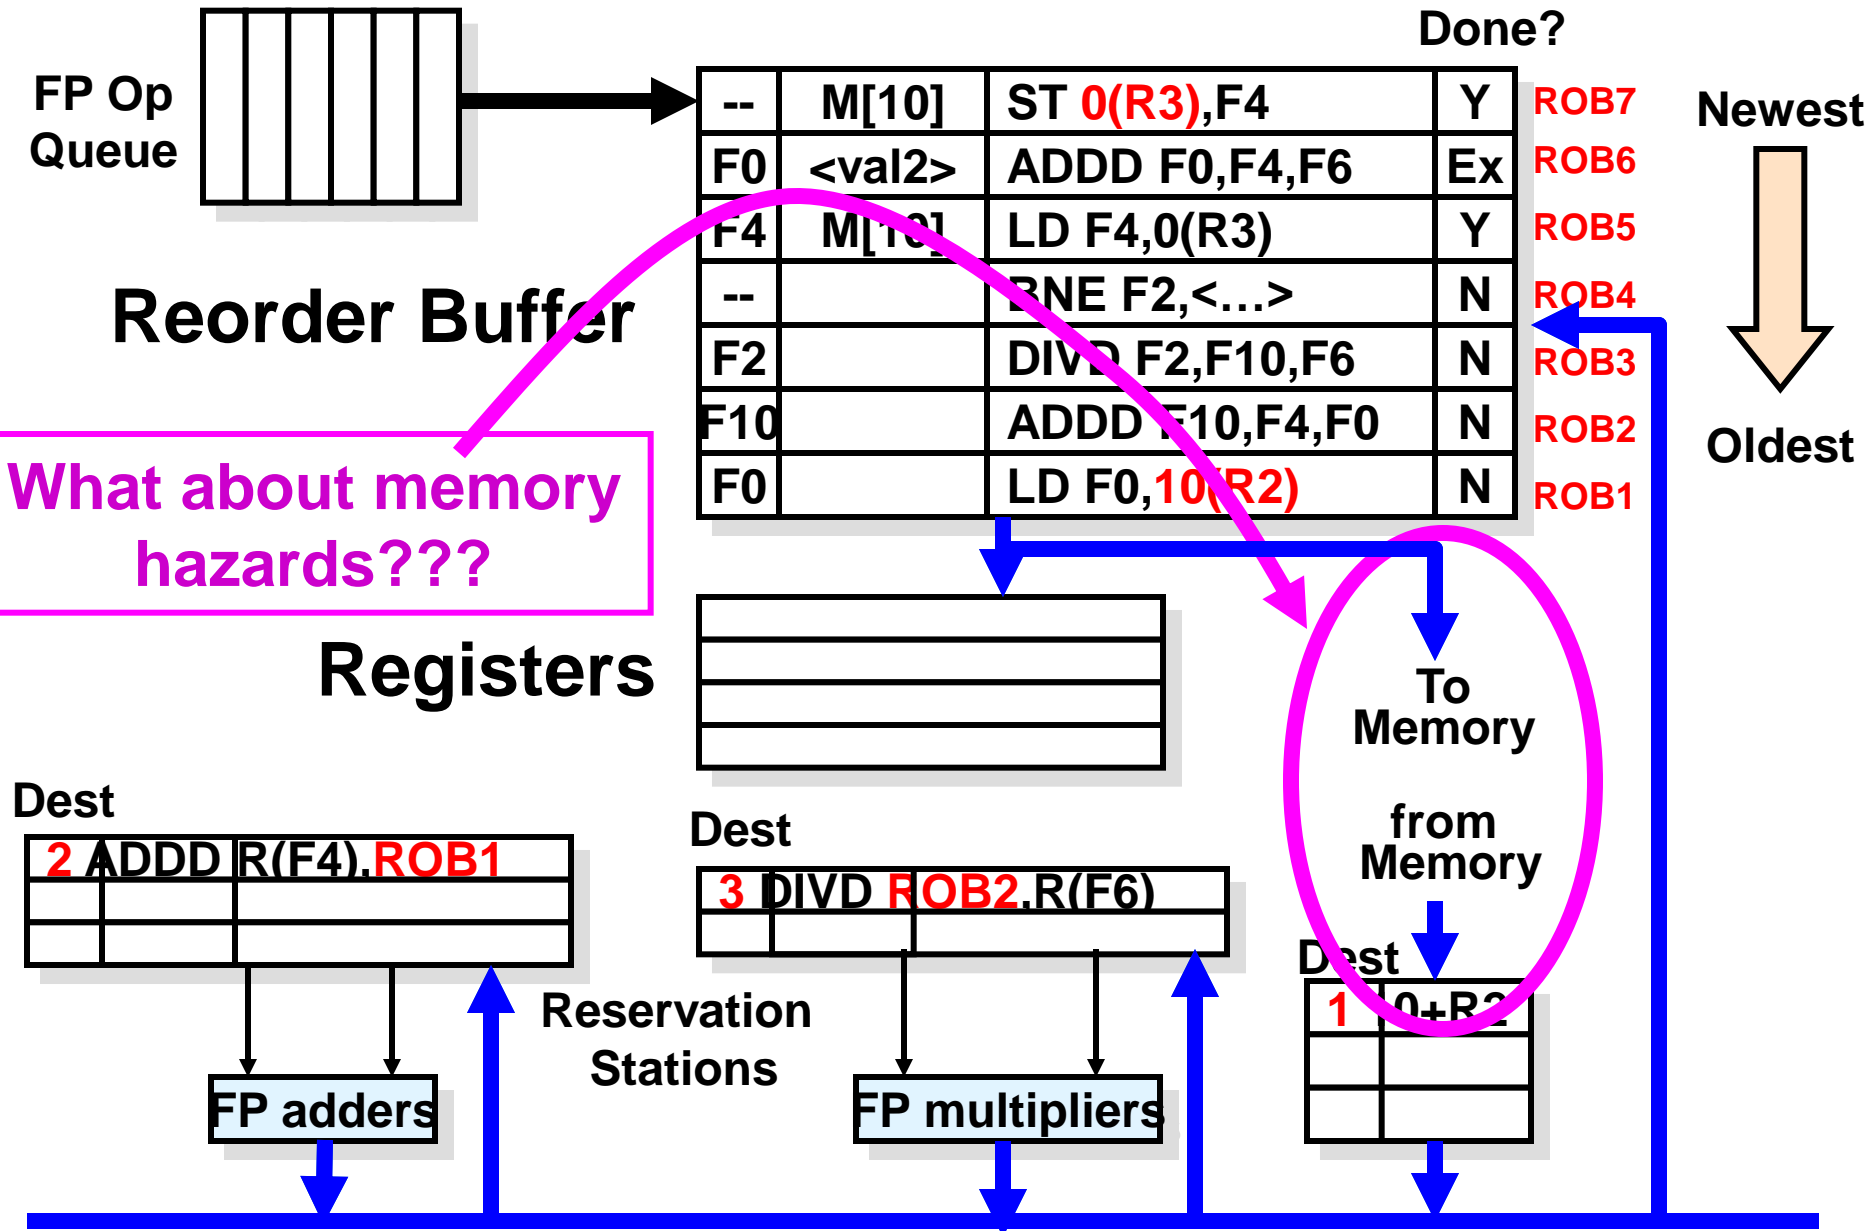
Tomasulo With Reorder buffer:



Tomasulo With Reorder buffer:



Tomasulo With Reorder buffer:



new problem--memory disambiguation

❑ Question: Given a load that follows a store in program order, are the two related?

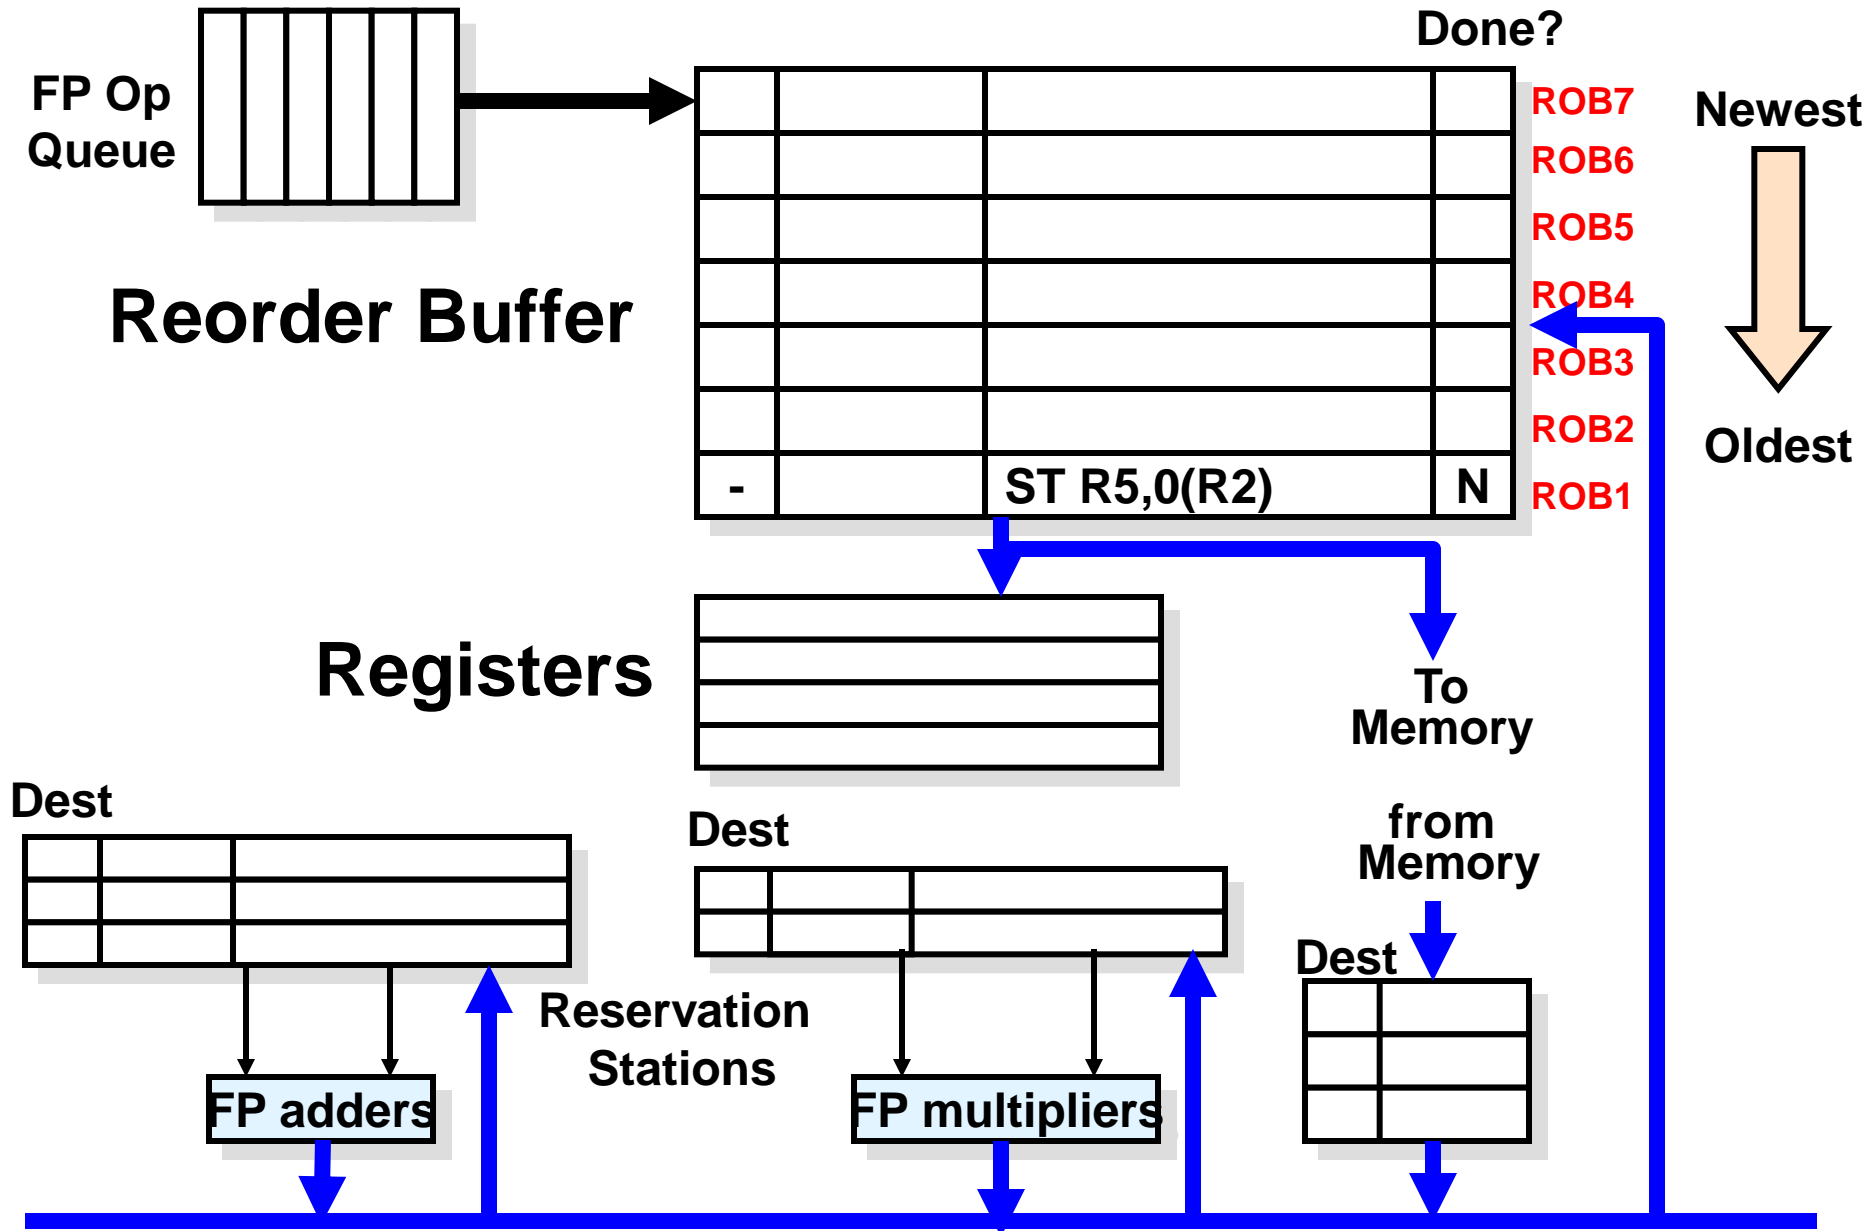
- (Alternatively: is there a RAW hazard between the store and the load)?

Eg: ST 0(R2), R5
 LD R4, 0(R3)

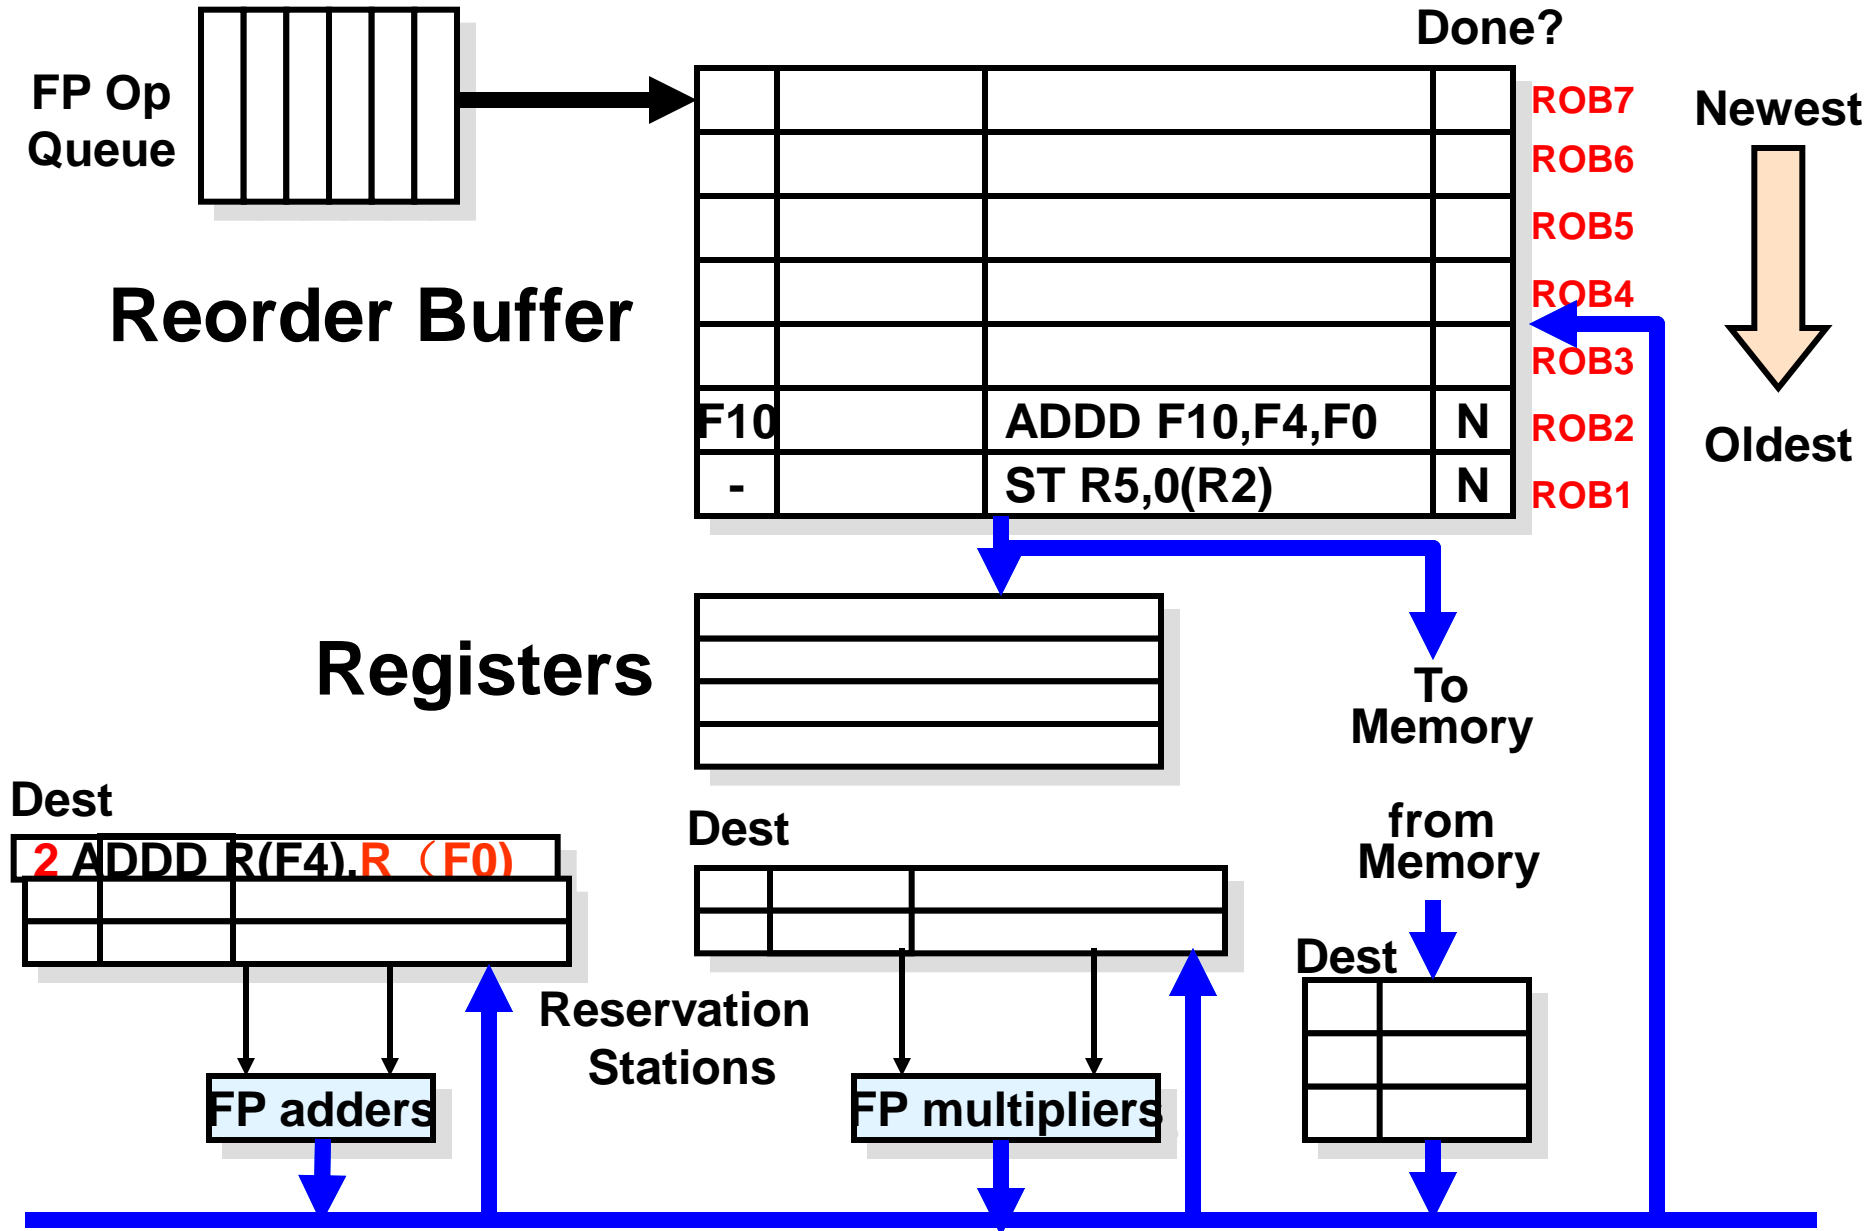
❑ Can we go ahead and start the load early?

- Store address could be delayed for a long time by some calculation that leads to R2 (divide?).
- We might want to issue/begin execution of both operations in same cycle.

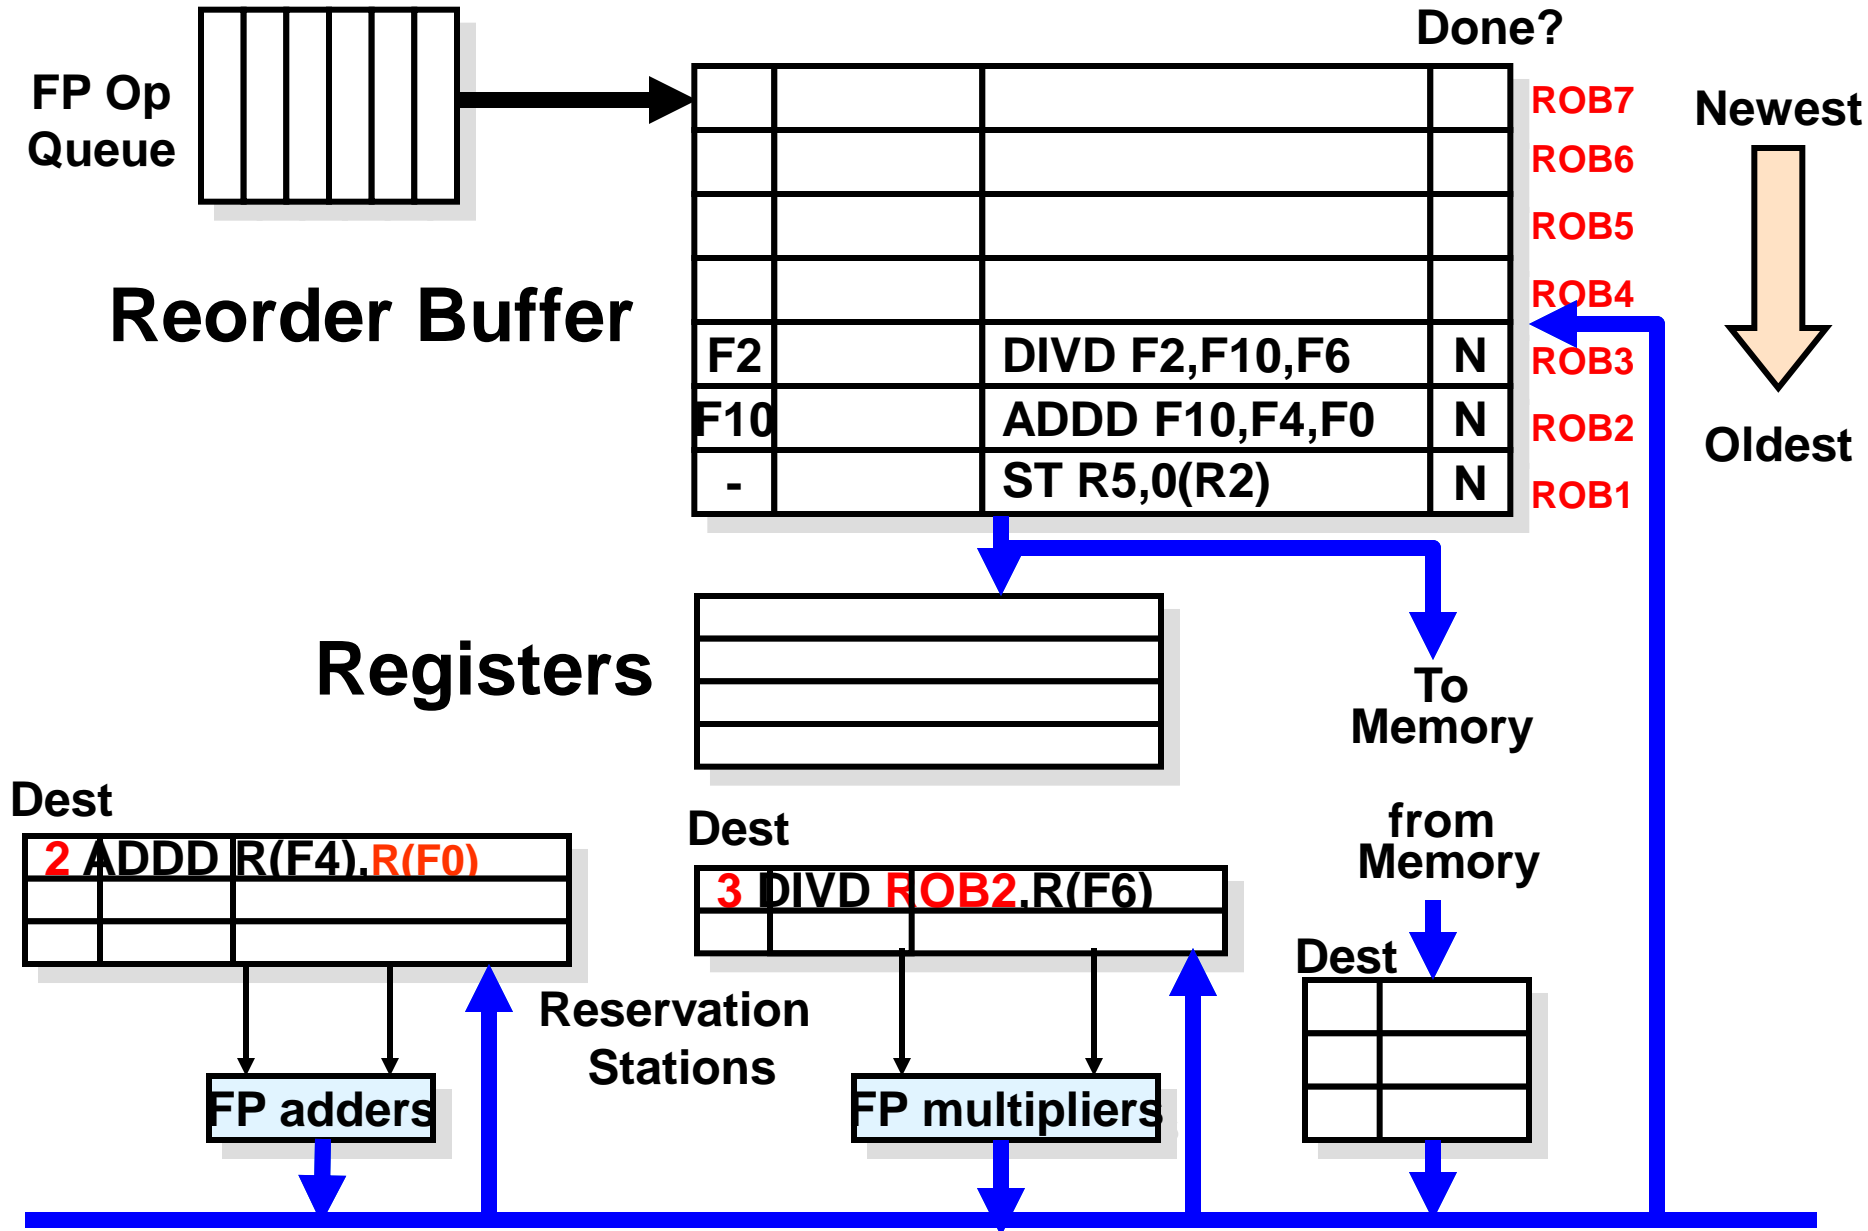
Tomasulo With Reorder buffer:



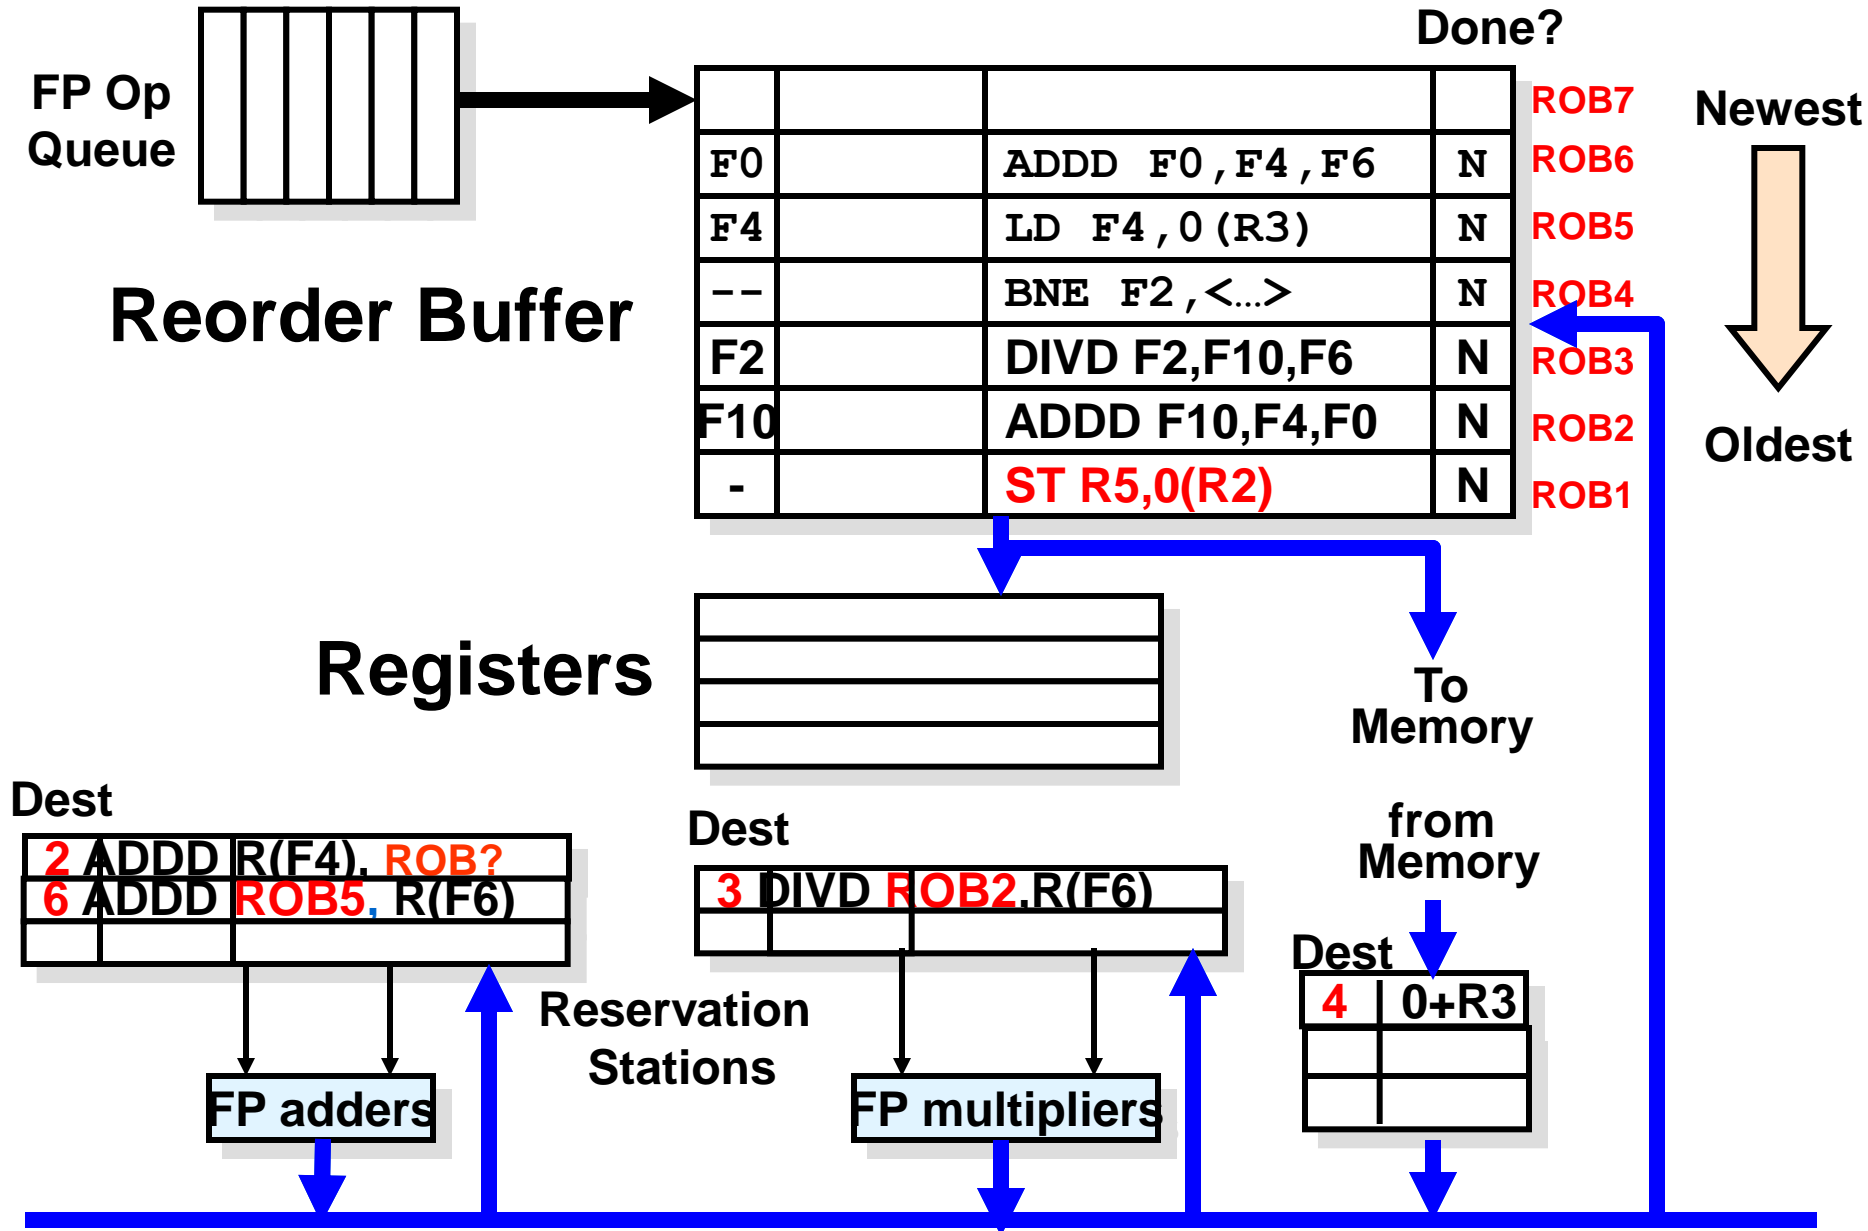
Tomasulo With Reorder buffer:



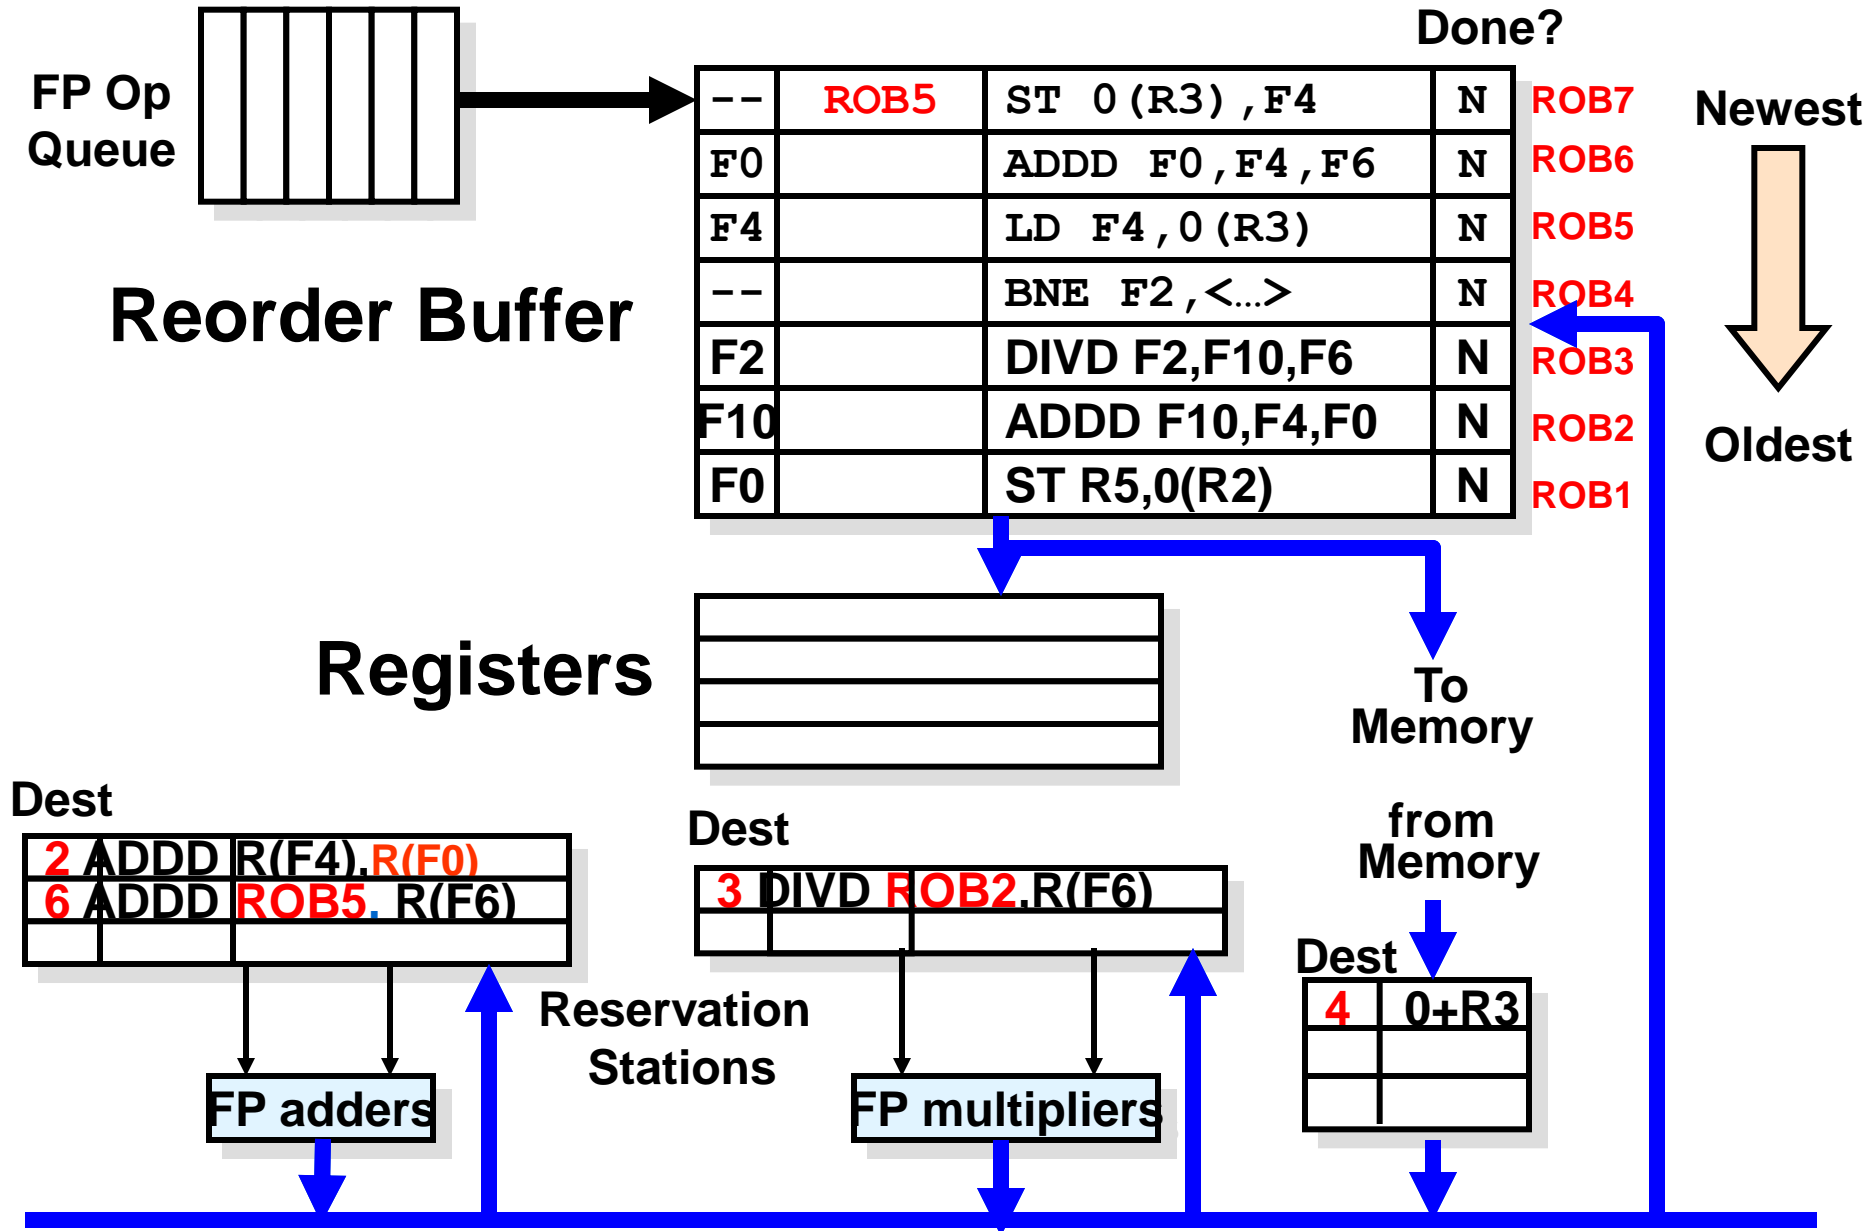
Tomasulo With Reorder buffer:



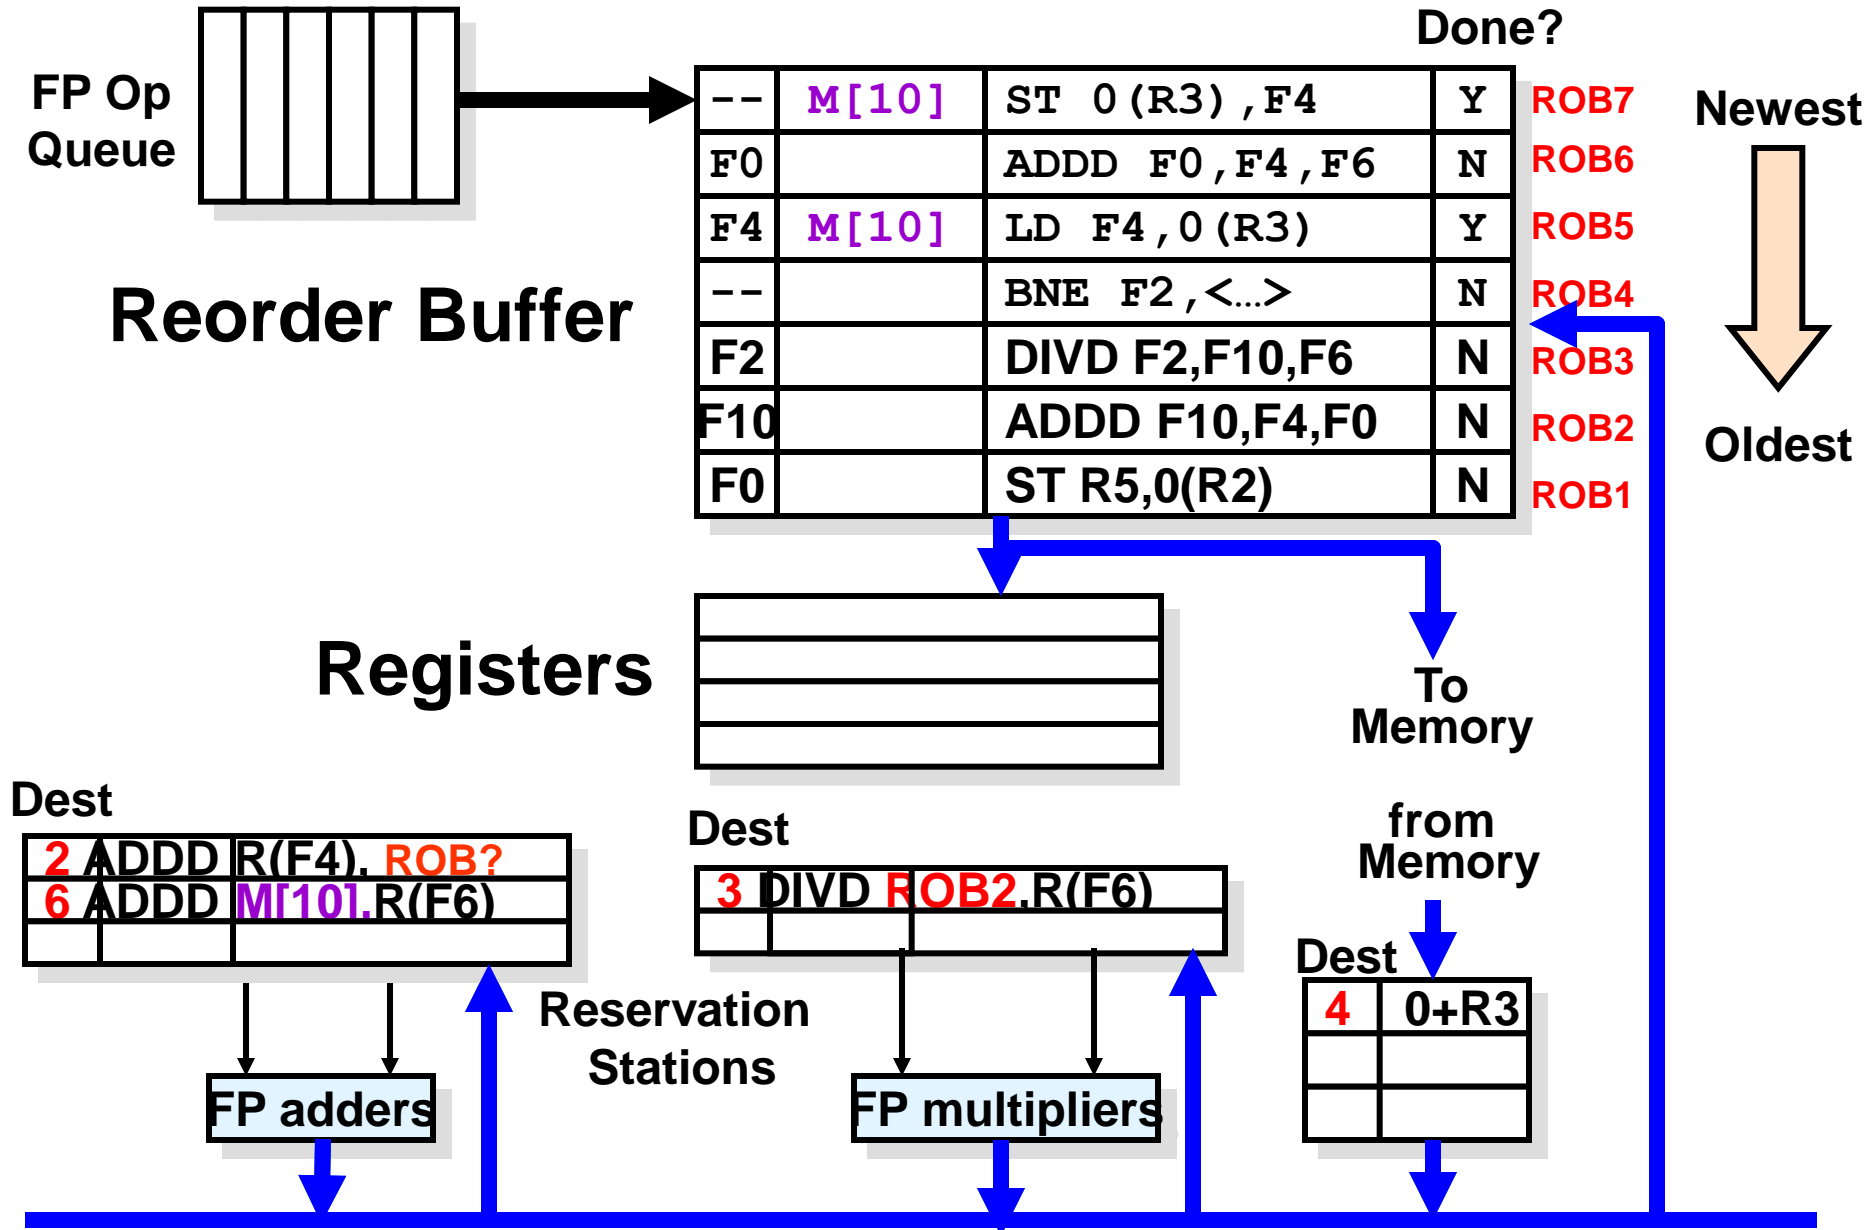
Tomasulo With Reorder buffer:



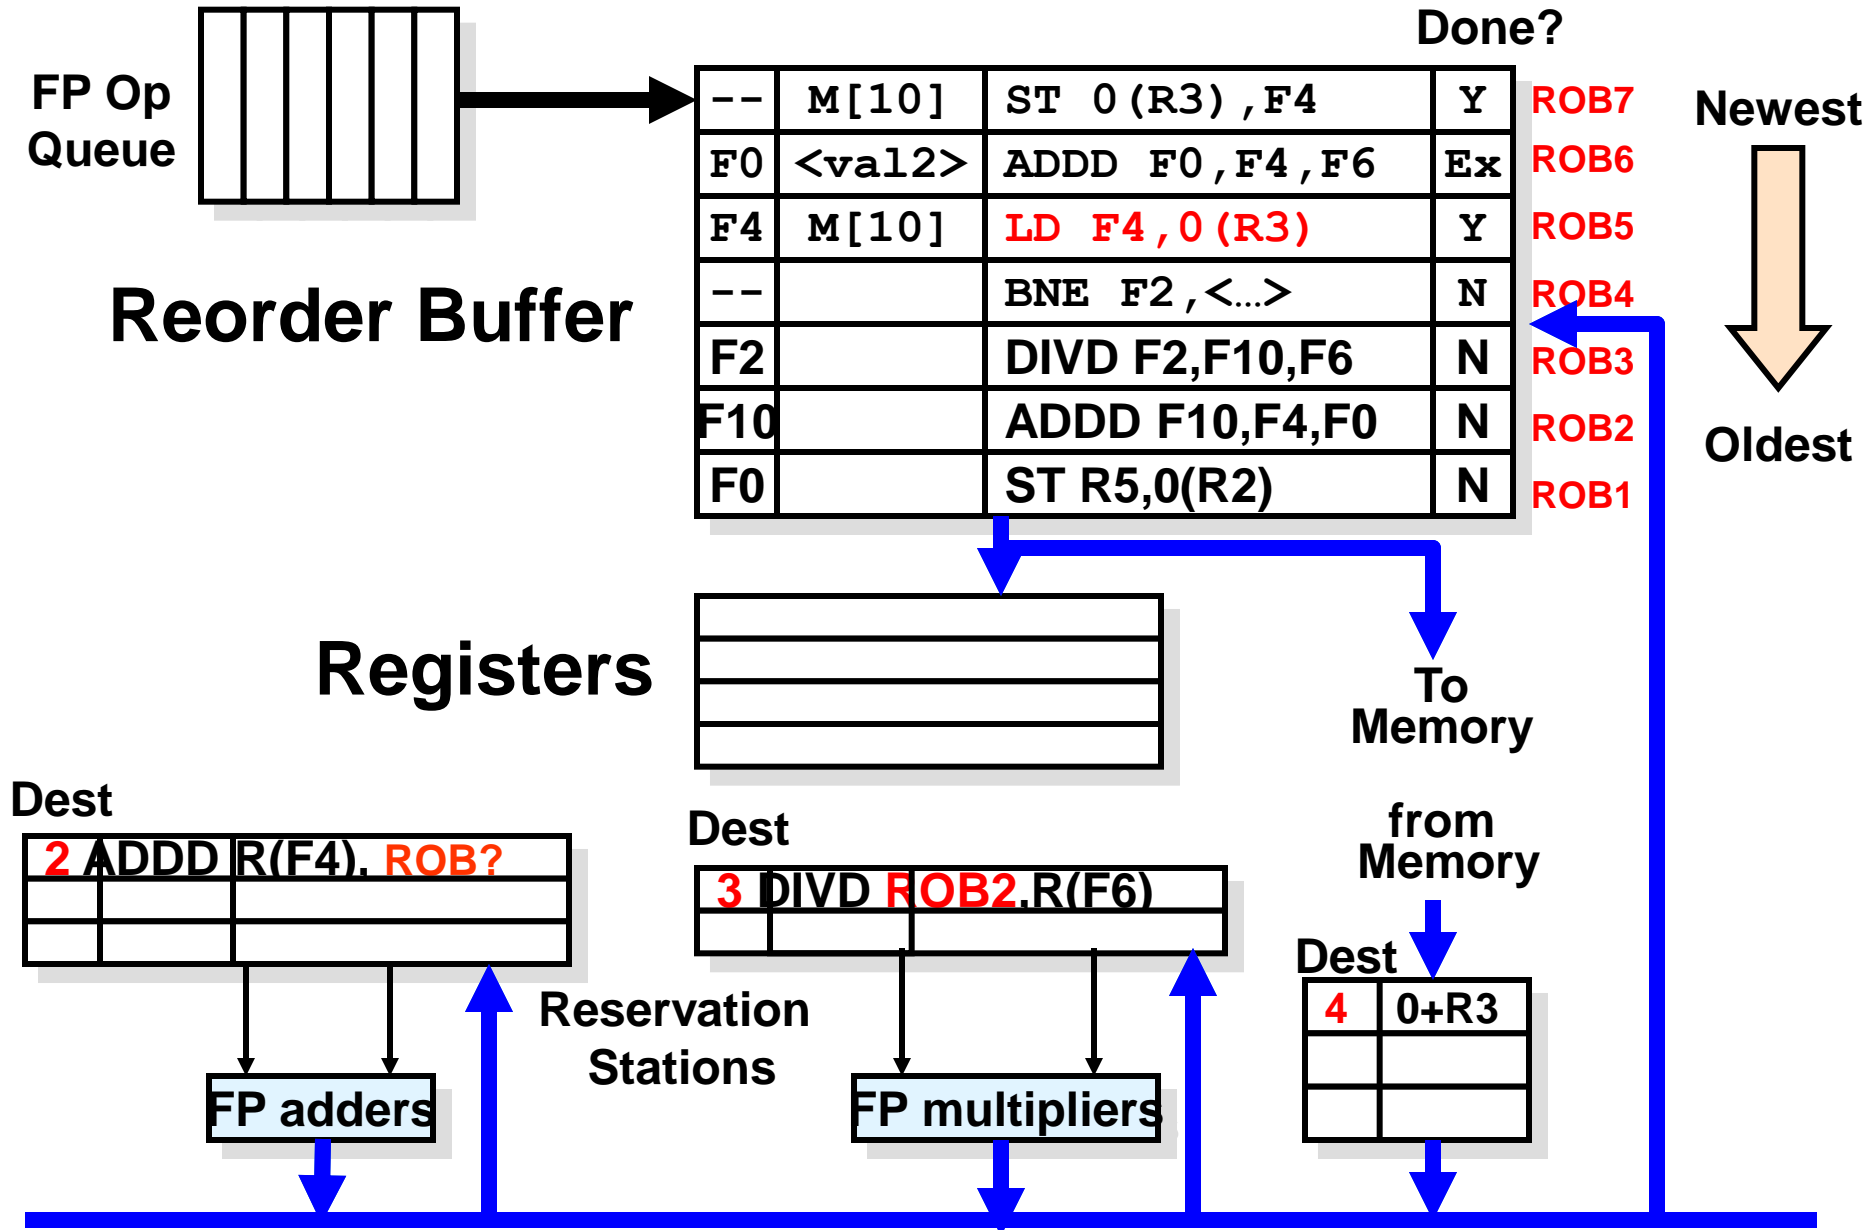
Tomasulo With Reorder buffer:



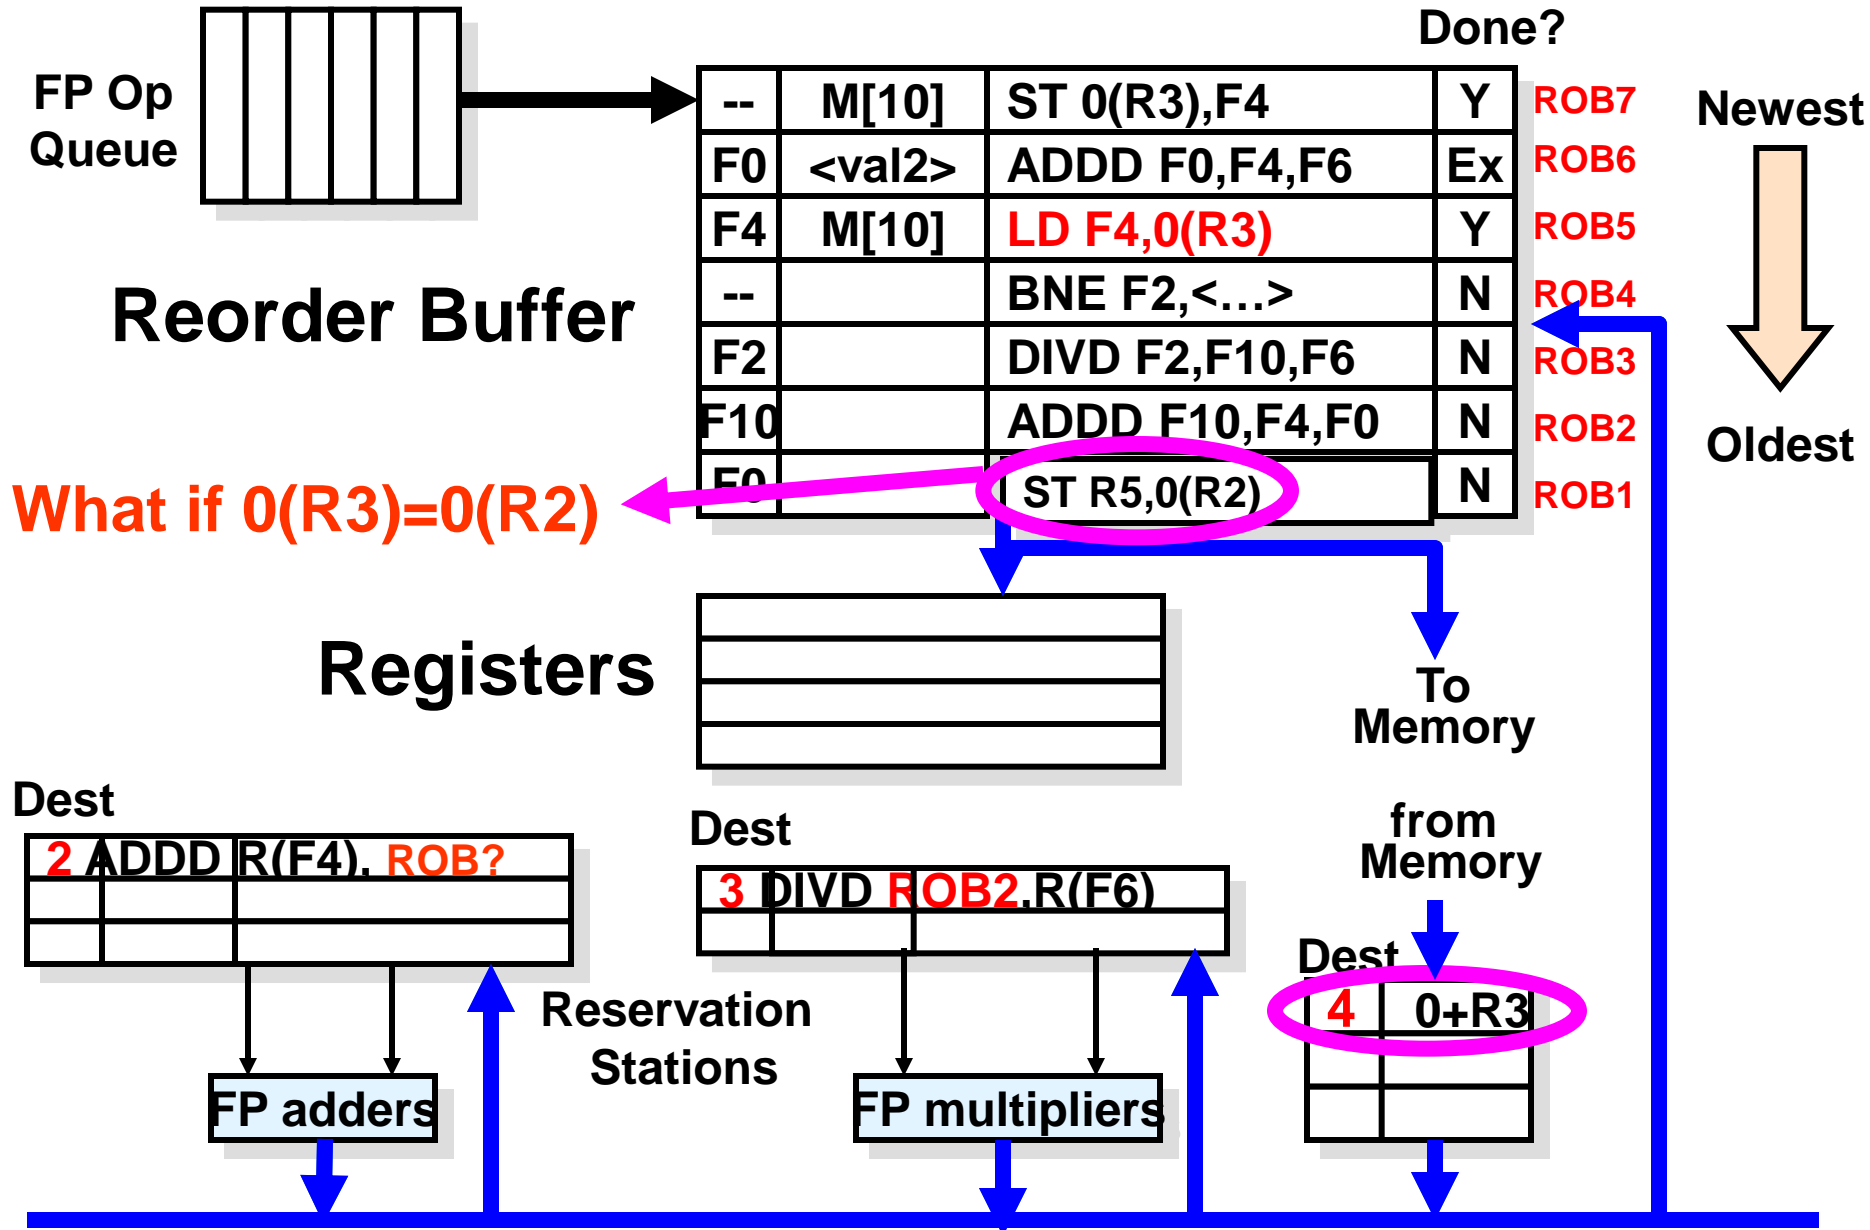
Tomasulo With Reorder buffer:



Tomasulo With Reorder buffer:



Tomasulo With Reorder buffer:



Memory Disambiguation:

Sorting out RAW Hazards in memory

•Question: Given a load that follows a store in program order, are the two related?

–(Alternatively: is there a RAW hazard between the store and the load)?

Eg: ST 0(R2), R5
 LD R6, 0(R3)

•Can we go ahead and start the load early?

–Store address could be delayed for a long time by some calculation that leads to R2 (divide?).

–We might want to issue/begin execution of both operations in same cycle.

–Today: Answer is that we are not allowed to start load until we know that address $0(R2) \neq 0(R3)$

Hardware Support for Memory Disambiguation

- ❑ Need buffer to keep track of all outstanding stores to memory, in program order.
 - Keep track of address (when becomes available) and value (when becomes available)
 - FIFO ordering: will retire stores from this buffer in program order
- ❑ When issuing a load, record current head of store queue (know which stores are ahead of you).

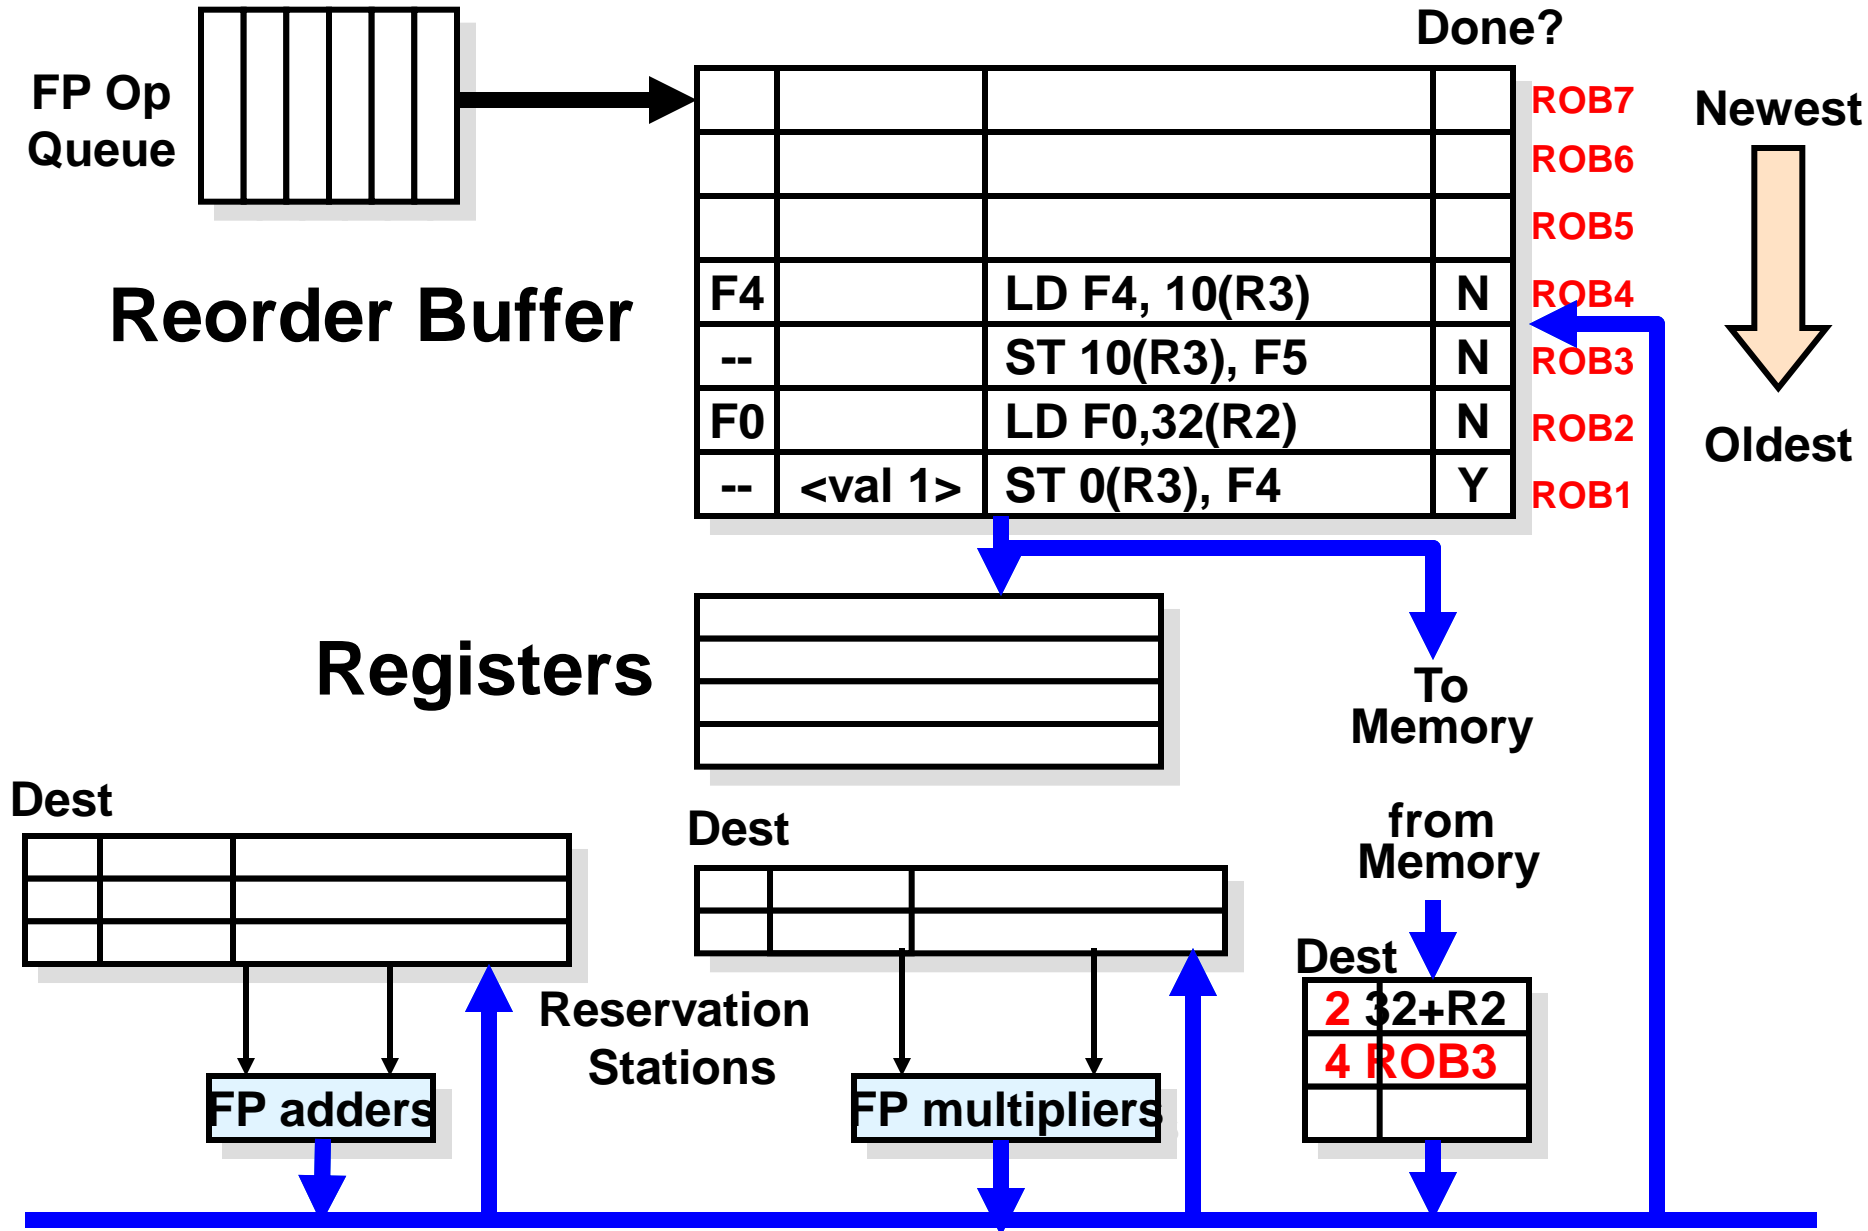
Hardware Support for Memory Disambiguation(2)

❑ When have address for load, check store queue:

- If *any* store prior to load is waiting for its address, stall load.
- If load address matches earlier store address (associative lookup), then we have a *memory-induced RAW hazard*:
 - store value available \Rightarrow return value
 - store value not available \Rightarrow return ROB number of source
- Otherwise, send out request to memory

❑ Actual stores commit in order, so no worry about WAR/WAW hazards through memory.

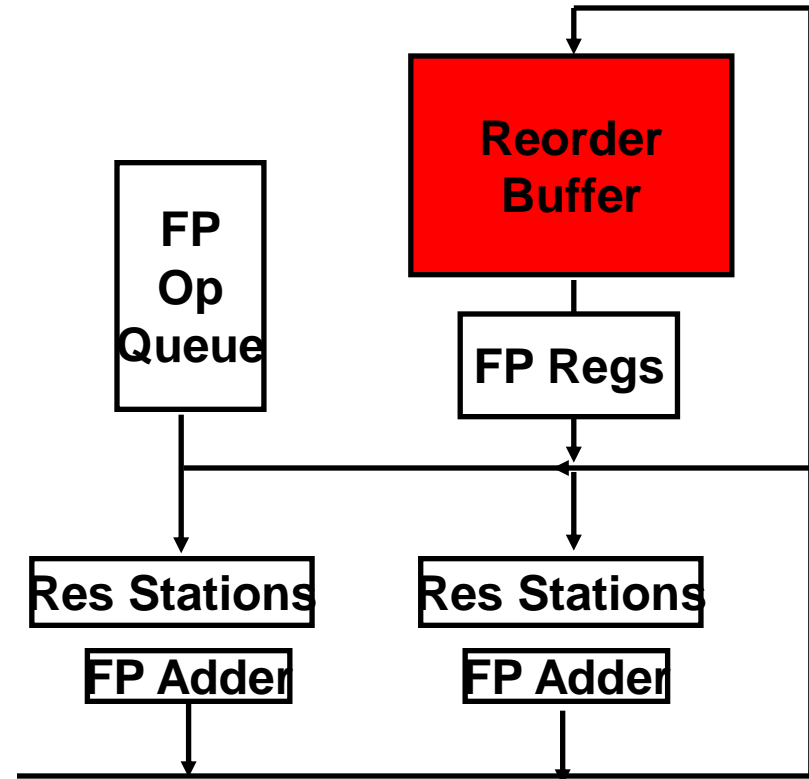
Memory Disambiguation:



HW support for precise interrupts

❑ Need HW buffer for results of uncommitted instructions: **reorder buffer**

- 3 fields: instr, destination, value
- Use **reorder buffer number** instead of reservation station when execution completes
- Supplies operands between execution complete & commit
- (Reorder buffer can be operand source => more registers like RS)
- Once instruction commits, result is put into register
- As a result, easy to undo speculated instructions on mispredicted branches or exceptions



| | 记分牌 | Tomasulo | 带投机的Tomasulo |
|--------|-------------------------------------------------------------------------------|--------------------------------------------------------------------------|-------------------------------------------------------------------------------------------------------------|
| 主要思想 | 动态调度、按序发射、乱序执行、乱序完成、集中控制。 | 动态调度、按序发射、乱序执行、乱序完成分散控制。增加RS（保留站）,CDB。 | 转移预测, 动态调度、投机执行、按序完成。在RS基础上再增加ROB(reorder buffer)。 |
| 解决何种竞争 | RAW数据竞争。 WAW,WAR时stall | 解决数据竞争,避免WW,WAR竞争;允许硬件支持的循环展开，不局限于程序基本块。 | 解决数据竞争,避免WW,WAR竞争;从数据竞争扩充到控制竞争。通过预测减少控制竞争stall。 |
| 指令节拍 | ISSUE：停顿于结构竞争和WAW竞争。 RO：等待两个操作数就绪，动态解决RAW。 EX：和普通流水线一样。 WB：停顿于WAR竞争。 | ISSUE：仅停顿于结构竞争（RS部件）。 EX：利用保留站（RS）进行换名，解决了WAW和WAR问题（即名字相关）。 WB：回写。 | ISSUE：仅停顿于结构竞争（RS部件和ROB部件）。 EX：等待操作数就绪，执行。 WB：结果写到ROB，其它指令可以引用该临时结果，但不更新寄存器。 COMMIT：保证按序完成，保证精确中断。 |
| 关键部件 | 记分牌、集中控制、指令状态、功能部件状态、寄存器结果状态。 | CDB(通用数据总线)：执行结果通过CDB传递，及时。 RS(保留站)：换名功能以及暂存临时数据。 | ROB(REORDER BUFFER)重构序缓冲器： 对移到顶部的指令进行判断。。。ROB负责寄存器重命名、暂存执行后和交付前的结果、 |
| 其它 | 从静态调度到动态调度的过渡算法。 | 一旦指令将临时结果放在保留站，则其它后续的指令都能立即在寄存器组中看见这个结果。 | 某指令将临时结果放在ROB后，其它指令并不能从寄存器组中看到这个结果（也就是说，寄存器组的值并没有更新），而是等待该指令提交（commit）后，寄存器组才更新。 |

ARM处理器体系结构

□ARM流水线的执行顺序：

- 取指令（Fetch）：从存储器读取指令；
- 译码（Decode）：译码以鉴别它是属于哪一条指令；
- 执行（Execute）：将操作数进行组合以得到结果或存储器地址；
- 缓冲/数据（Buffer/data）：如果需要，则访问存储器以存储数据；
- 回写：（Write-back）：将结果写回到寄存器组中；

基于ARMv8的鲲鹏流水线技术



Taishan coreV110 Pipeline Architecture

基于ARMv8的鲲鹏流水线技术

- **Branch**预测和取指流水线解耦设计，取指流水线每拍最多可提供**32Bytes**指令供译码，分支预测流水线可以不受取指流水停顿影响，超前进行预测处理；
- 定浮点流水线分开设计，解除定浮点相互反压，每拍可为后端执行部件提供**4条整型微指令**及**3条浮点微指令**；
- 整型运算单元支持每拍**4条ALU运算**（含**2条跳转**）及**1条乘除运算**；
- 浮点及**SIMD**运算单元支持每拍**2条ARM Neon 128bits**浮点及**SIMD运算**；
- 访存单元支持每拍**2条读或写访存操作**，读操作最快**4拍**完成，每拍访存带宽为**2x128bits**读及**1x128bits**写；